

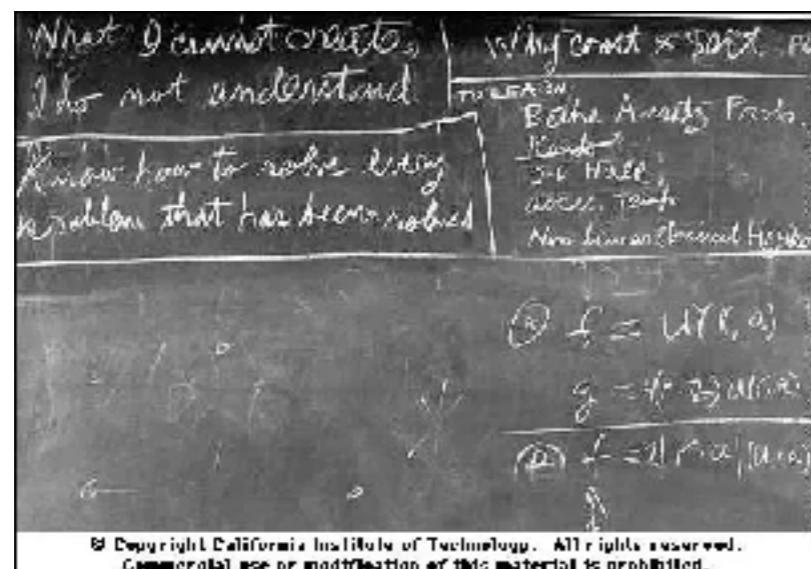
Recent Trends in Deep Learning

Sungjoon Choi
CPSLAB, SNU



Recent Trends in Deep Learning

Generative Model (14)



Domain Adaptation (7)

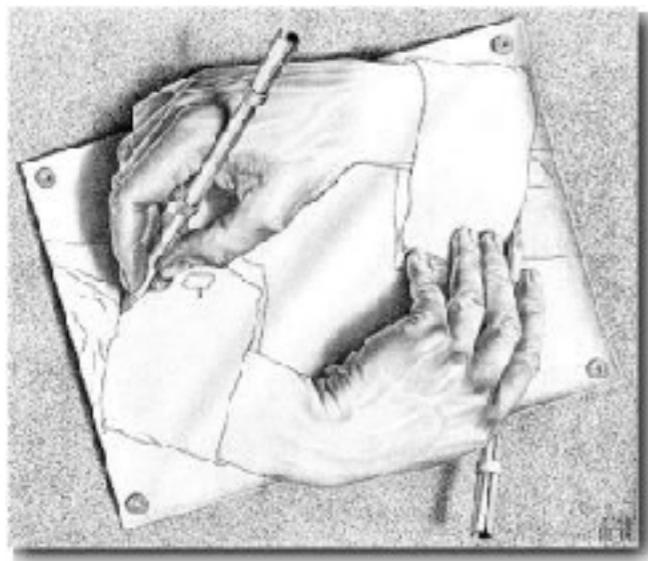
Training Images



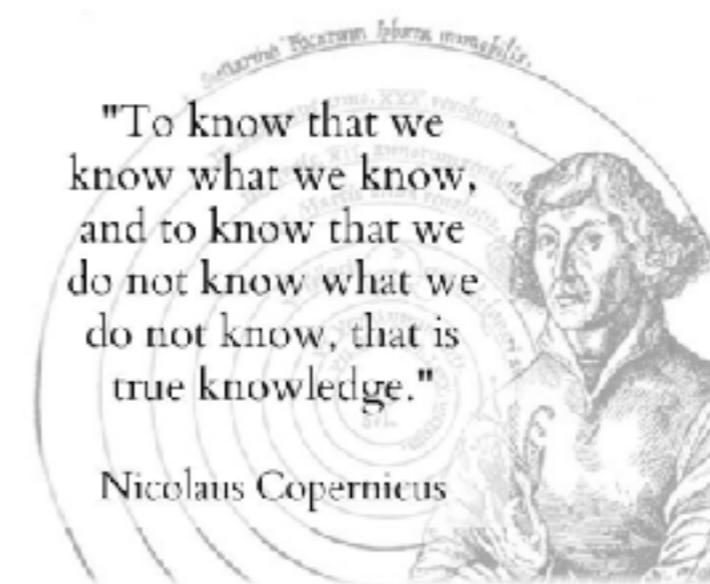
Real World Scenarios



Meta Learning (8)



Uncertainty in Deep Learning (6)



Generative Model



Generative Model

What I cannot create, I do not understand

What I cannot create,
I do not understand.

Know how to solve every
problem that has been solved

Why can't you sort PC

TO LEARN:

- Belief Anxiety Probs.
- Kondo
- 2-D Hull
- vector. Temp
- Non linear Classical Hydro

(1) $f = u(r, \alpha)$

$g = \psi(r, z) u(r, z)$

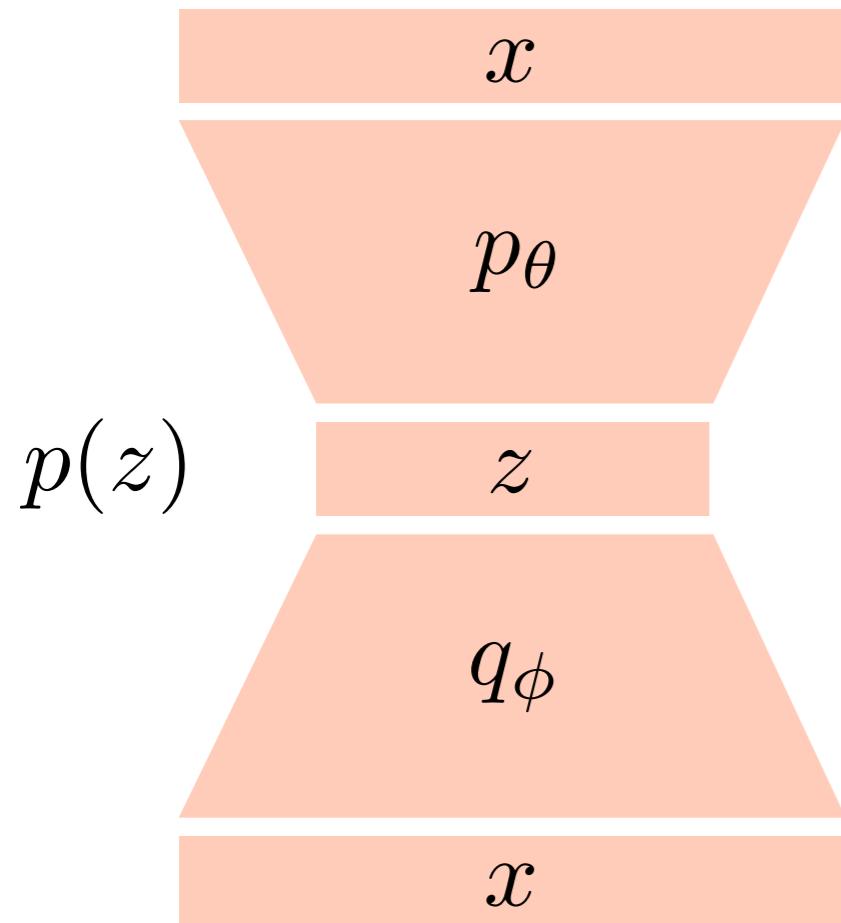
(2) $f = \psi(r, \alpha) u(r, \alpha)$

© Copyright California Institute of Technology. All rights reserved.
Commercial use or modification of this material is prohibited.

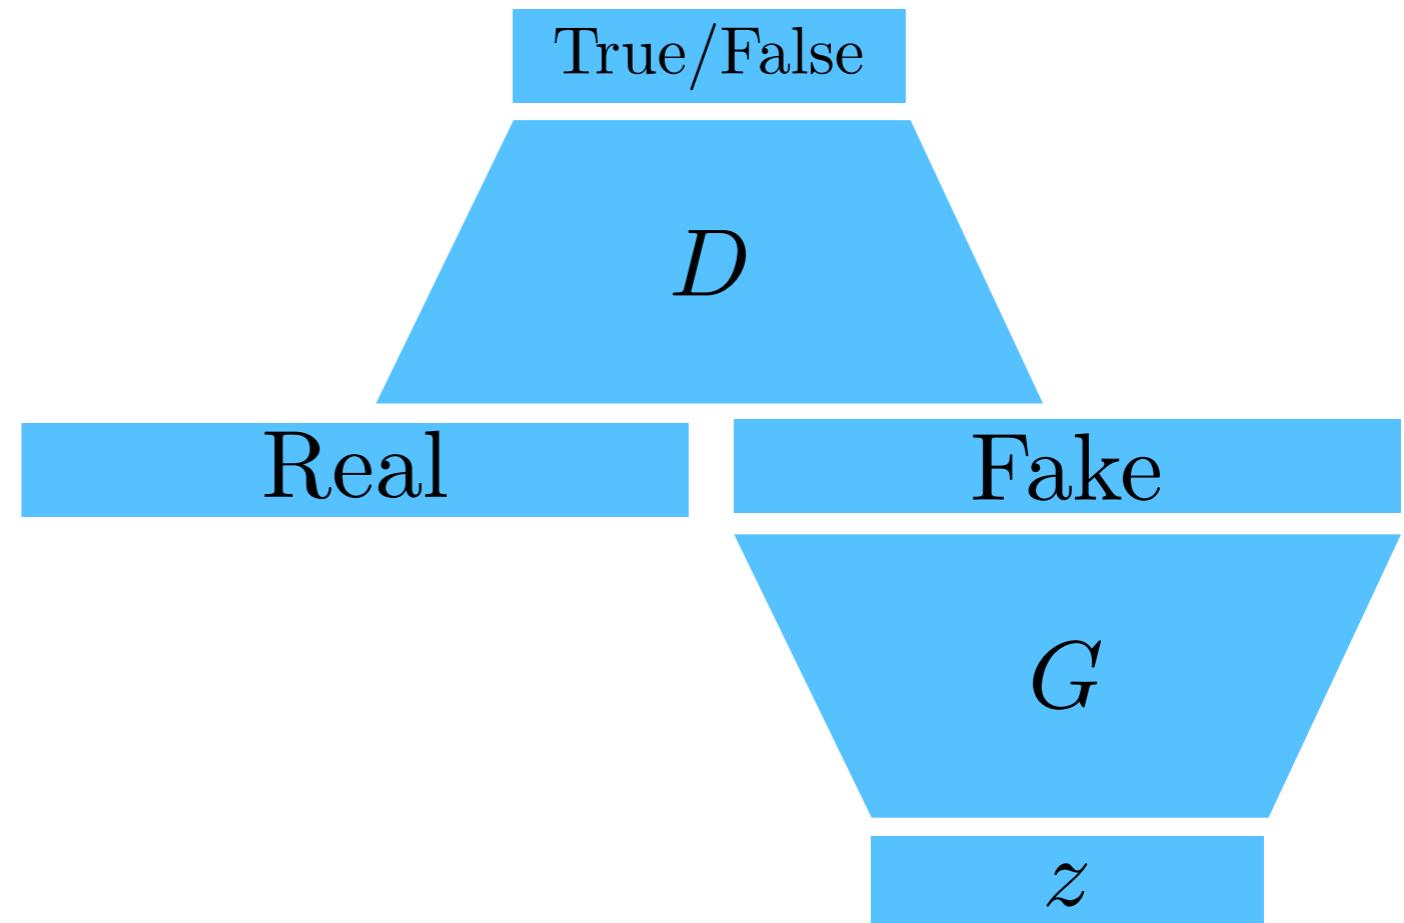


Generative Model

Generative Adversarial Nets vs. Variational Autoencoders



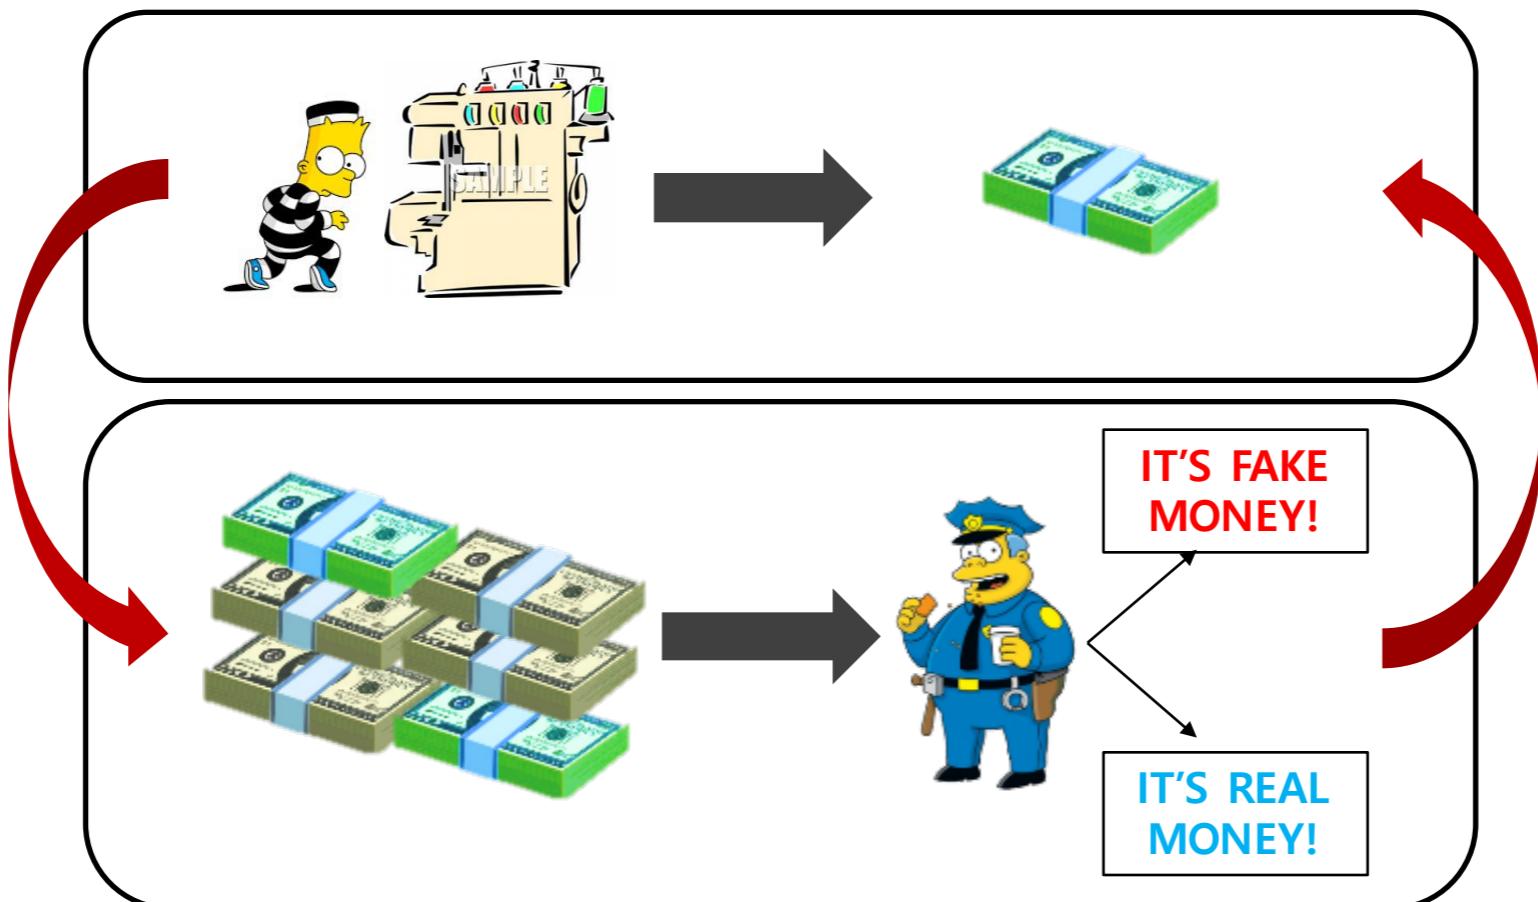
Variational Autoencoder



Generative Adversarial Networks

Generative Model

Generative Adversarial Network (GAN)



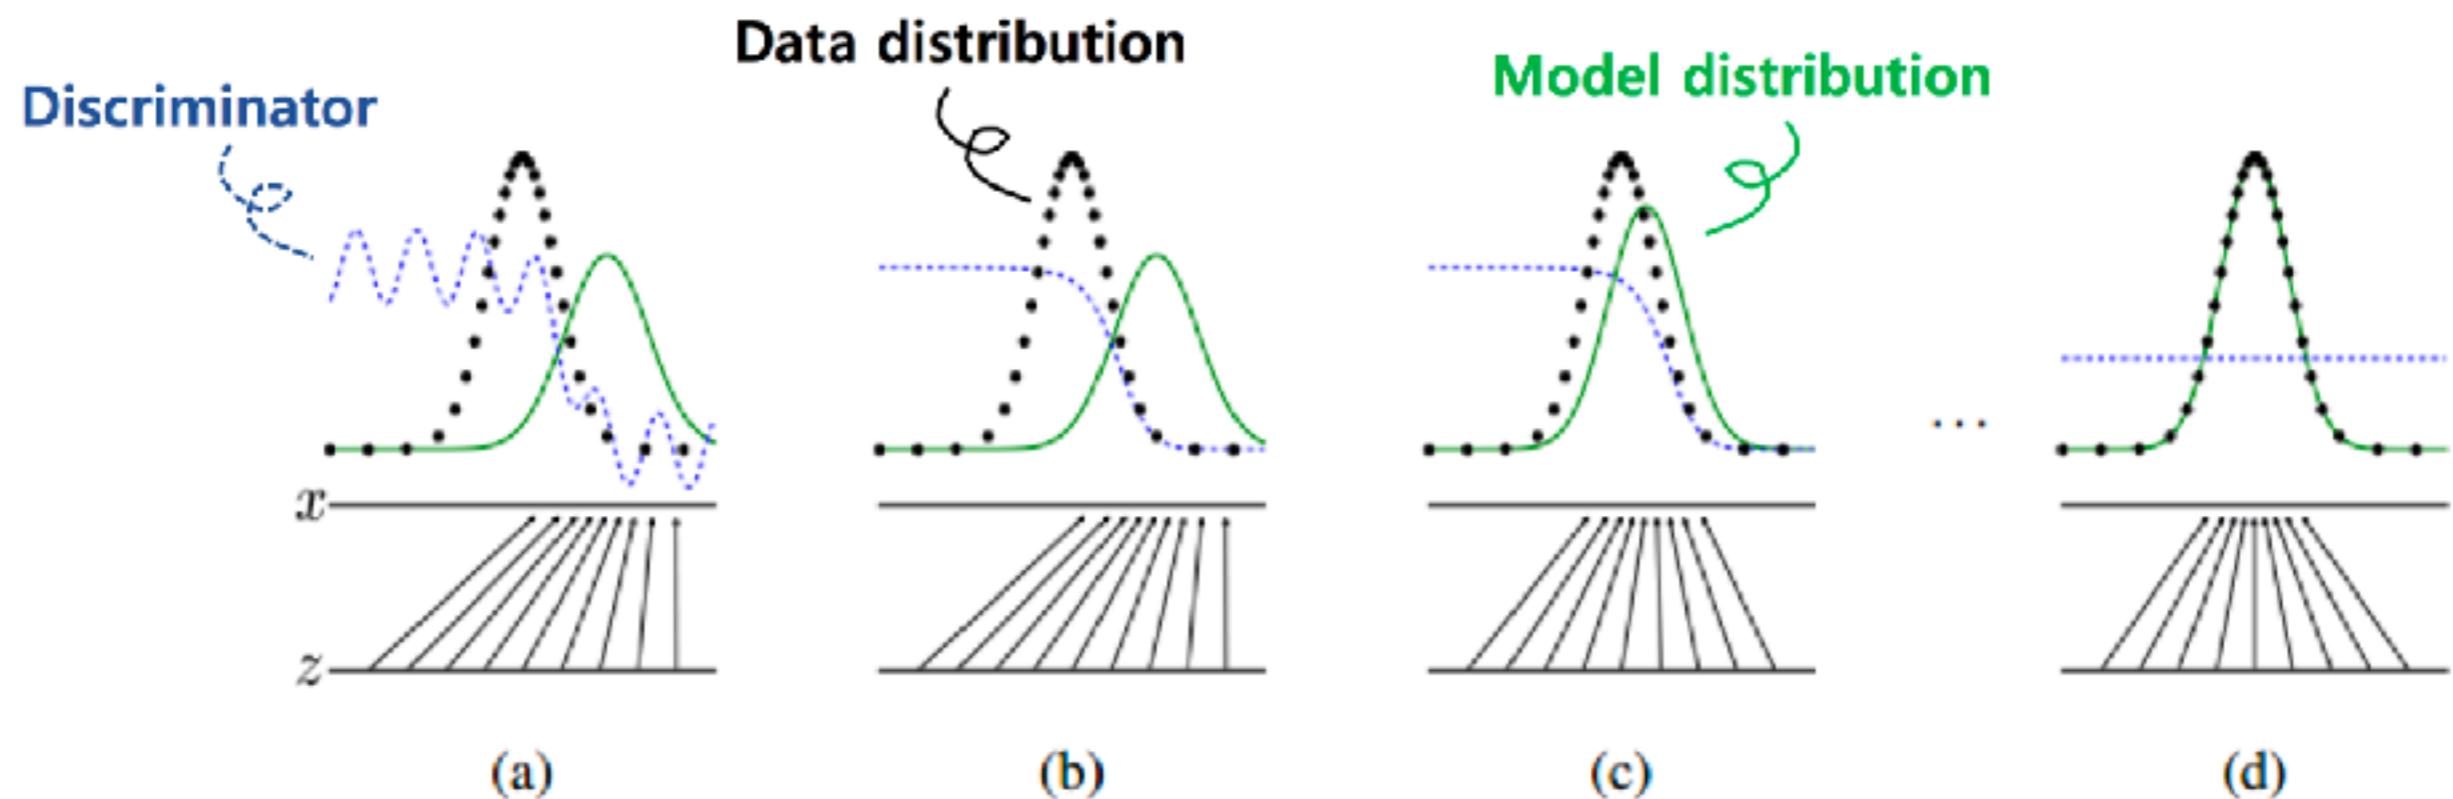
$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))].$$

<https://arxiv.org/abs/1406.2661>



Generative Model

Generative Adversarial Network (GAN)

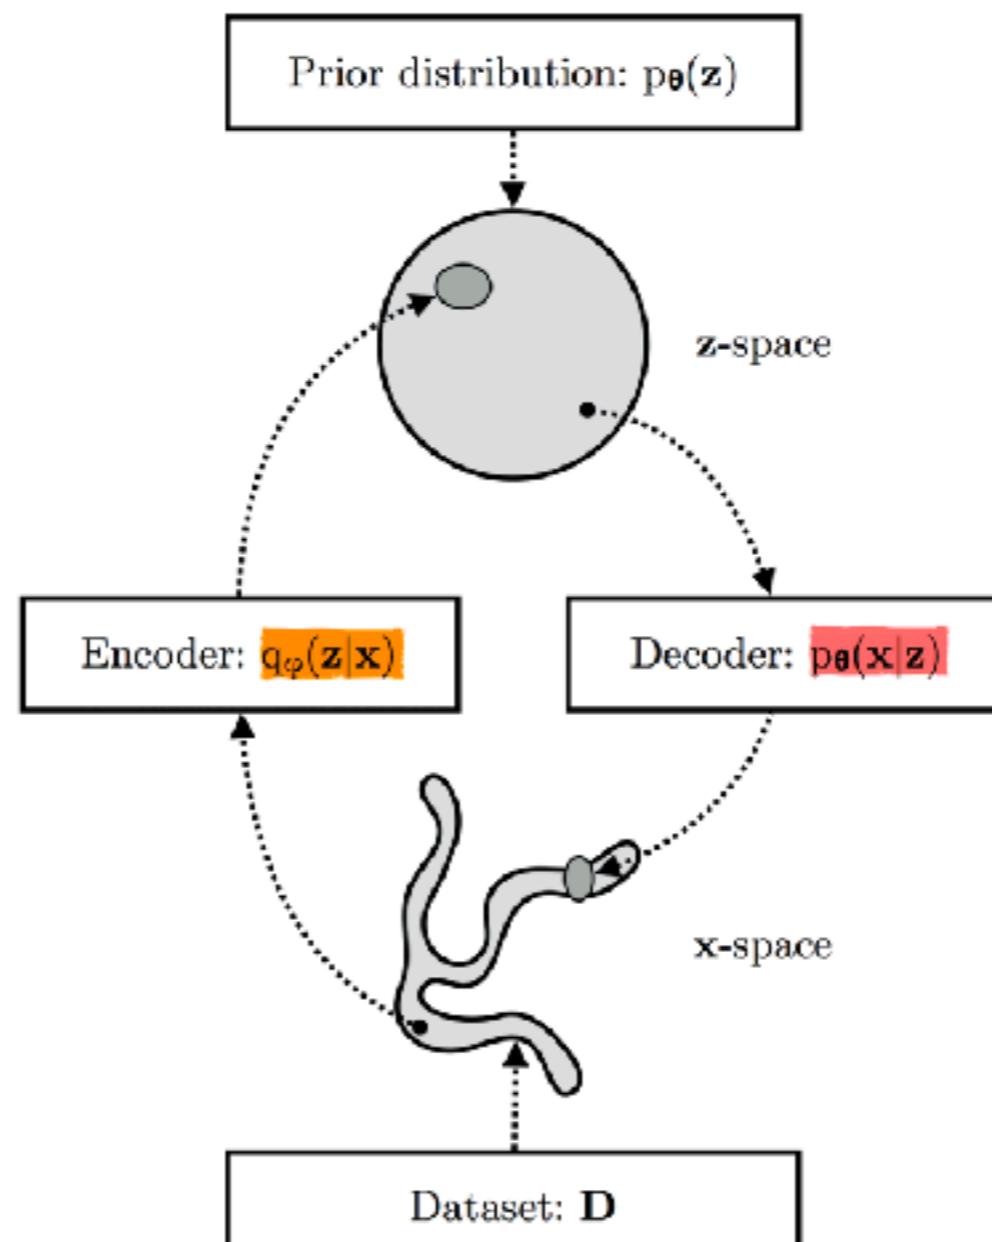


<https://arxiv.org/abs/1406.2661>



Generative Model

Variational Autoencoder (VAE)



Kingma's Thesis



Generative Model

Variational Autoencoder (VAE)

$$\begin{aligned}\log p_{\theta}(\mathbf{x}) &= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x})] \\&= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right] \right] \\&= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right] \right] \\&= \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \right] \right]}_{=\mathcal{L}_{\theta, \phi}(\mathbf{x}) \text{ (ELBO)}} + \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right] \right]}_{=D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z}|\mathbf{x}))}\end{aligned}$$

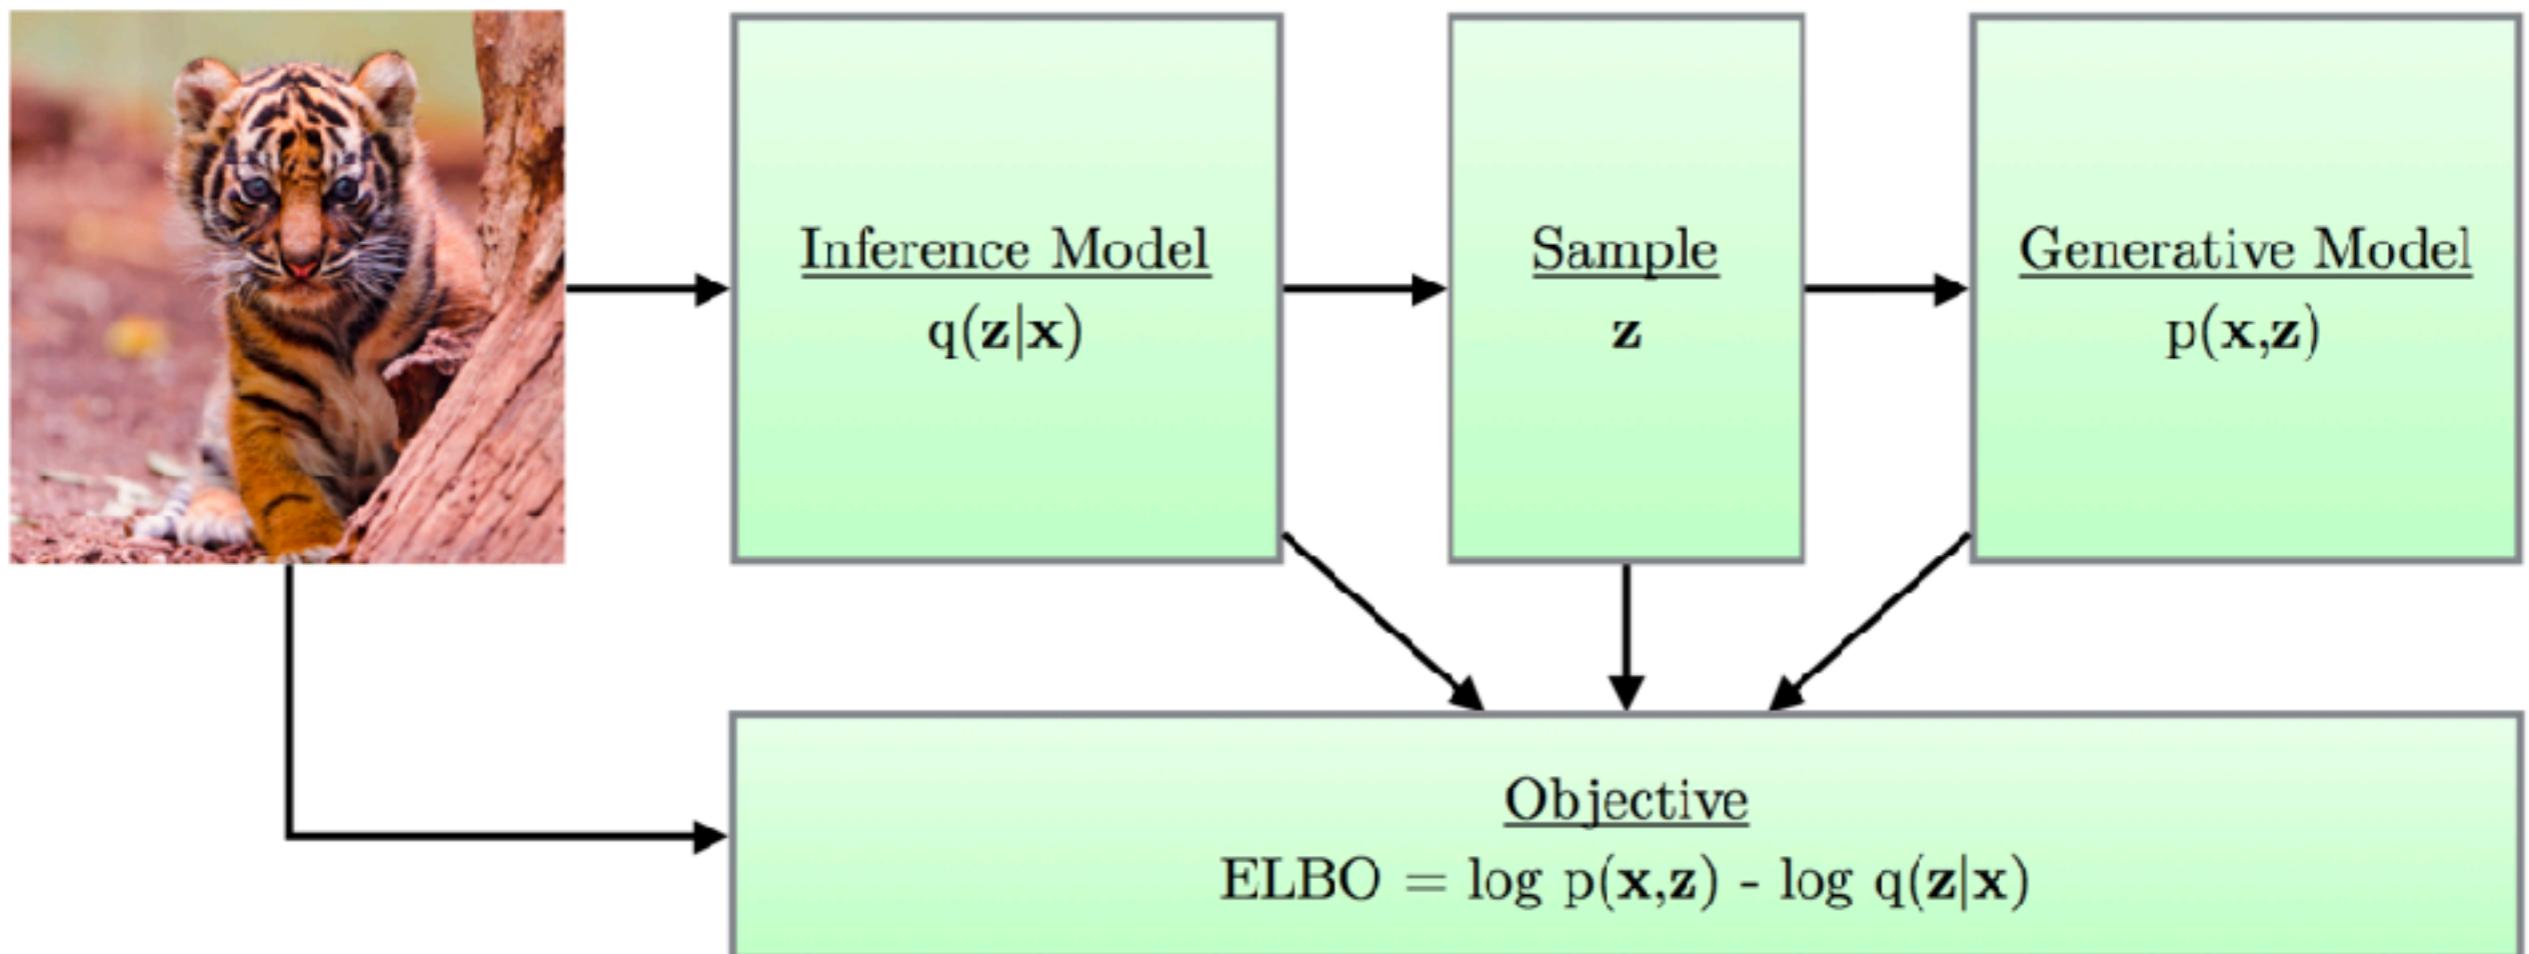
Kingma's Thesis



Generative Model

Variational Autoencoder (VAE)

Datapoint

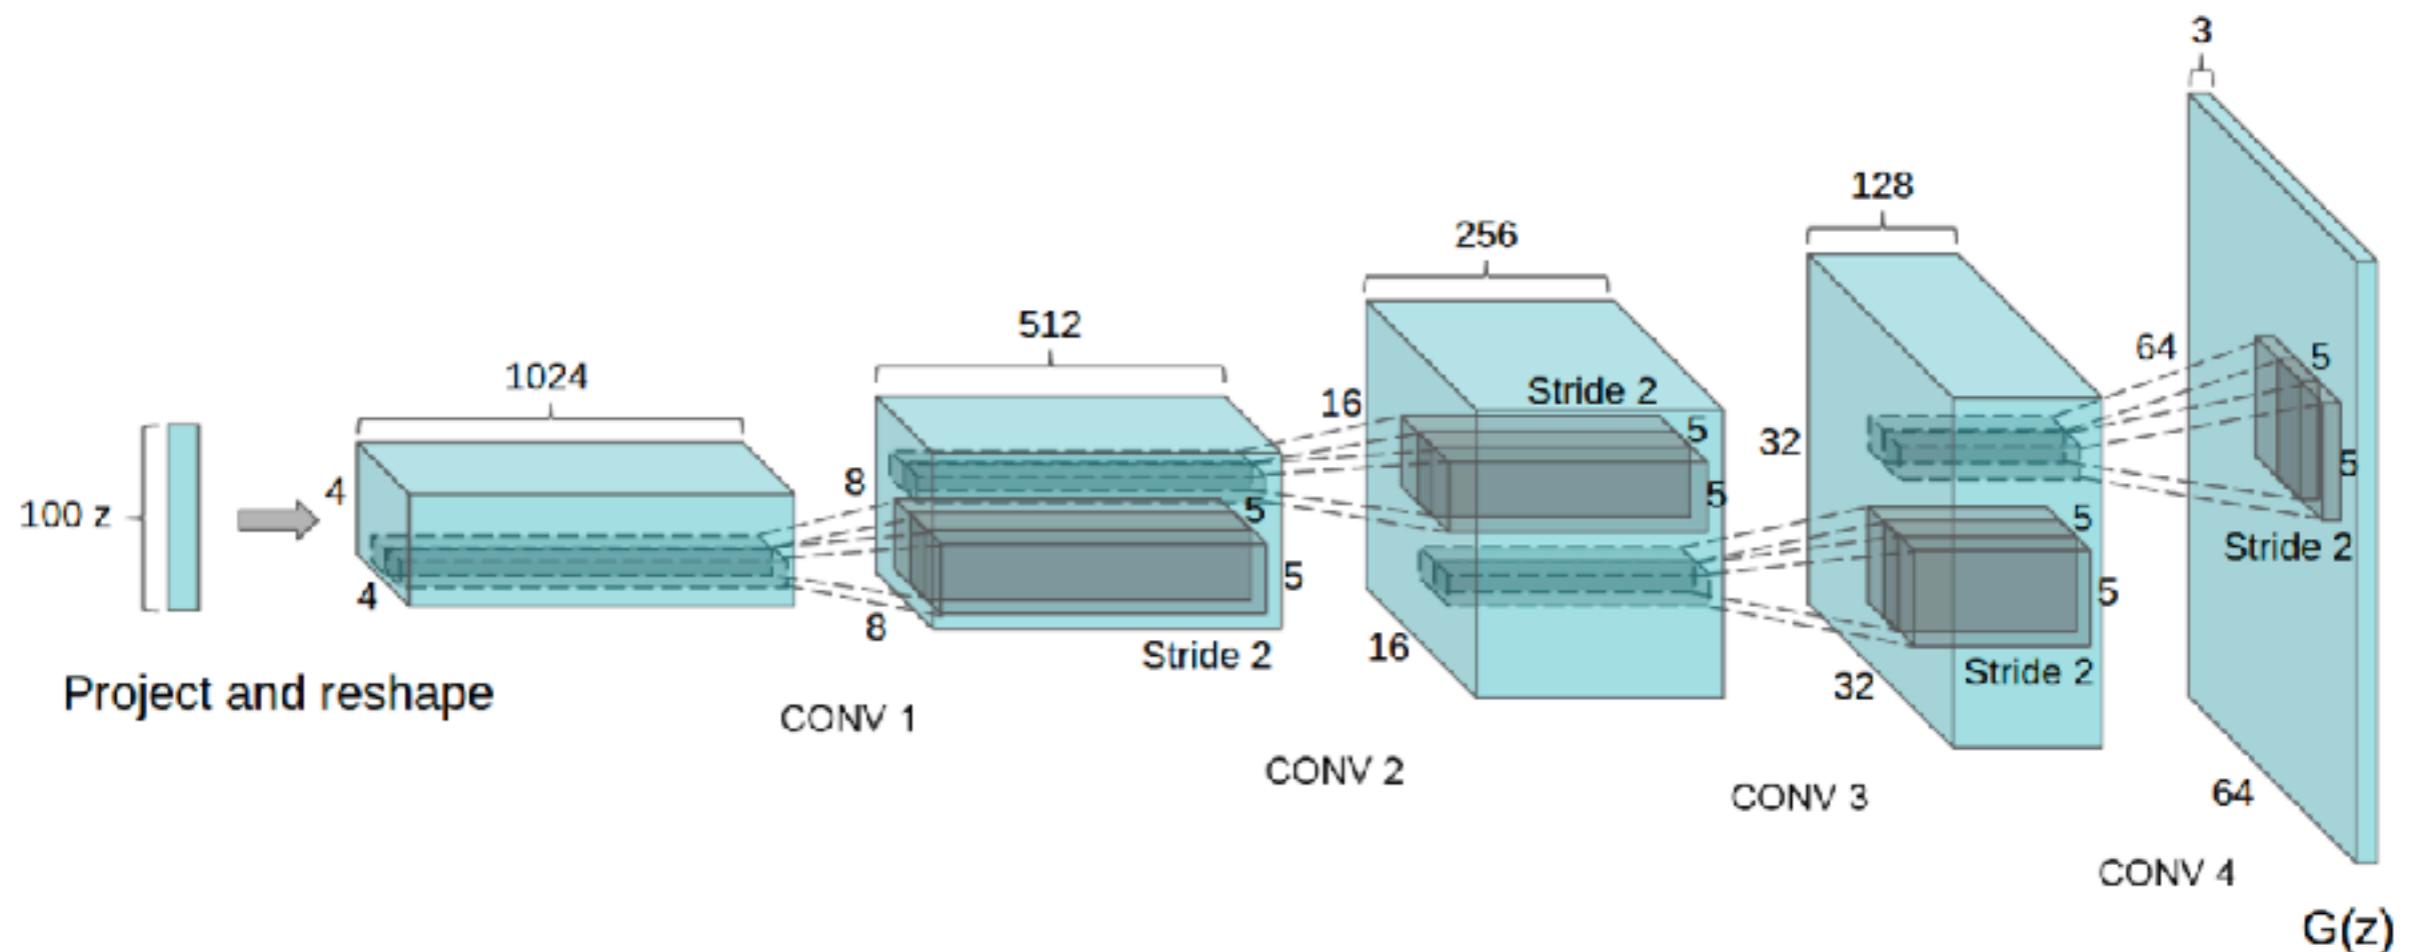


Kingma's Thesis



Generative Model

DCGAN



<https://arxiv.org/abs/1511.06434>



Generative Model

Info GAN



(a) Varying c_1 on InfoGAN (Digit type)



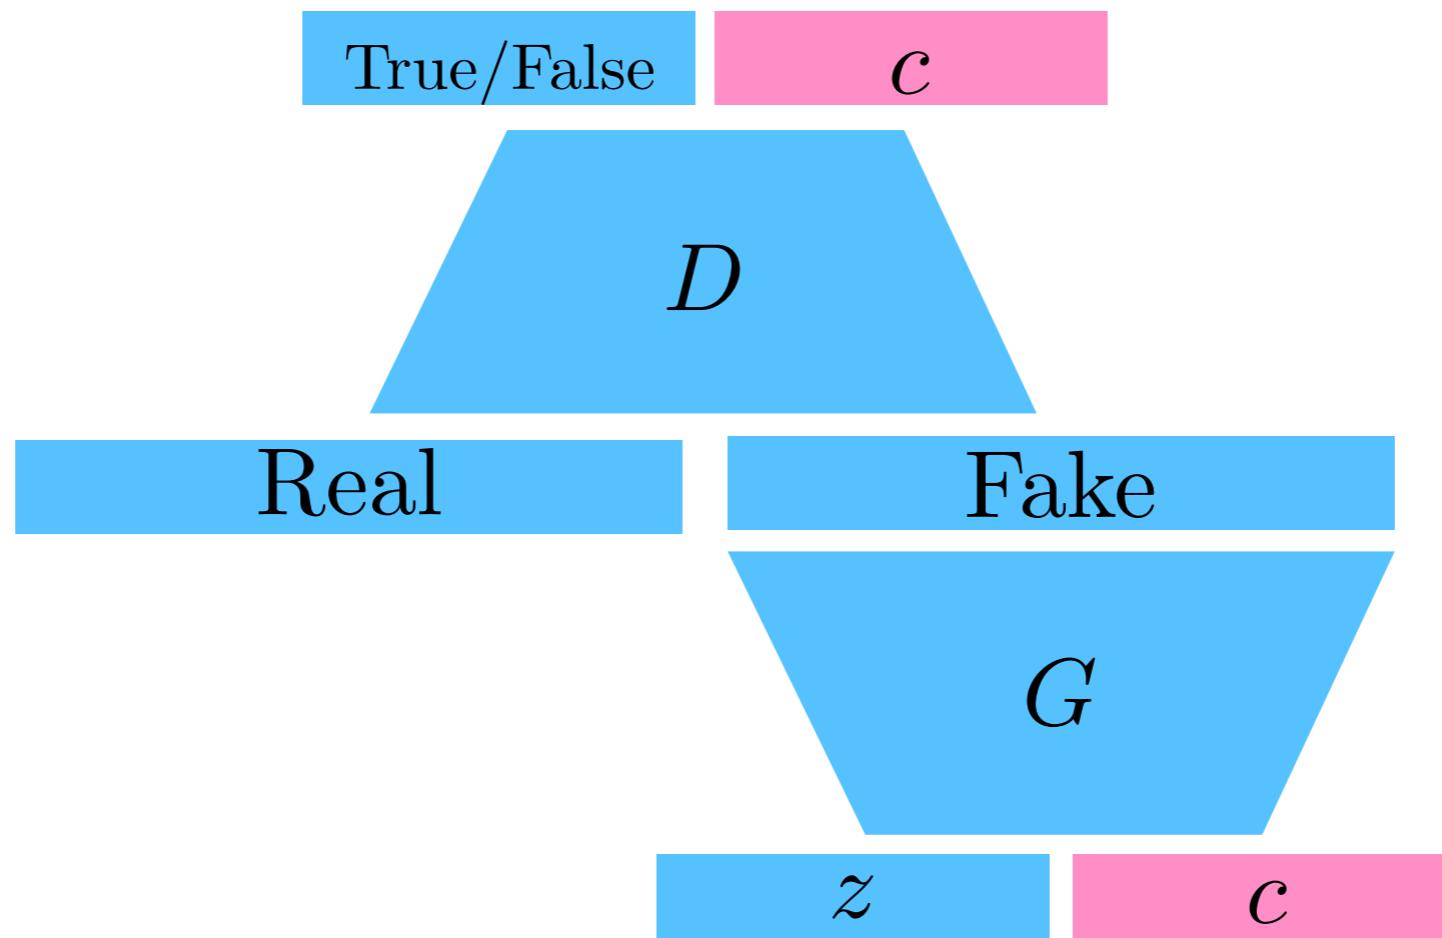
(c) Varying c_2 from -2 to 2 on InfoGAN (Rotation)

<https://arxiv.org/abs/1606.03657>



Generative Model

Info GAN

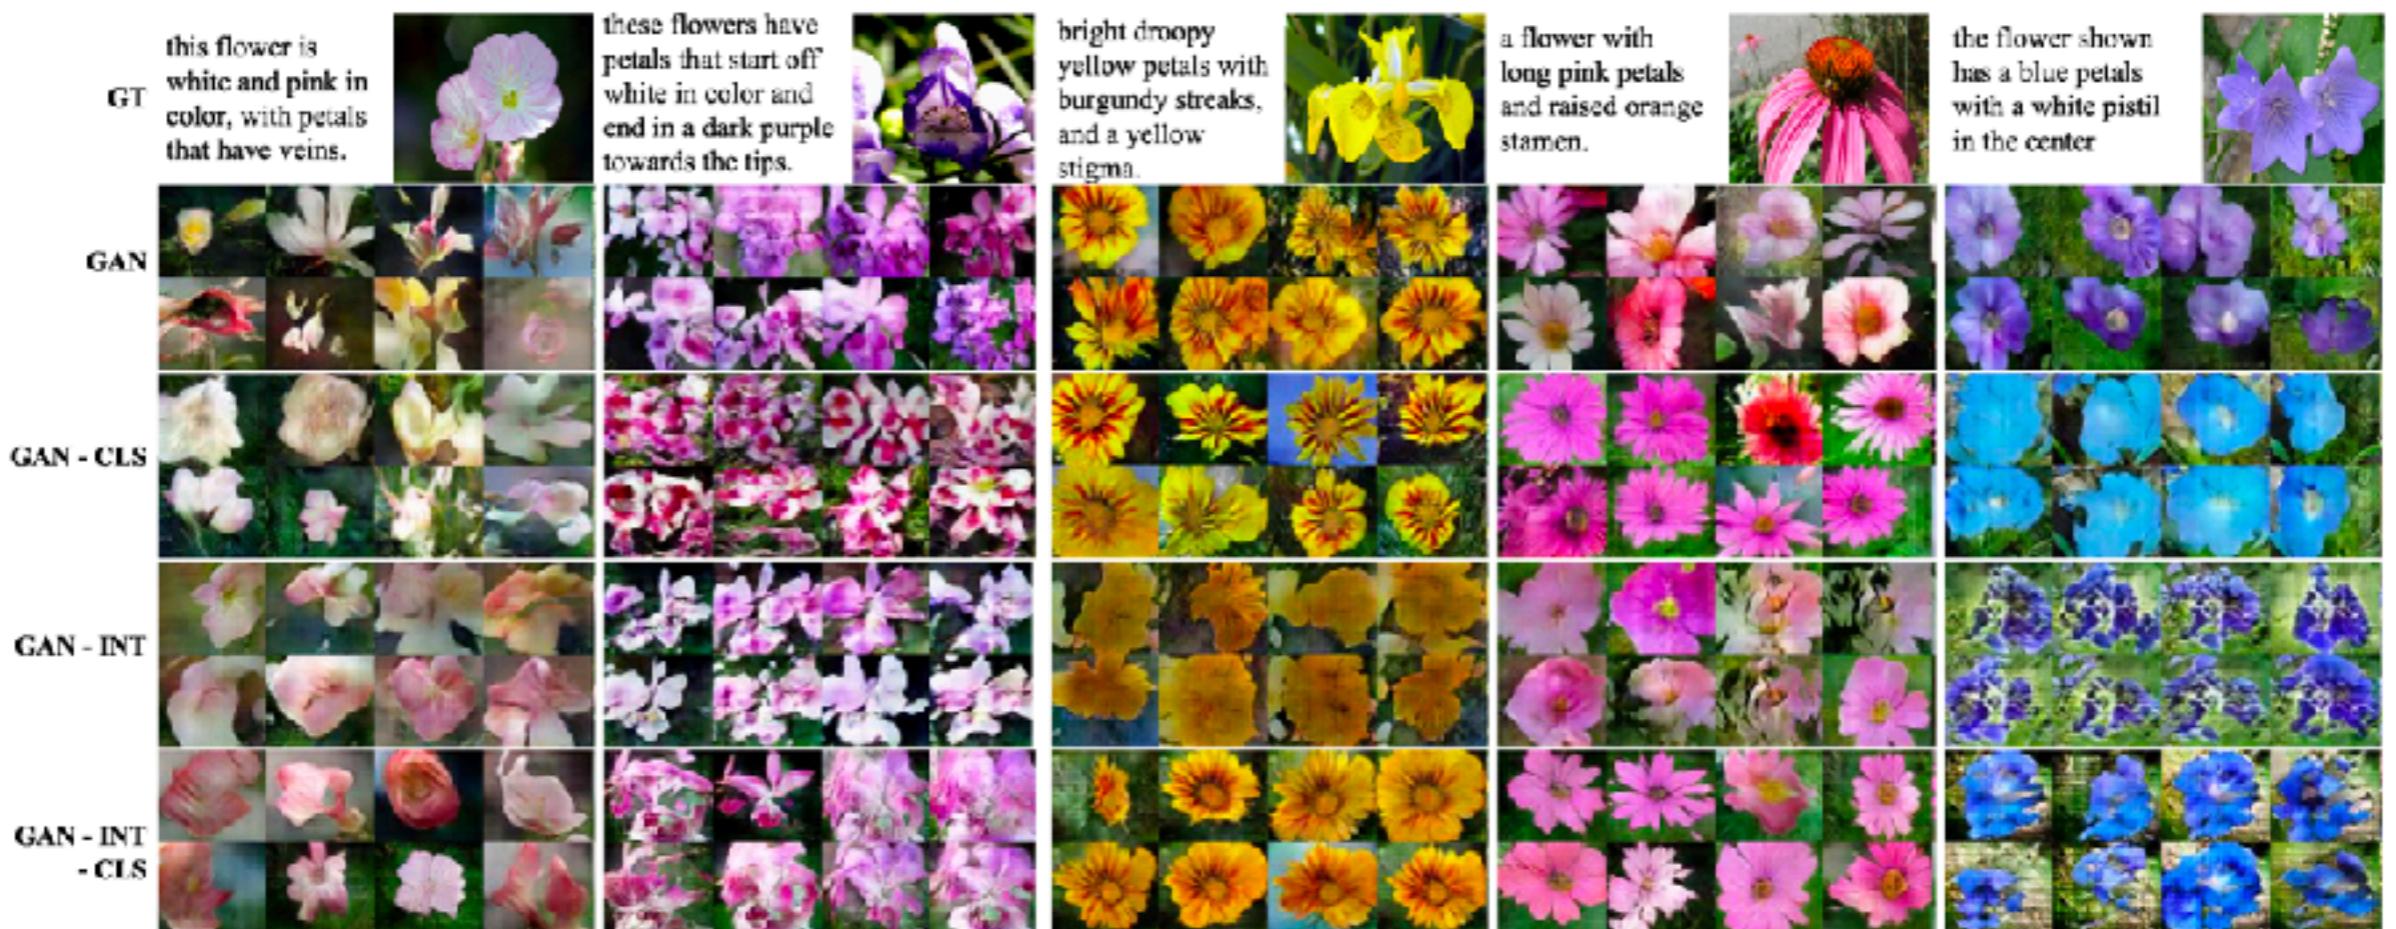


<https://arxiv.org/abs/1606.03657>



Generative Model

Text2Image

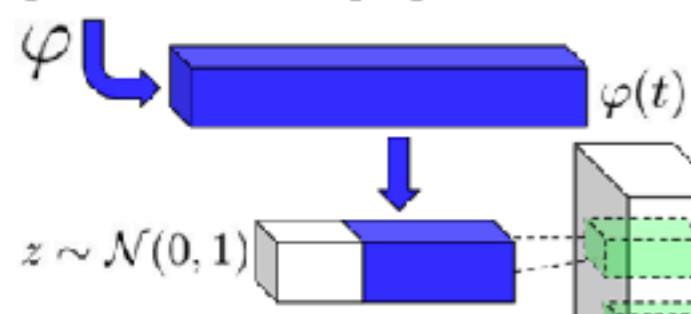


<https://arxiv.org/abs/1605.05396>

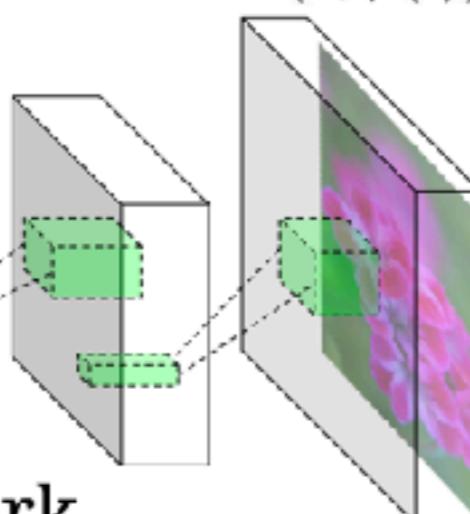
Generative Model

Text2Image

This flower has small, round violet petals with a dark purple center

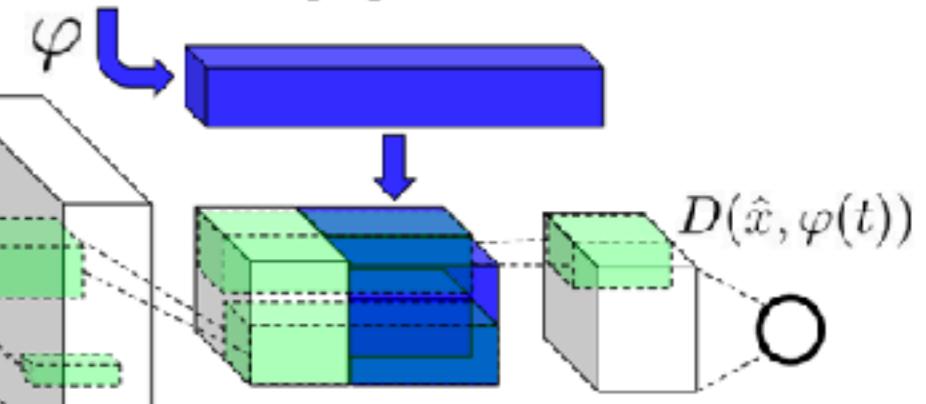


$$\hat{x} := G(z, \varphi(t))$$



Generator Network

This flower has small, round violet petals with a dark purple center

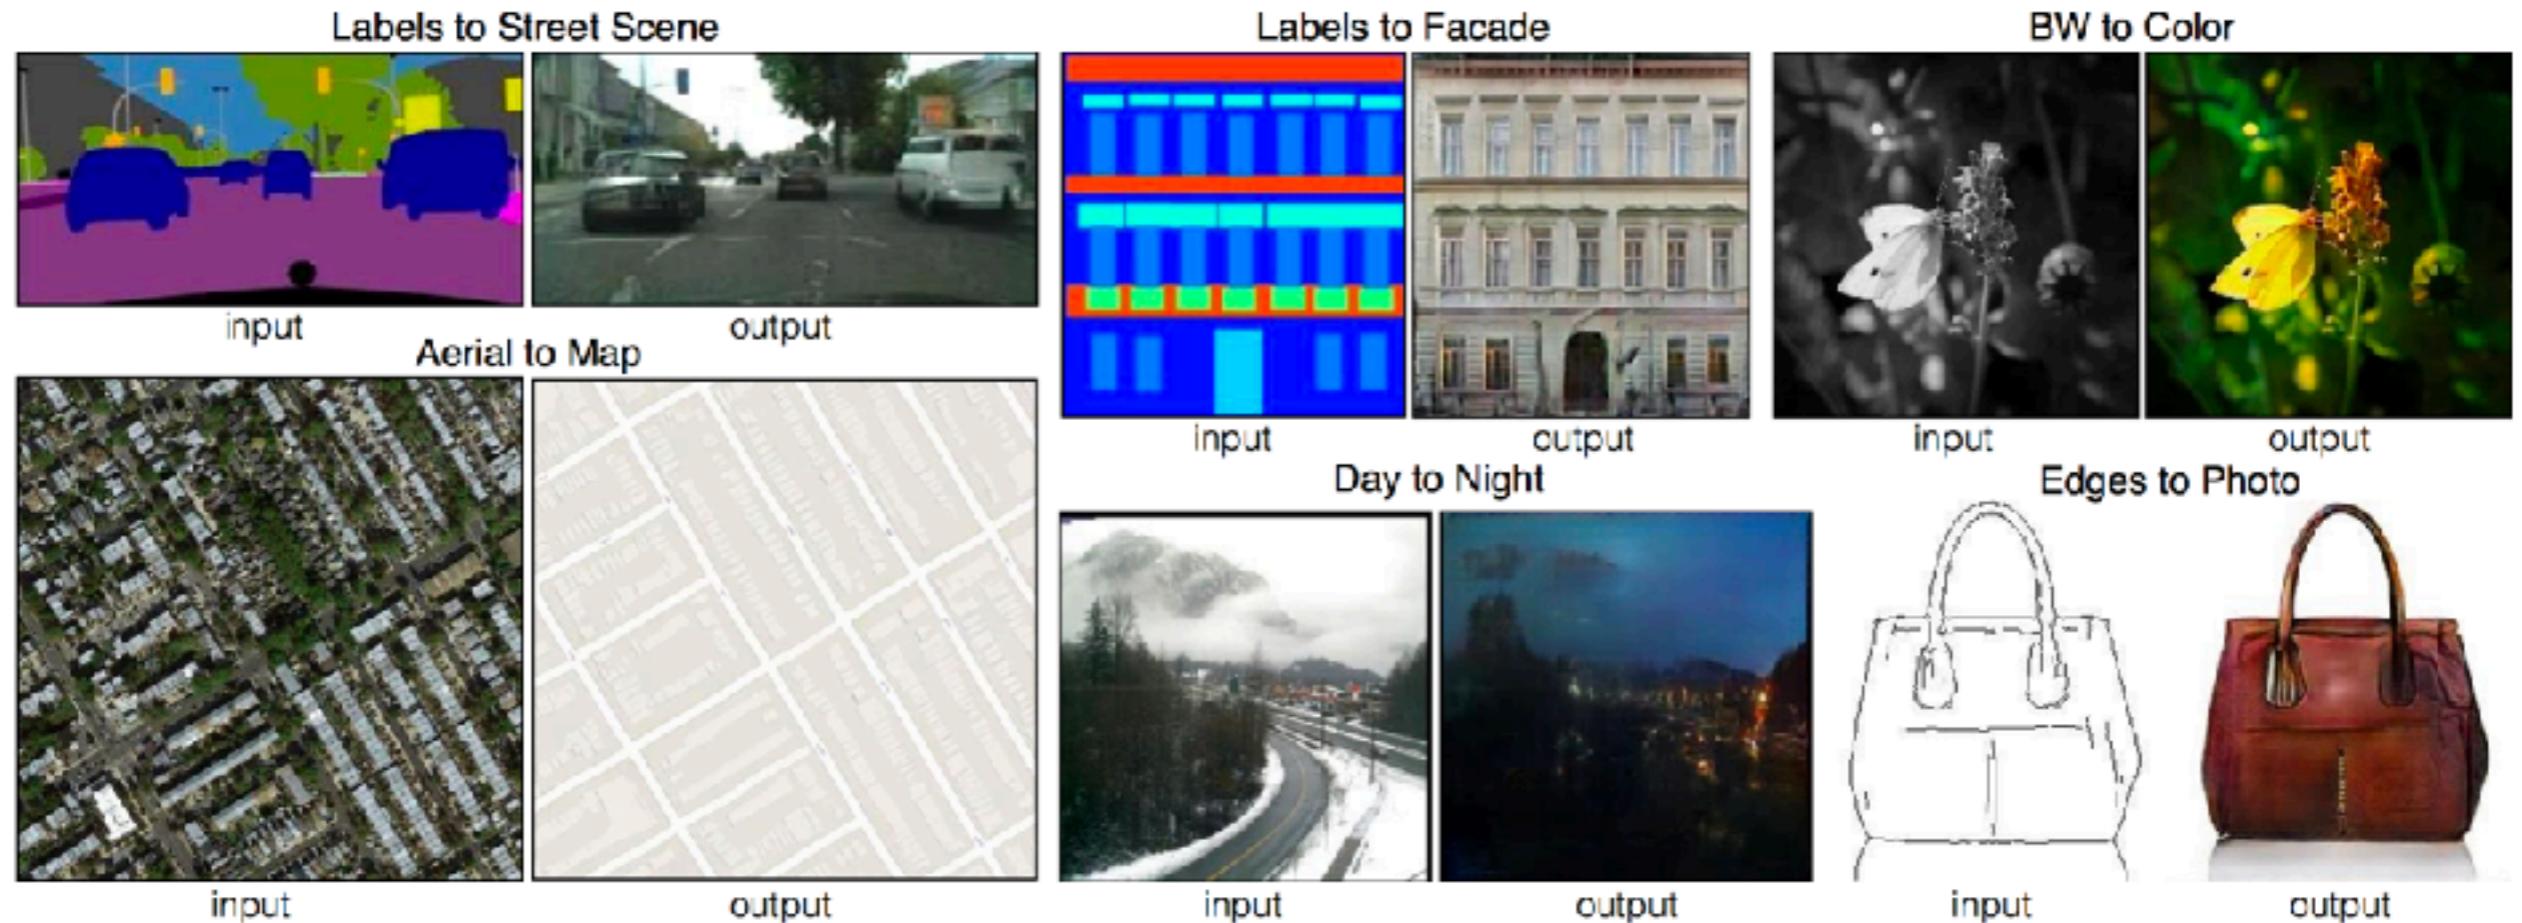


Discriminator Network

<https://arxiv.org/abs/1605.05396>

Generative Model

Pix2pix

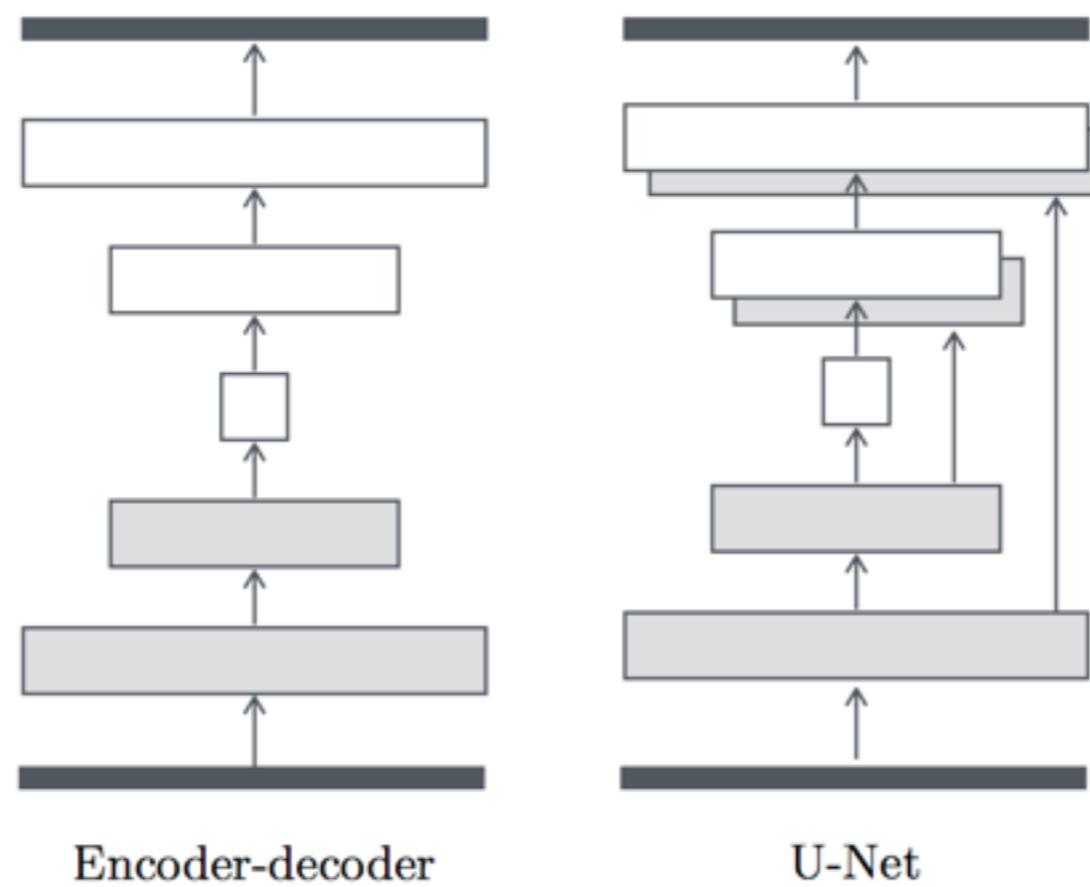


<https://arxiv.org/abs/1611.07004>



Generative Model

Pix2pix

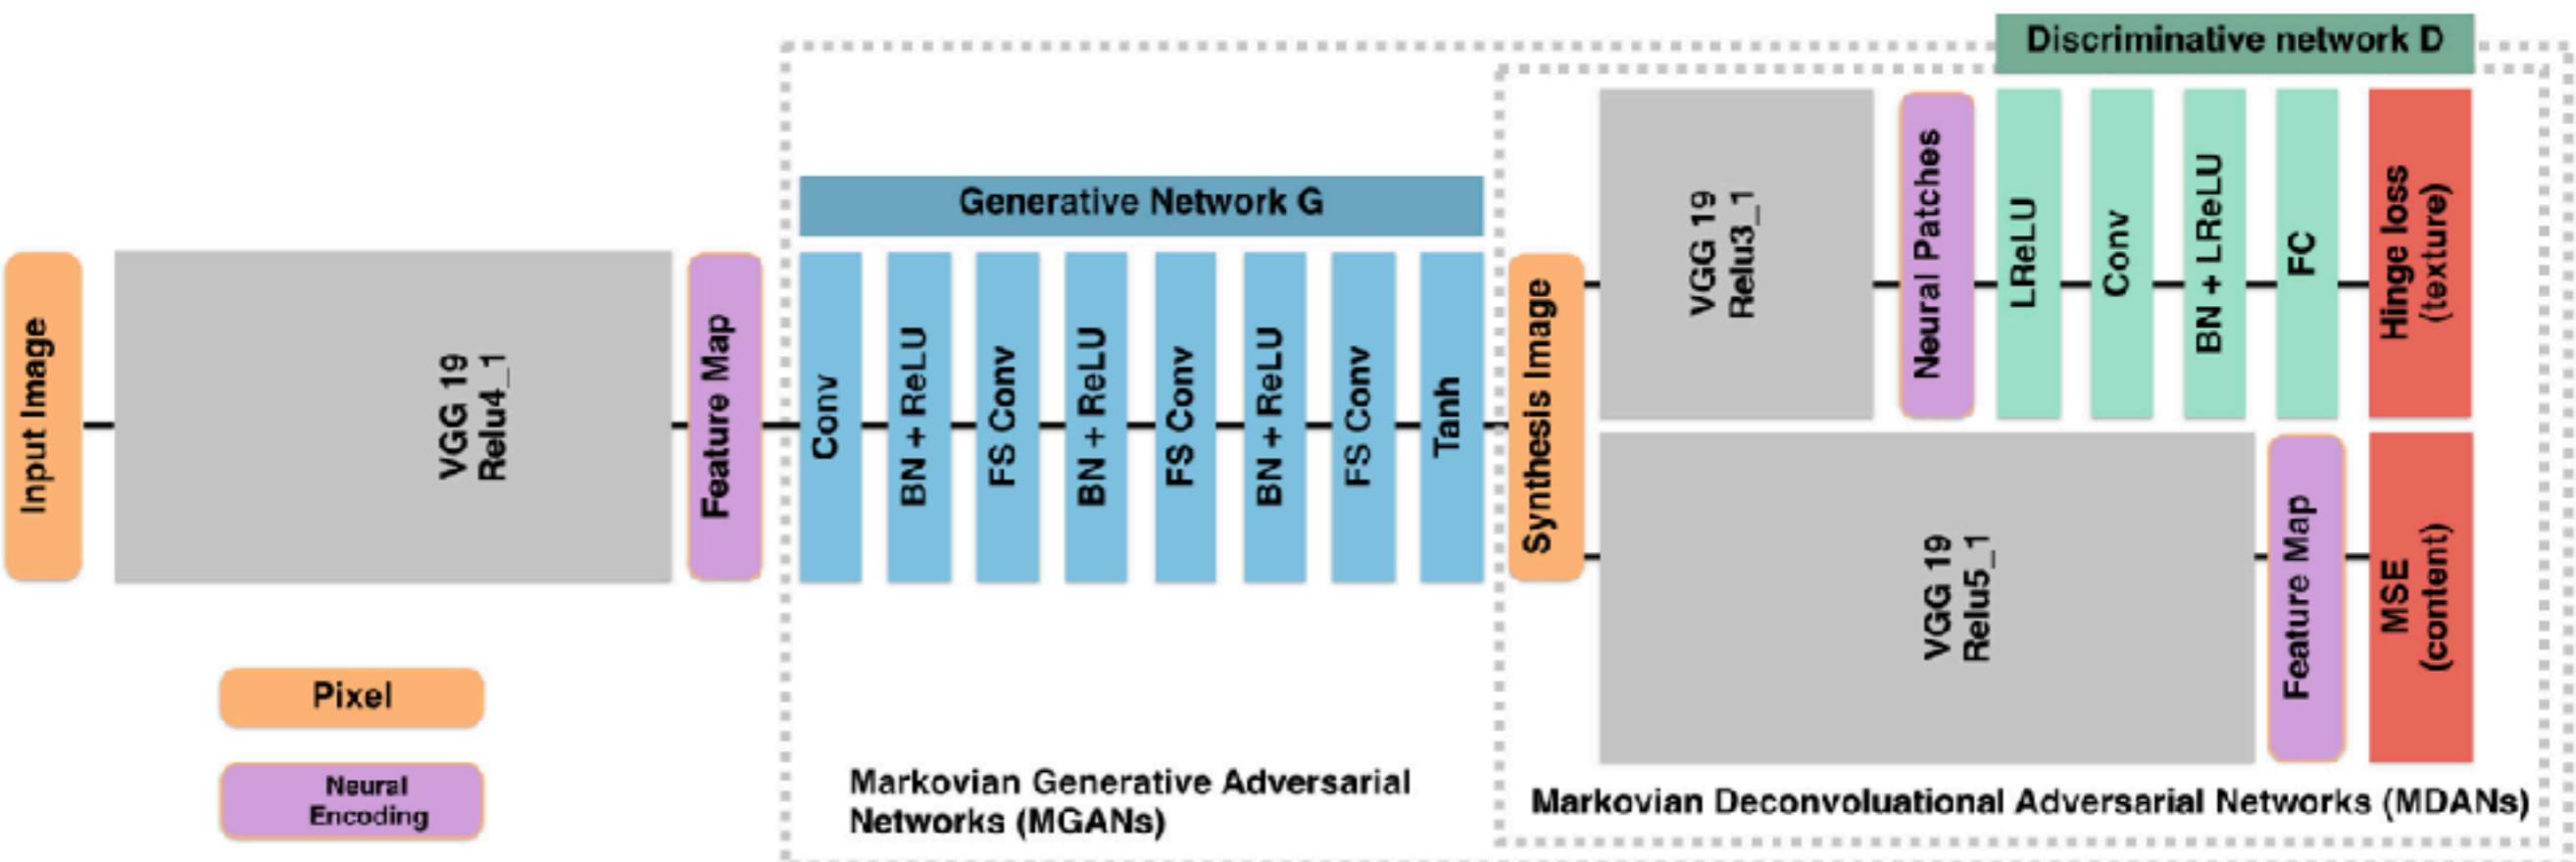


<https://arxiv.org/abs/1611.07004>



Generative Model

PatchGAN

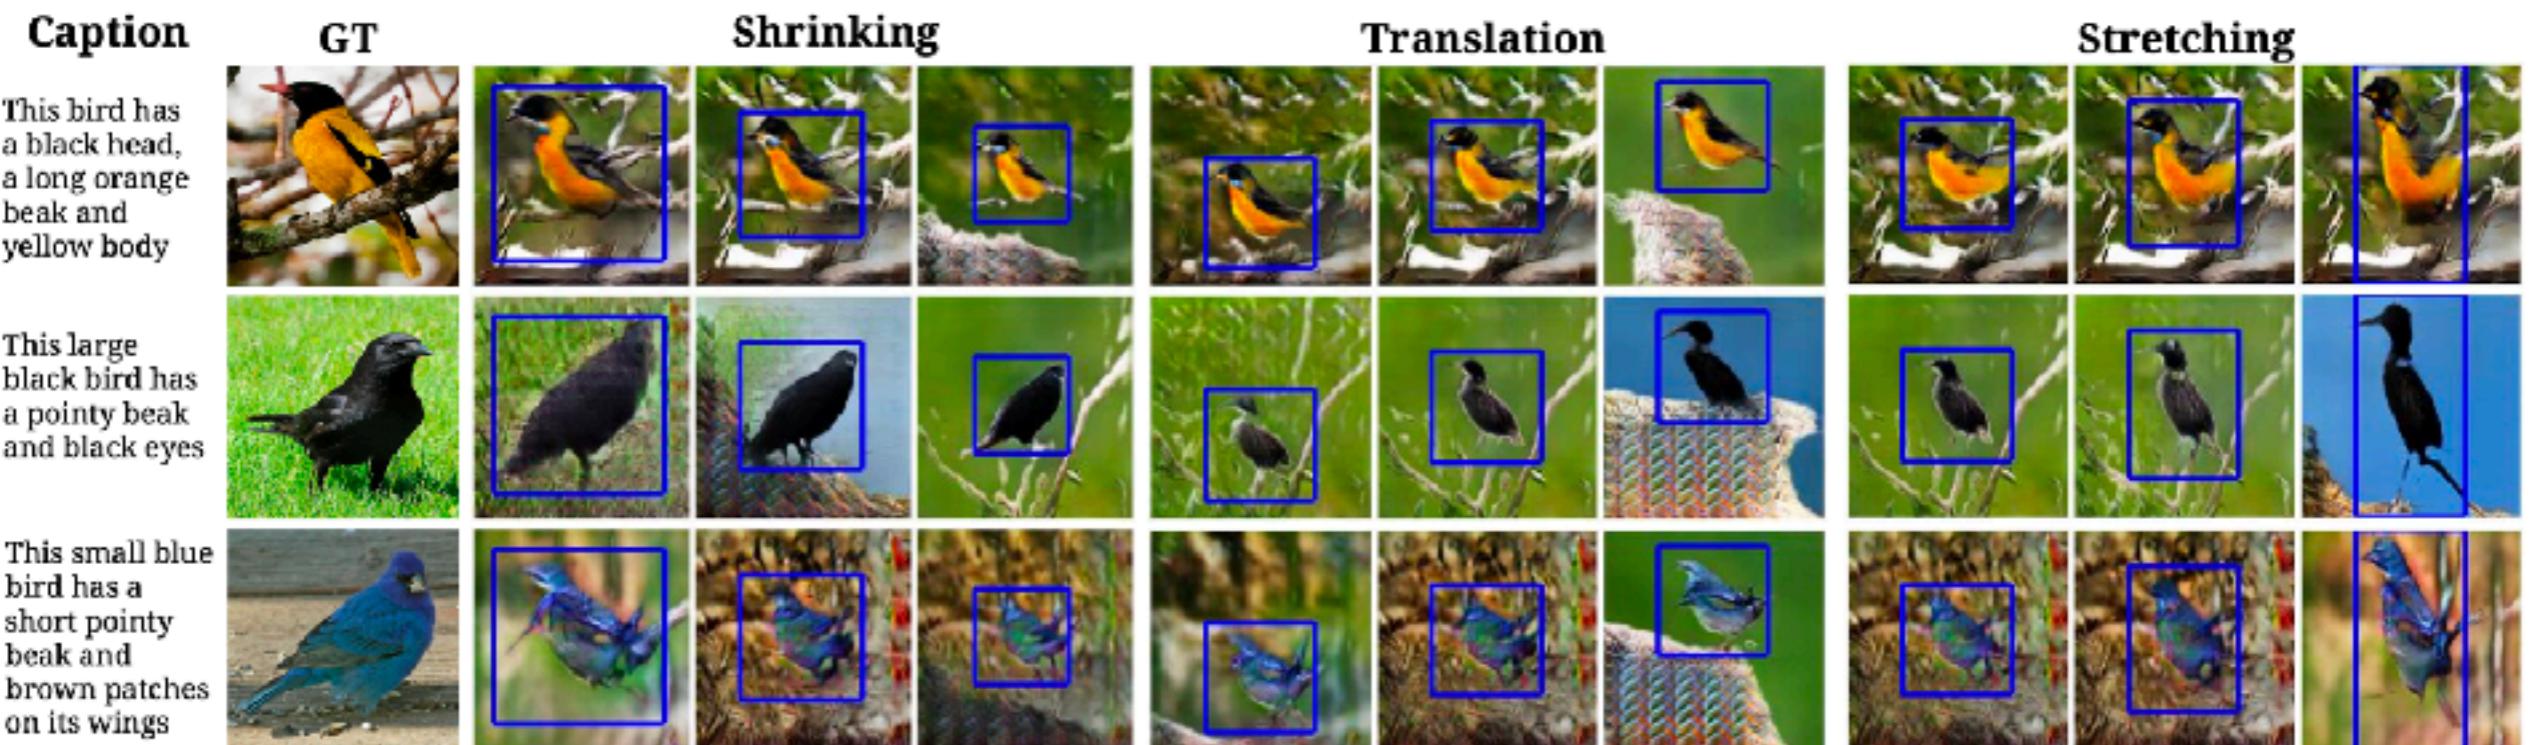


<https://arxiv.org/abs/1604.04382>



Generative Model

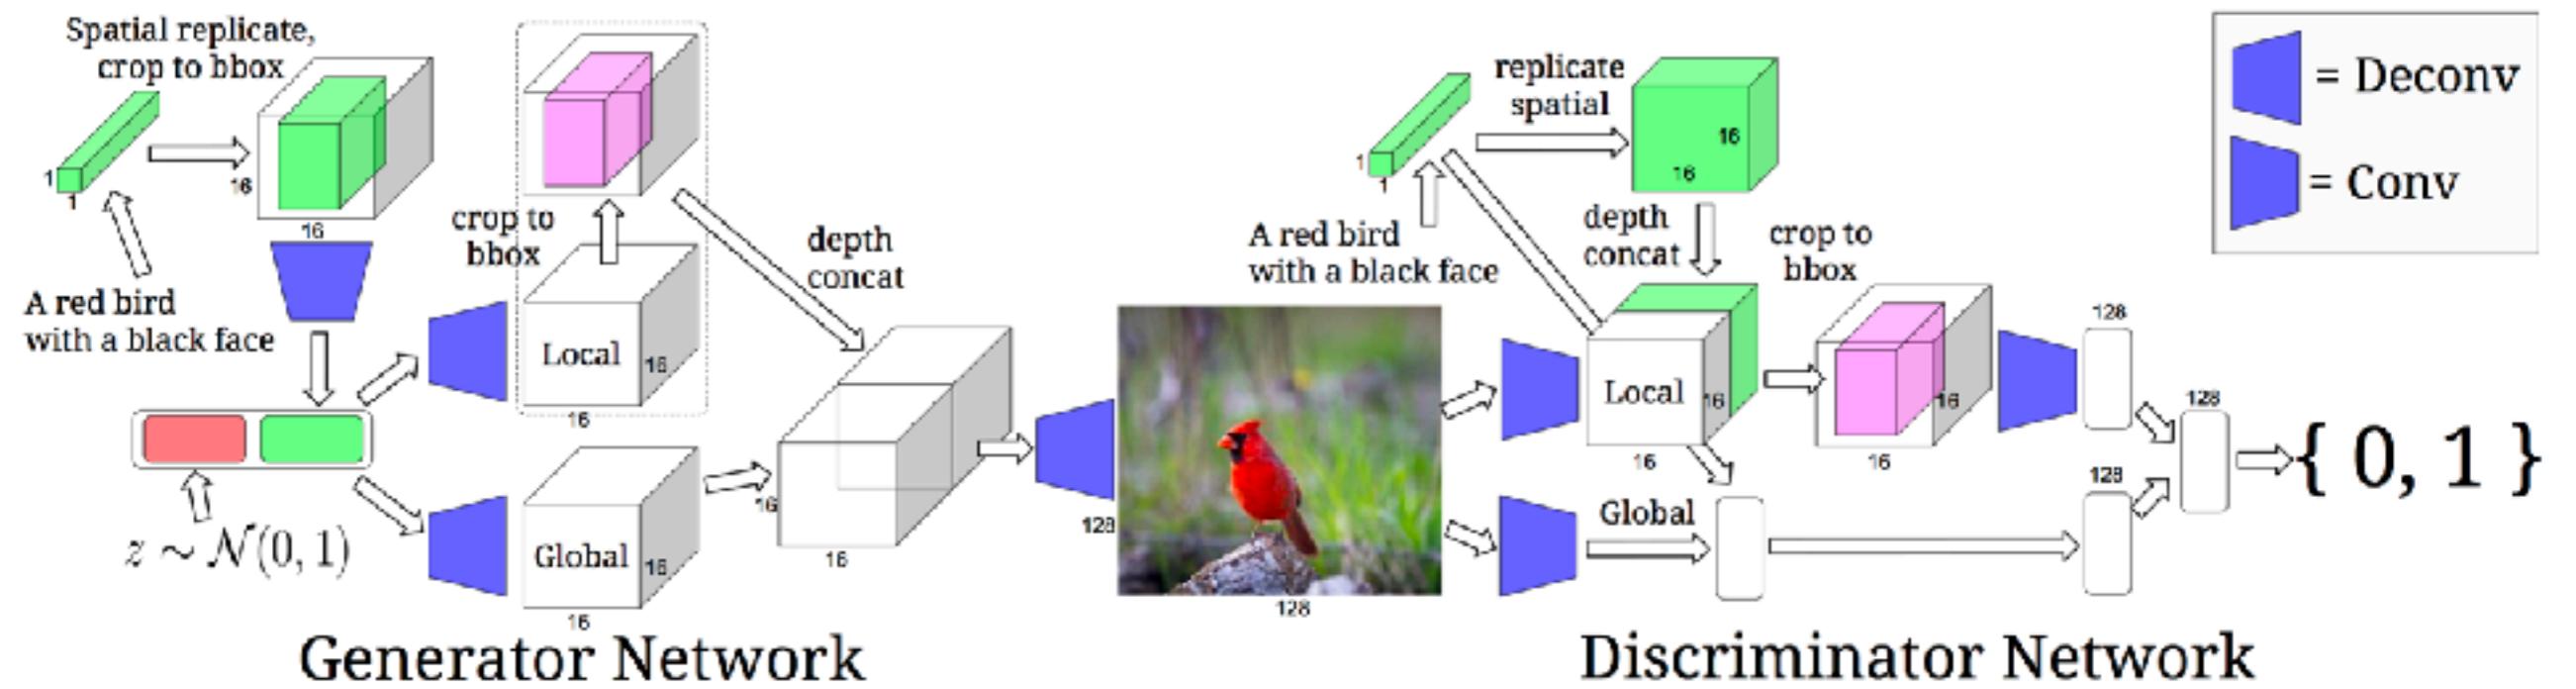
Generative Adversarial What-Where Networks (GAWWN)



<https://arxiv.org/abs/1610.02454>

Generative Model

Generative Adversarial What-Where Networks (GAWWN)



<https://arxiv.org/abs/1610.02454>



Generative Model

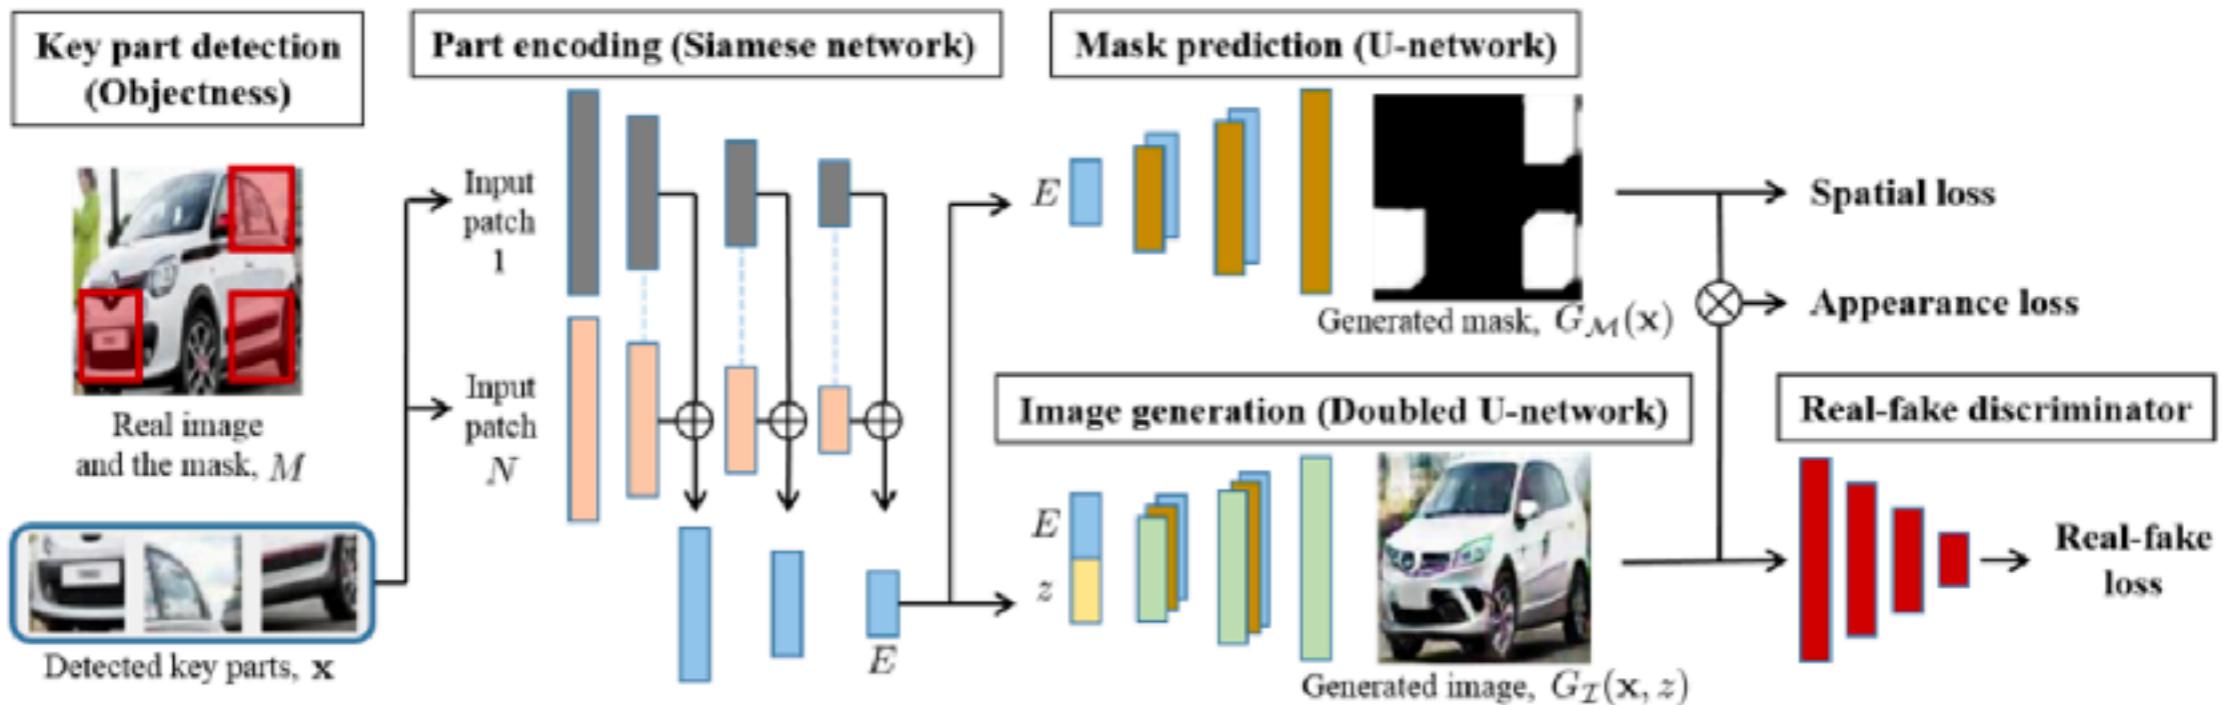
Puzzle-GAN



<https://arxiv.org/abs/1703.10730>

Generative Model

Puzzle-GAN

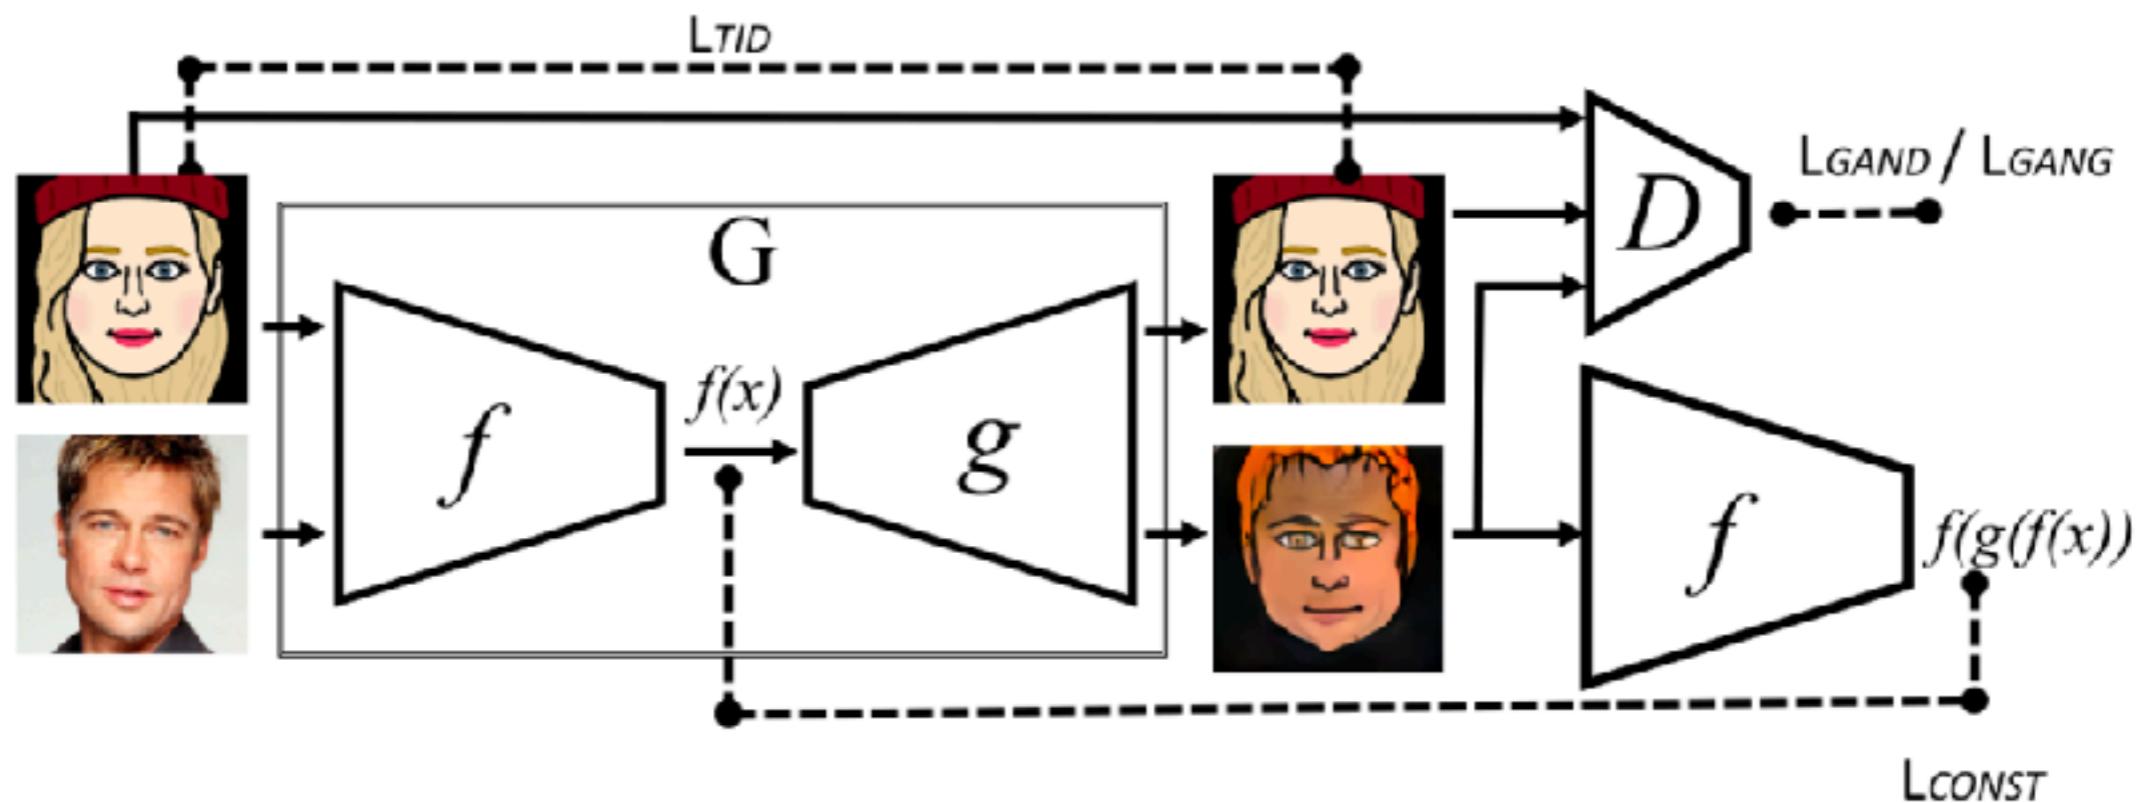


<https://arxiv.org/abs/1703.10730>



Generative Model

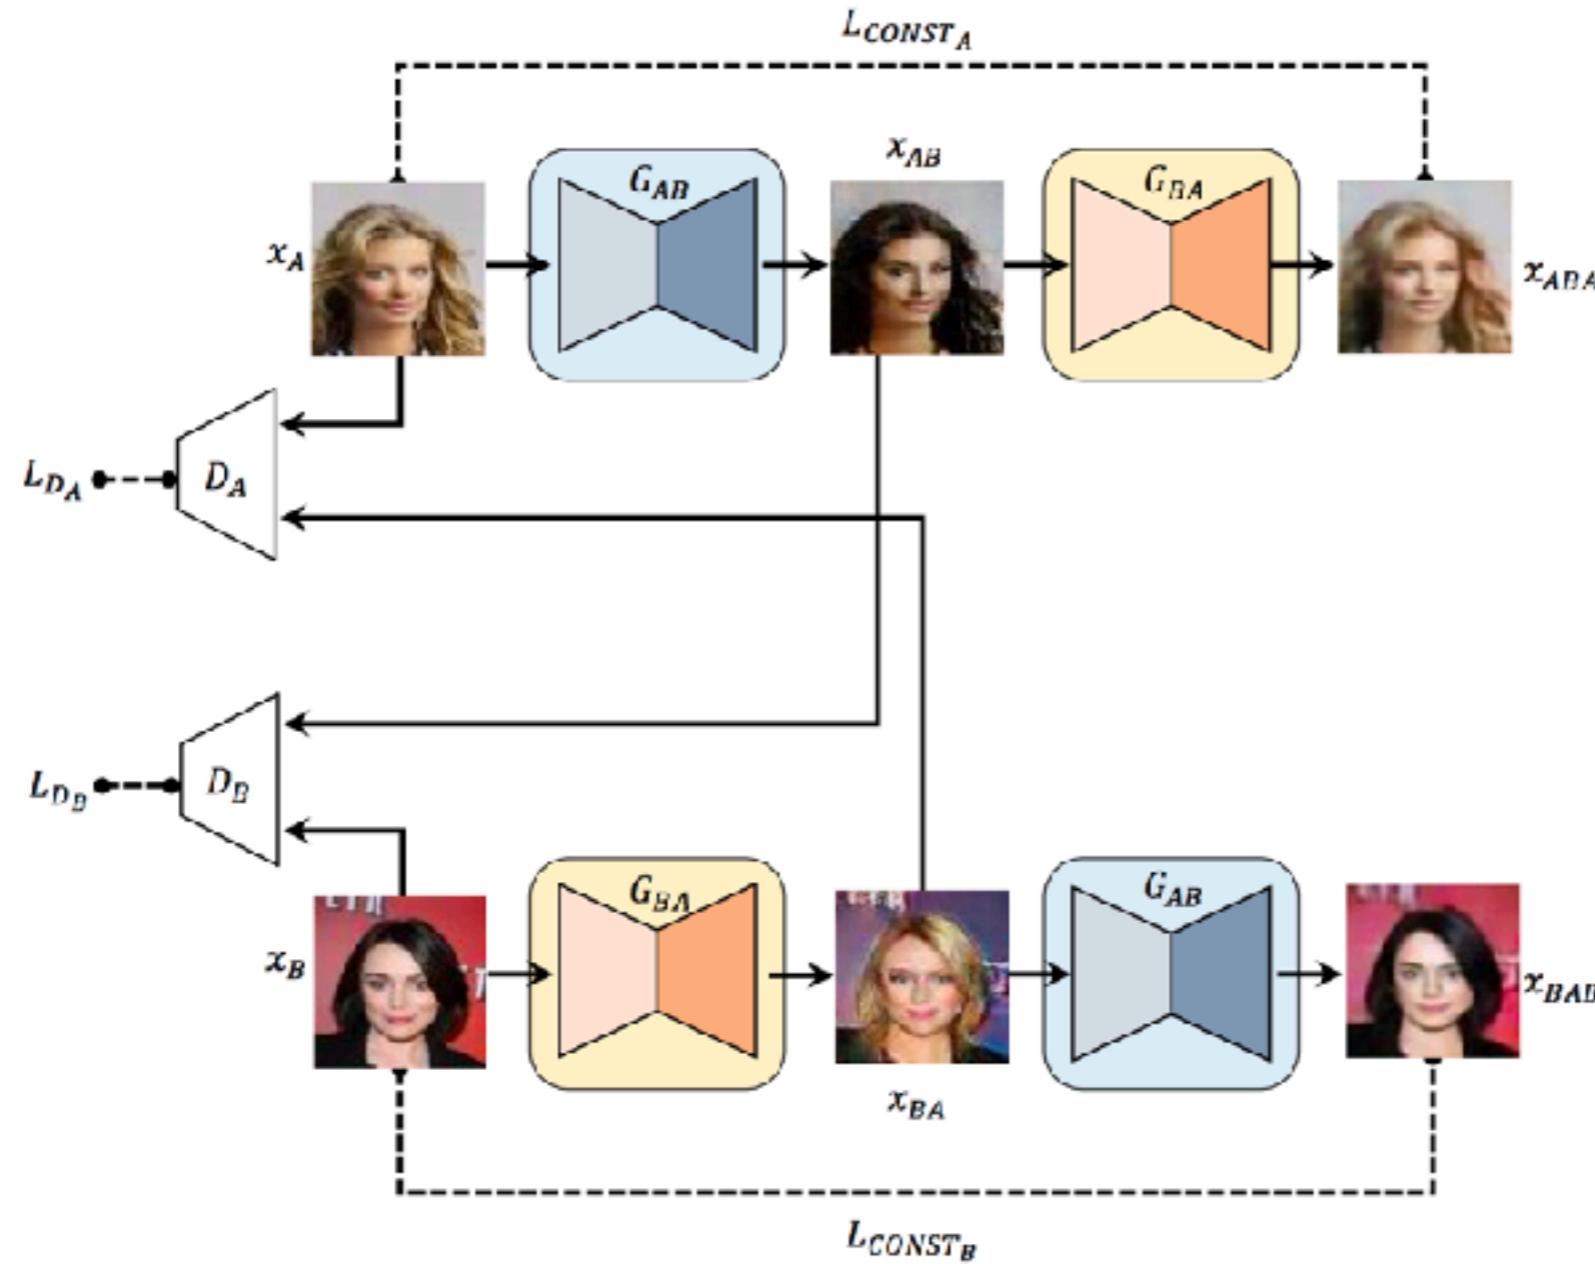
Domain Transfer Network



<https://arxiv.org/abs/1611.02200>

Generative Model

DiscoGAN



<https://arxiv.org/abs/1703.05192>



Generative Model

DiscoGAN

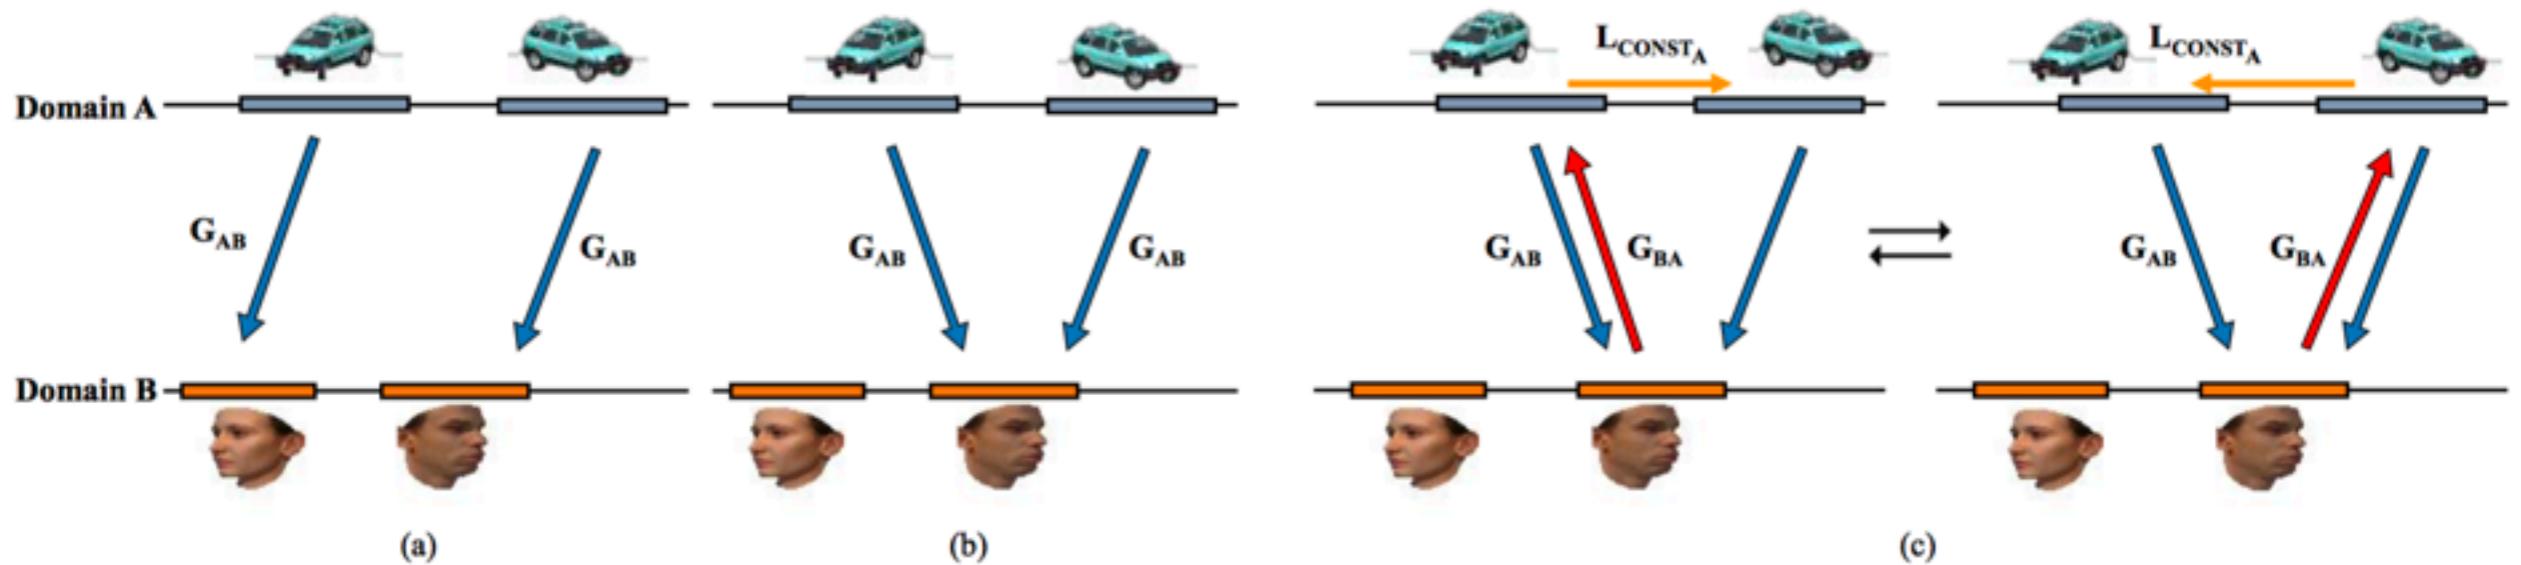


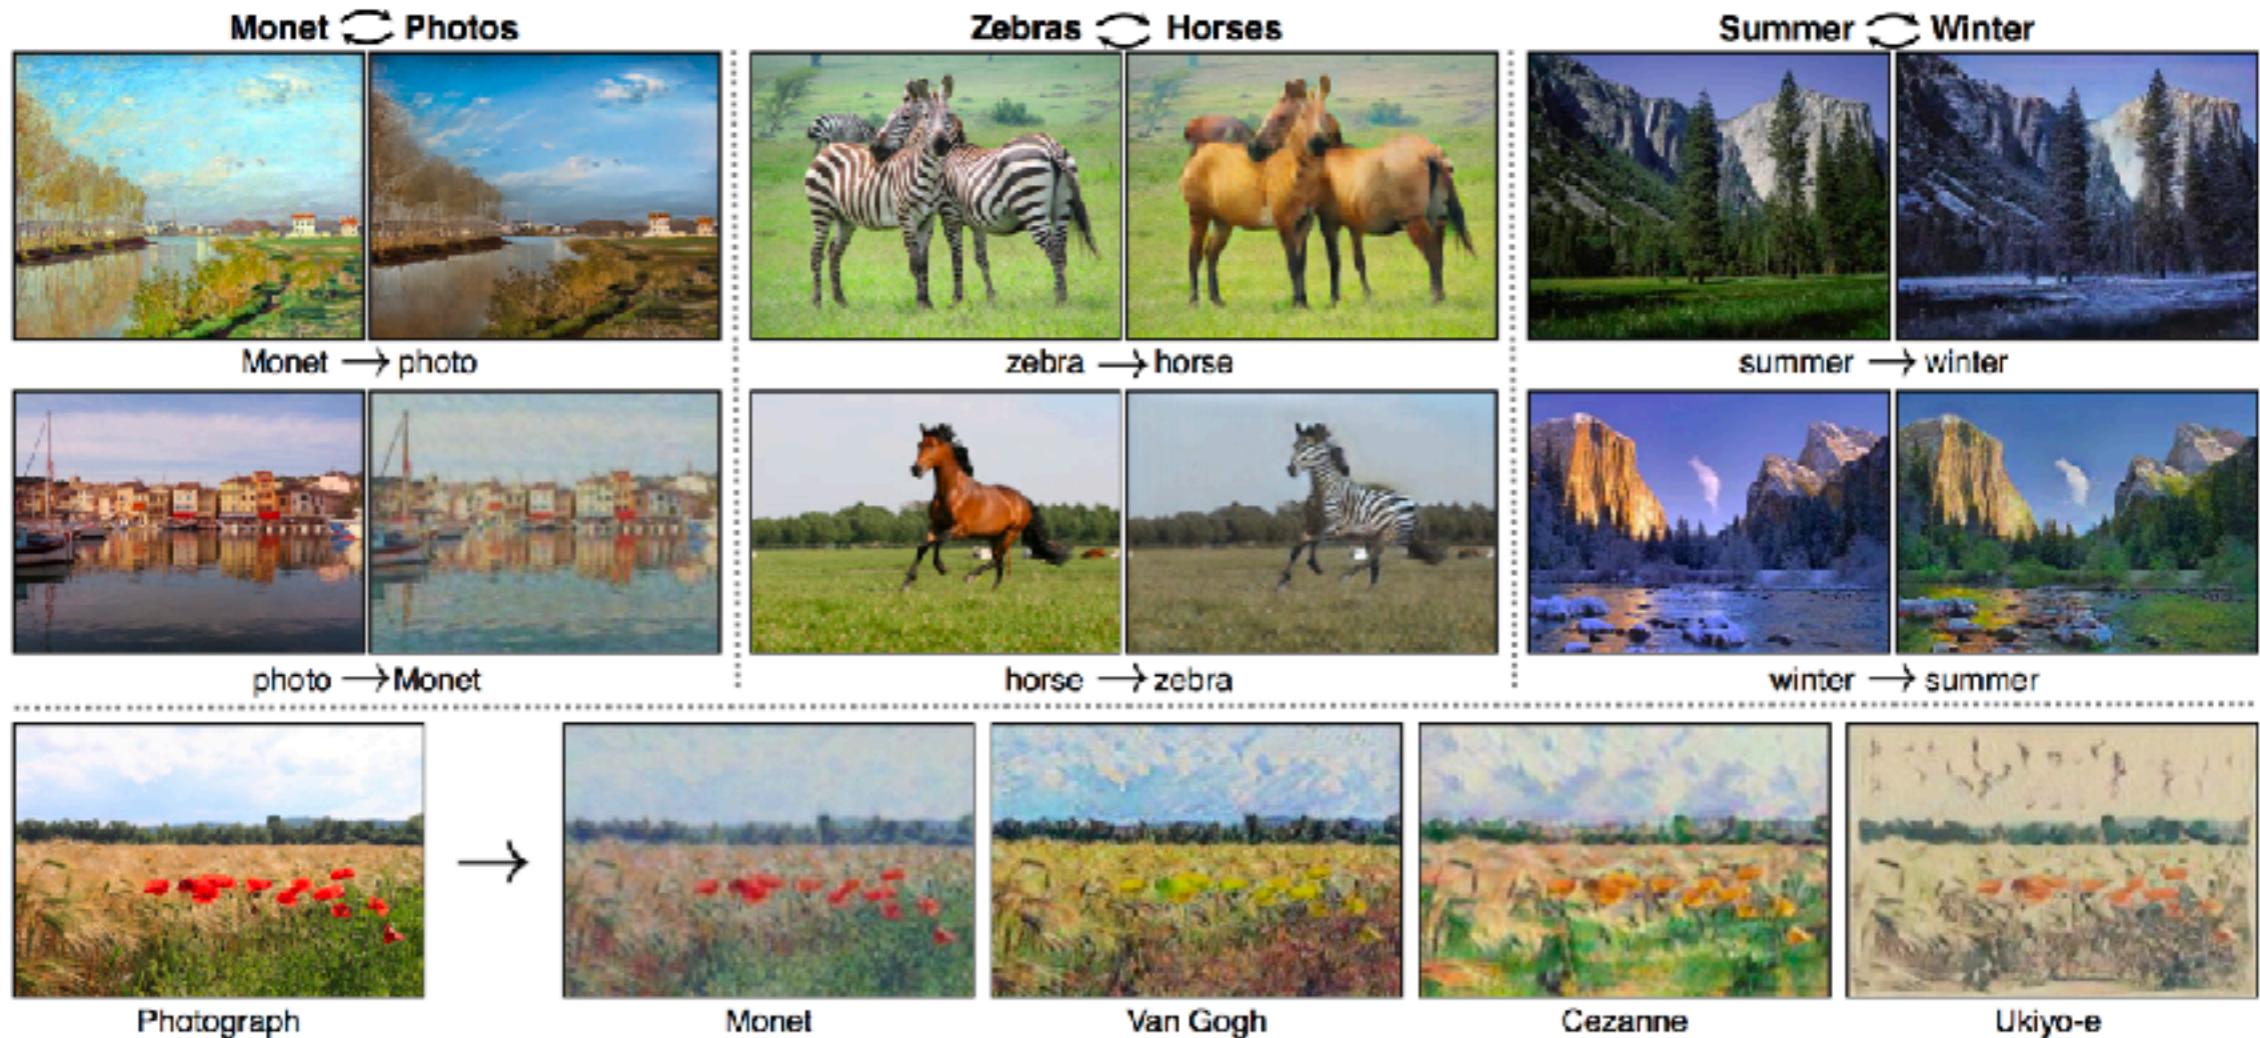
Figure 3. Illustration of our models on simplified one dimensional domains. (a) ideal mapping from domain A to domain B in which the two domain A modes map to two different domain B modes, (b) GAN model failure case, (c) GAN with reconstruction model failure case.

<https://arxiv.org/abs/1703.05192>



Generative Model

CycleGAN

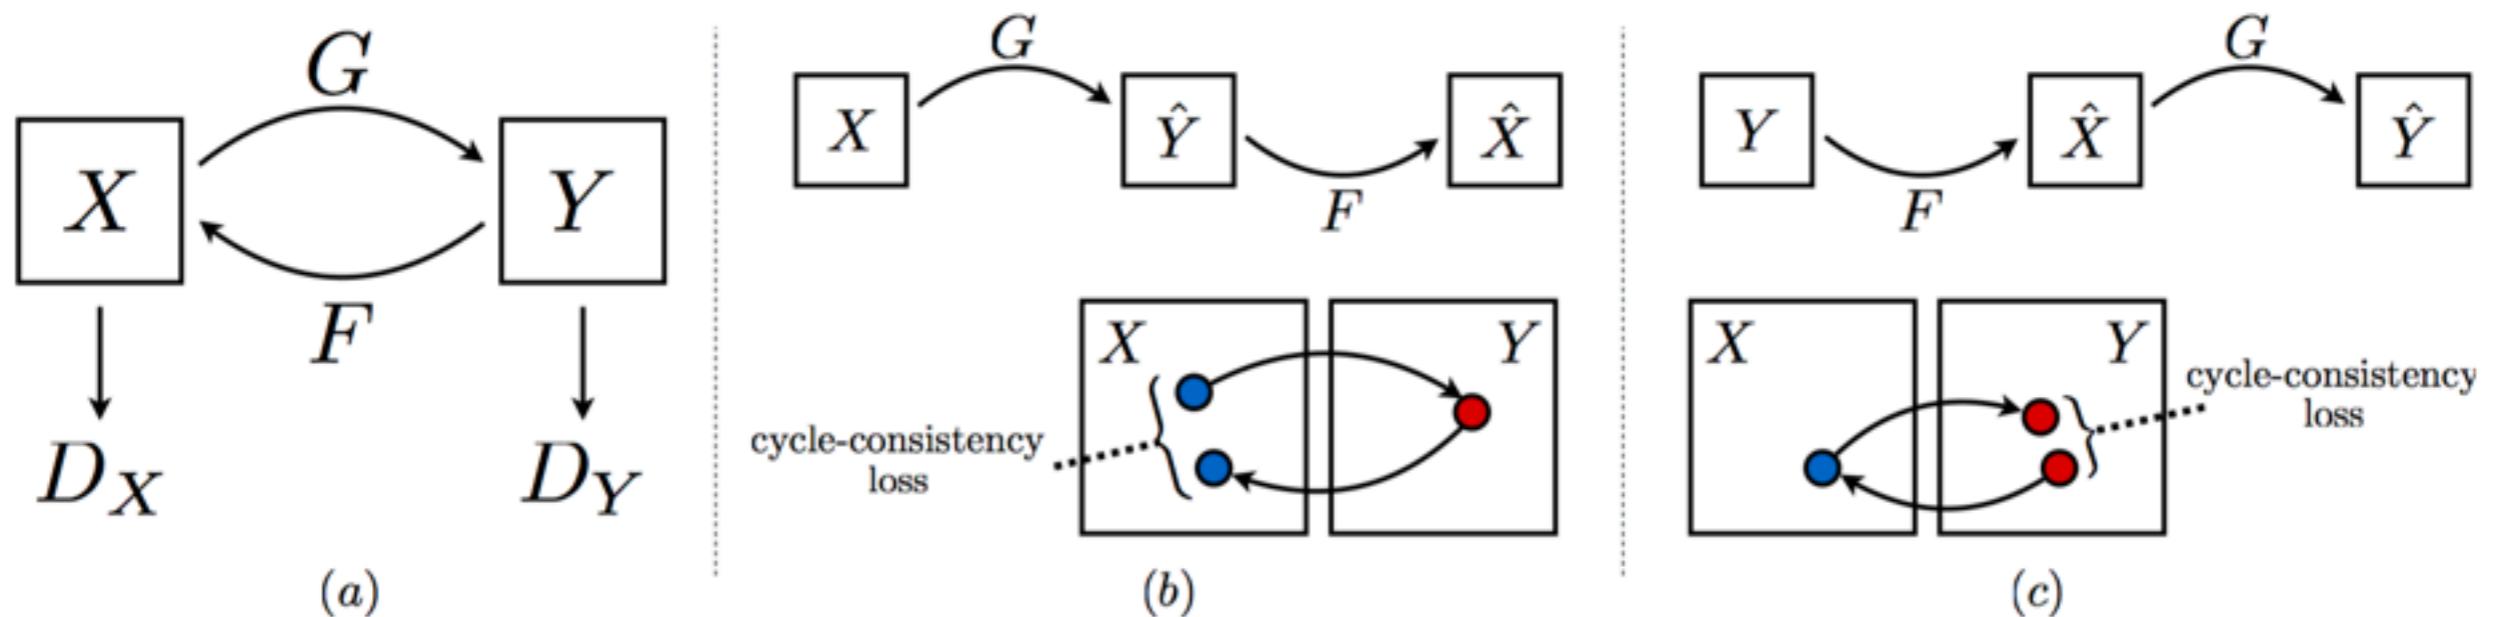


<https://arxiv.org/abs/1703.10593>



Generative Model

CycleGAN



<https://arxiv.org/abs/1703.10593>

Generative Model

StarGAN

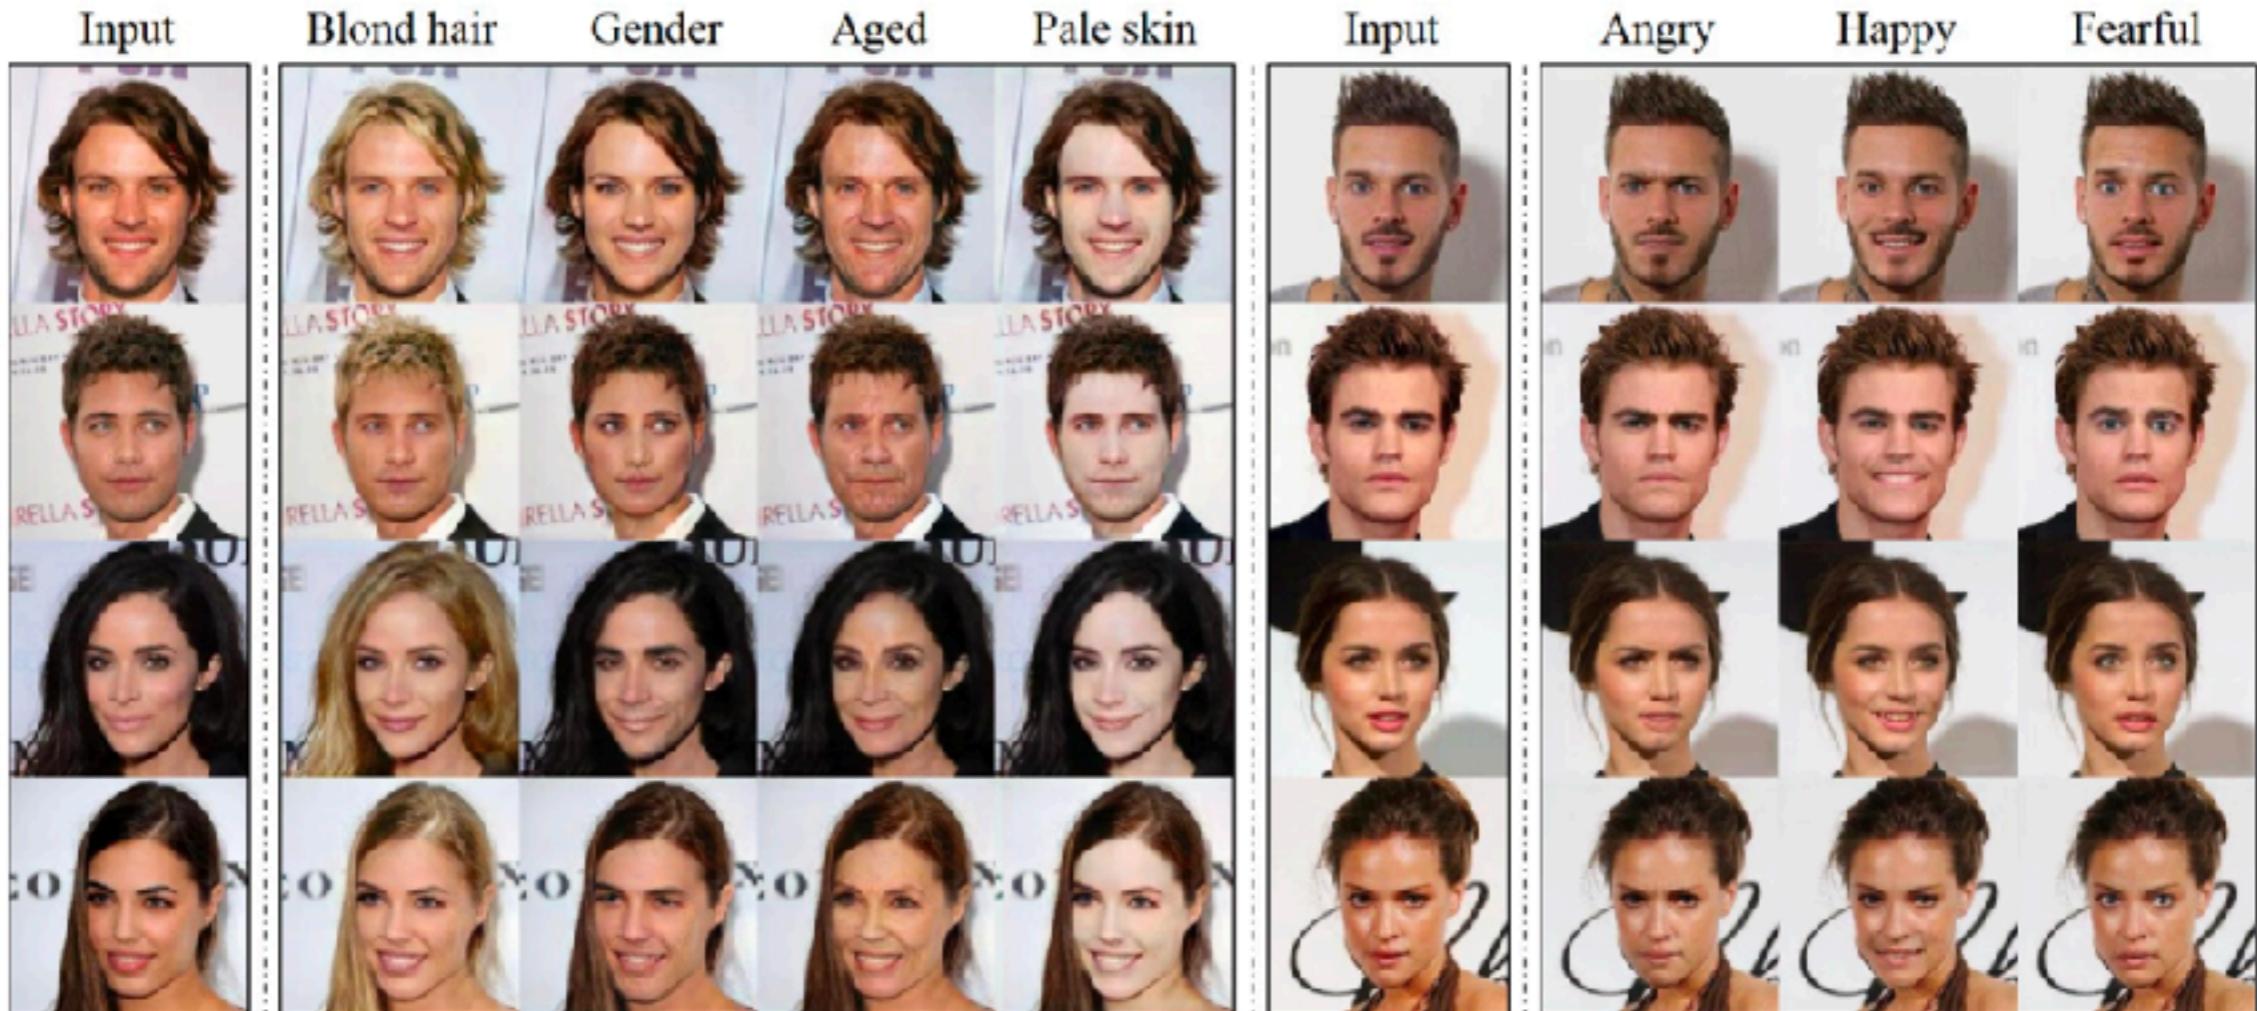


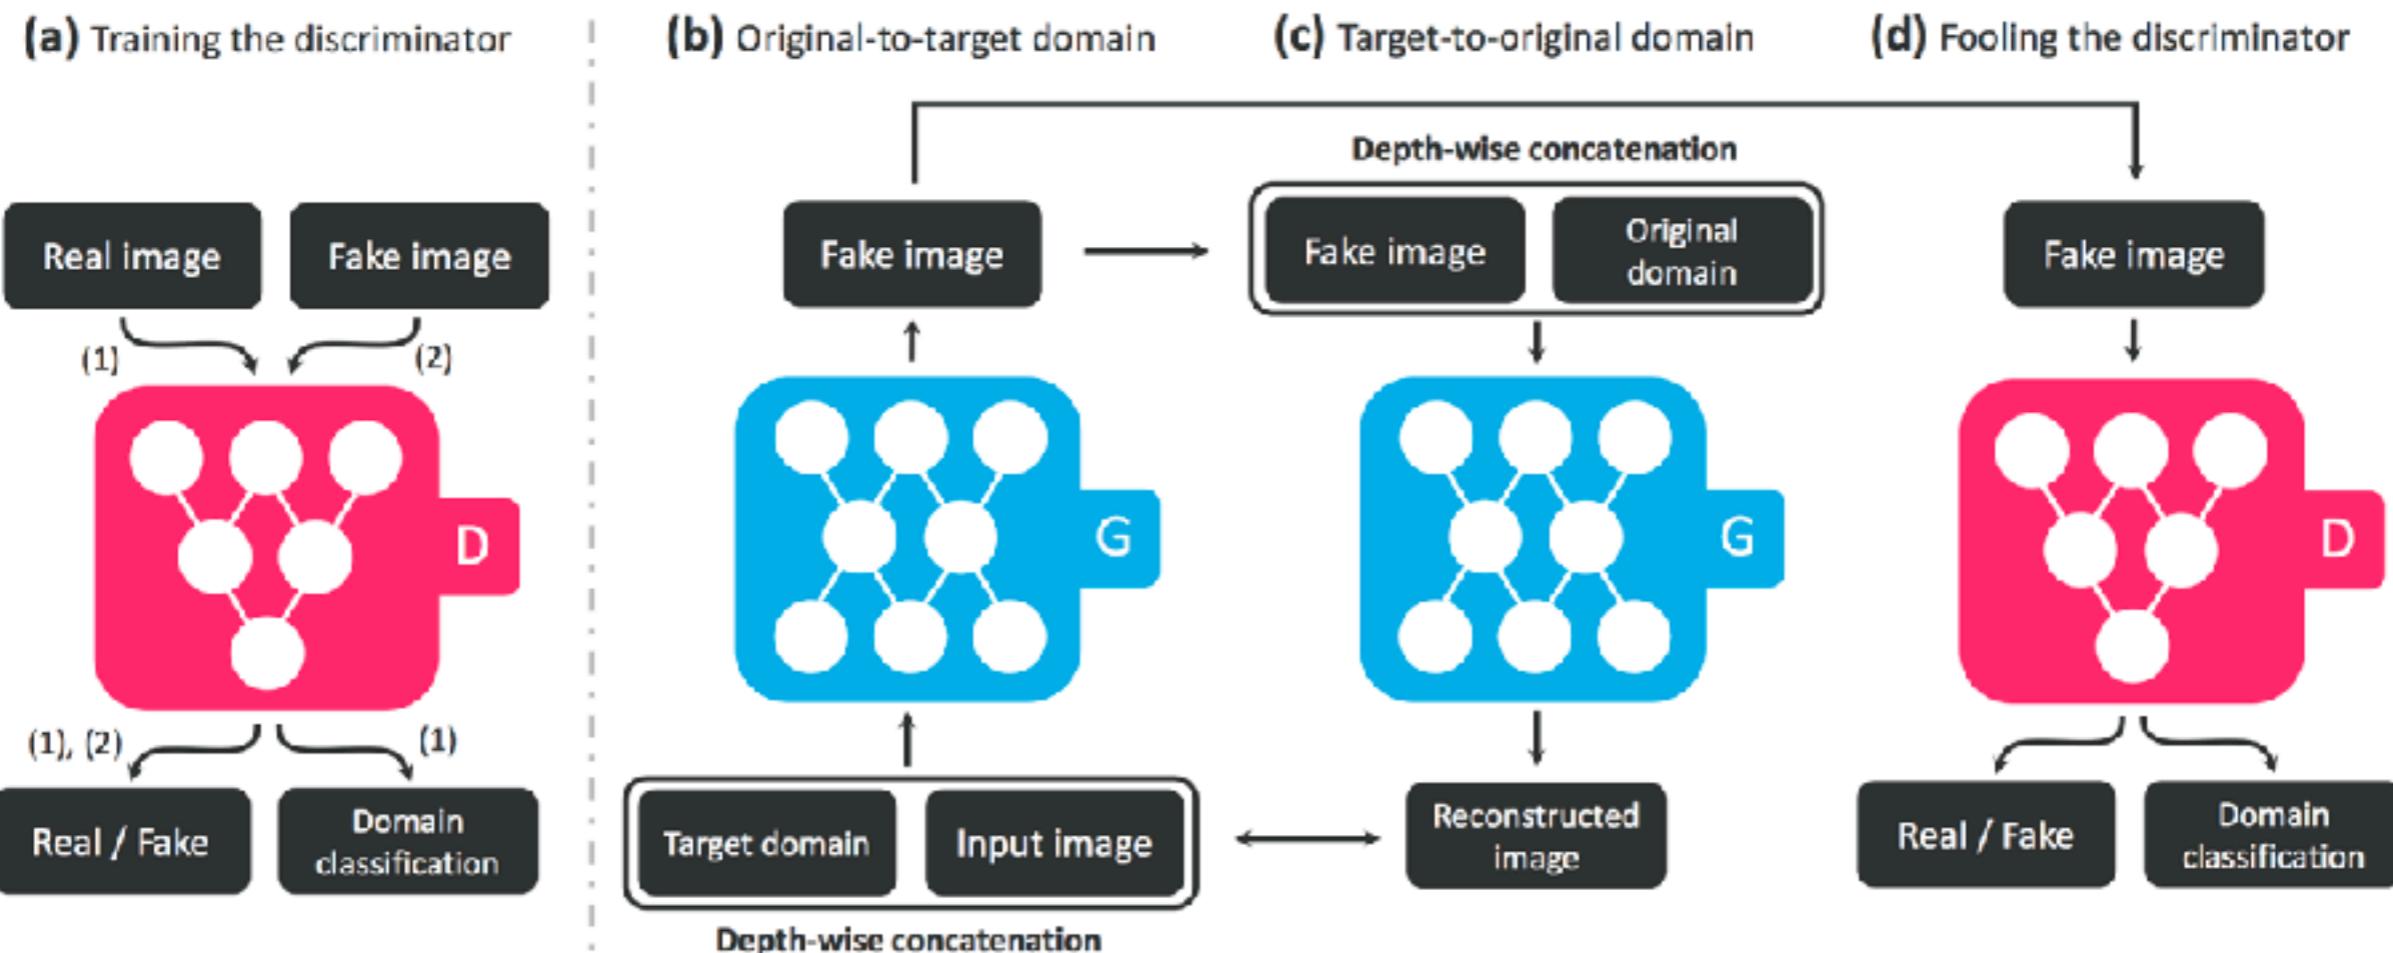
Figure 1. Multi-domain image-to-image translation results on the CelebA dataset via transferring knowledge learned from the RaFD dataset. The first and sixth columns show input images while the remaining columns are images generated by StarGAN. Note that the images are generated by a single generator network, and facial expression labels such as angry, happy, and fearful are from RaFD, not CelebA.

<https://arxiv.org/abs/1711.09020>



Generative Model

StarGAN



<https://arxiv.org/abs/1711.09020>



Generative Model

Progressive GAN

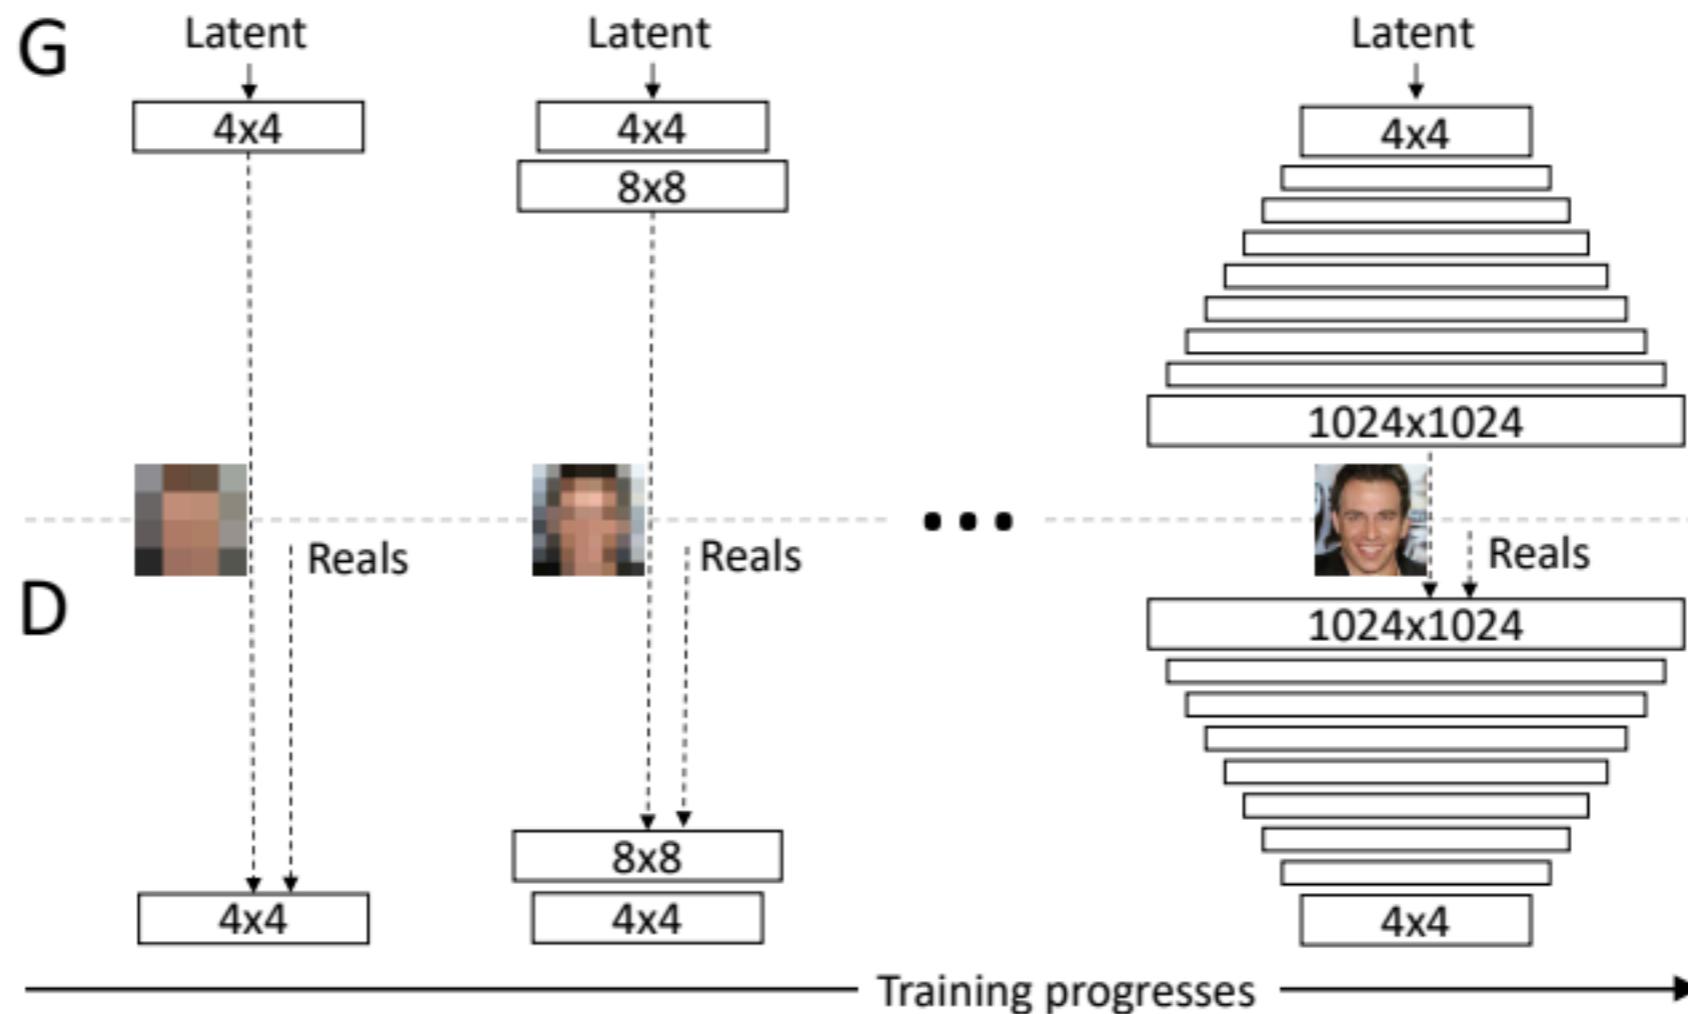


<https://arxiv.org/abs/1710.10196>



Generative Model

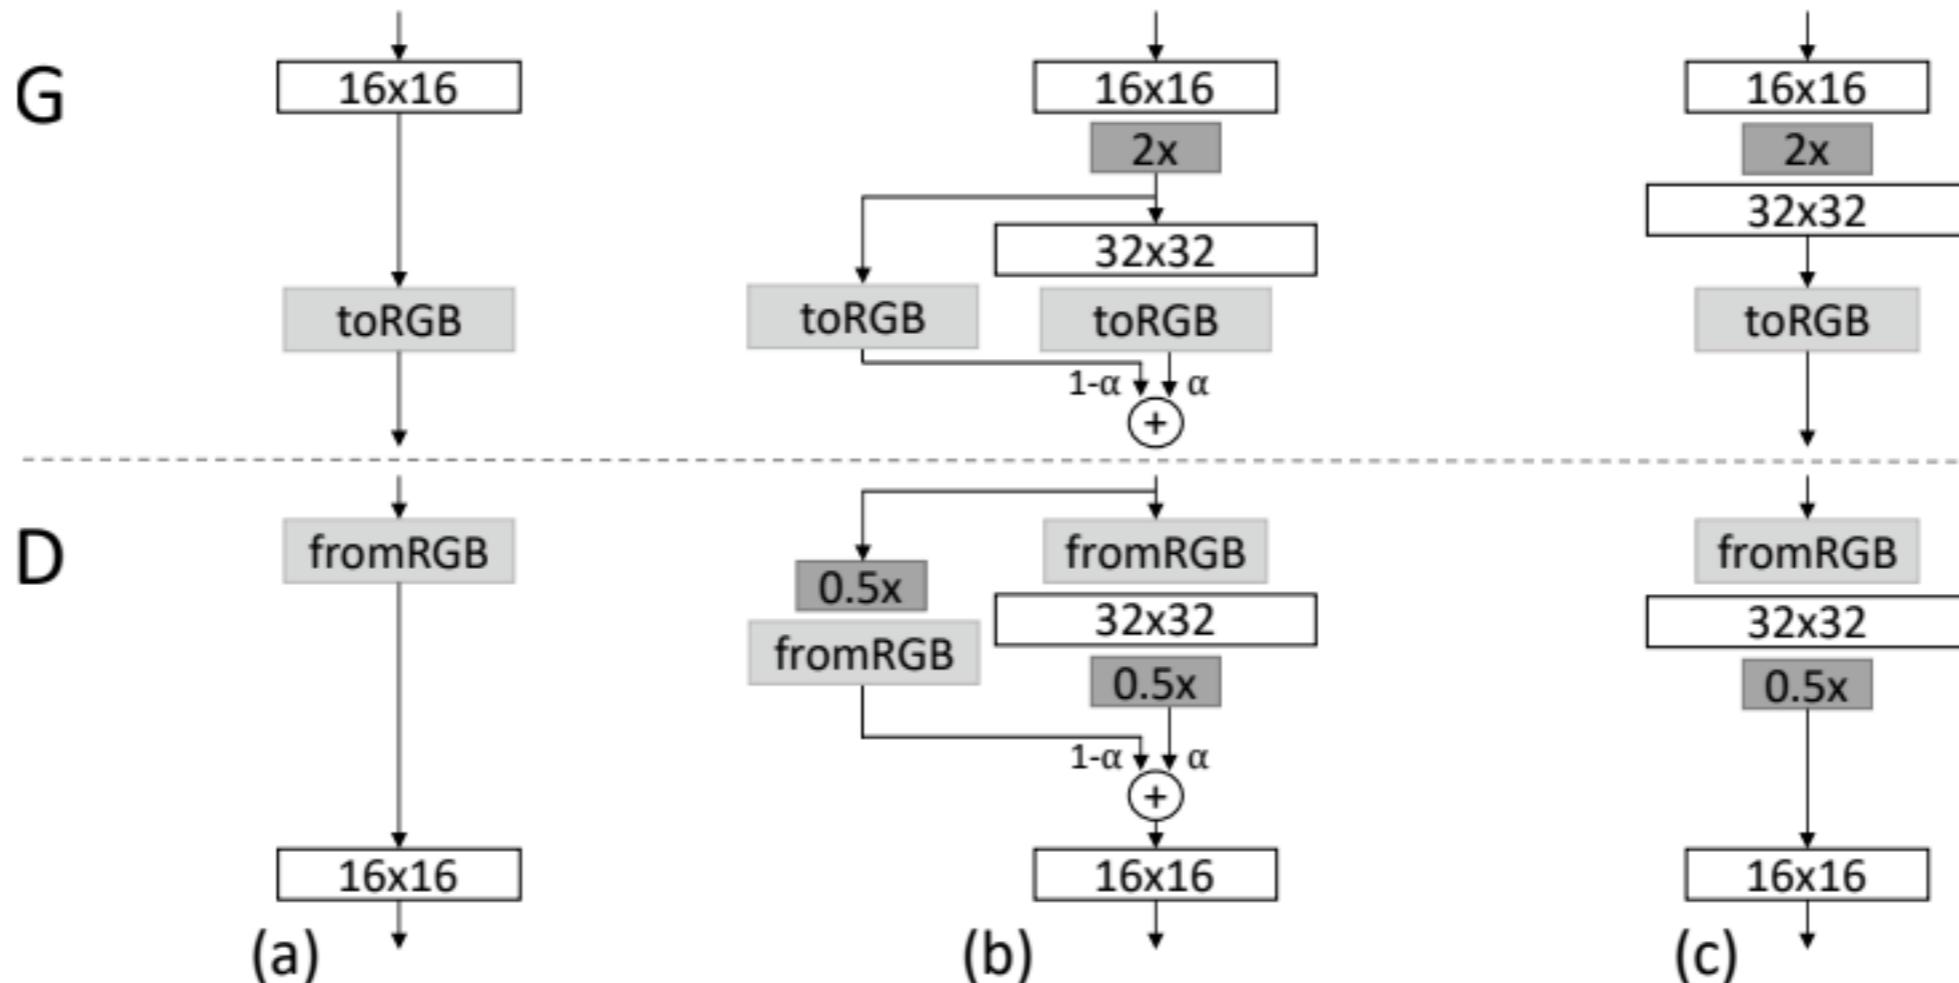
Progressive GAN



<https://arxiv.org/abs/1710.10196>

Generative Model

Progressive GAN



<https://arxiv.org/abs/1710.10196>

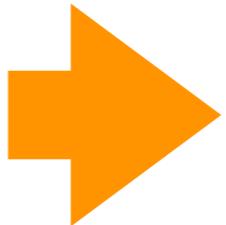


Domain Adaptation



Domain Adaptation

A step towards practical deployment



Domain Adaptation

Analysis for domain adaptation

Theorem 2 Let \mathcal{R} be a fixed representation function from \mathcal{X} to \mathcal{Z} and \mathcal{H} be a hypothesis space of VC-dimension d .

If a random labeled sample of size m is generated by applying \mathcal{R} to a \mathcal{D}_S - i.i.d. sample labeled according to f , and $\tilde{\mathcal{U}}_S, \tilde{\mathcal{U}}_T$ are unlabeled samples of size m' each, drawn from $\tilde{\mathcal{D}}_S$ and $\tilde{\mathcal{D}}_T$ respectively, then with probability at least $1 - \delta$ (over the choice of the samples), for every $h \in \mathcal{H}$:

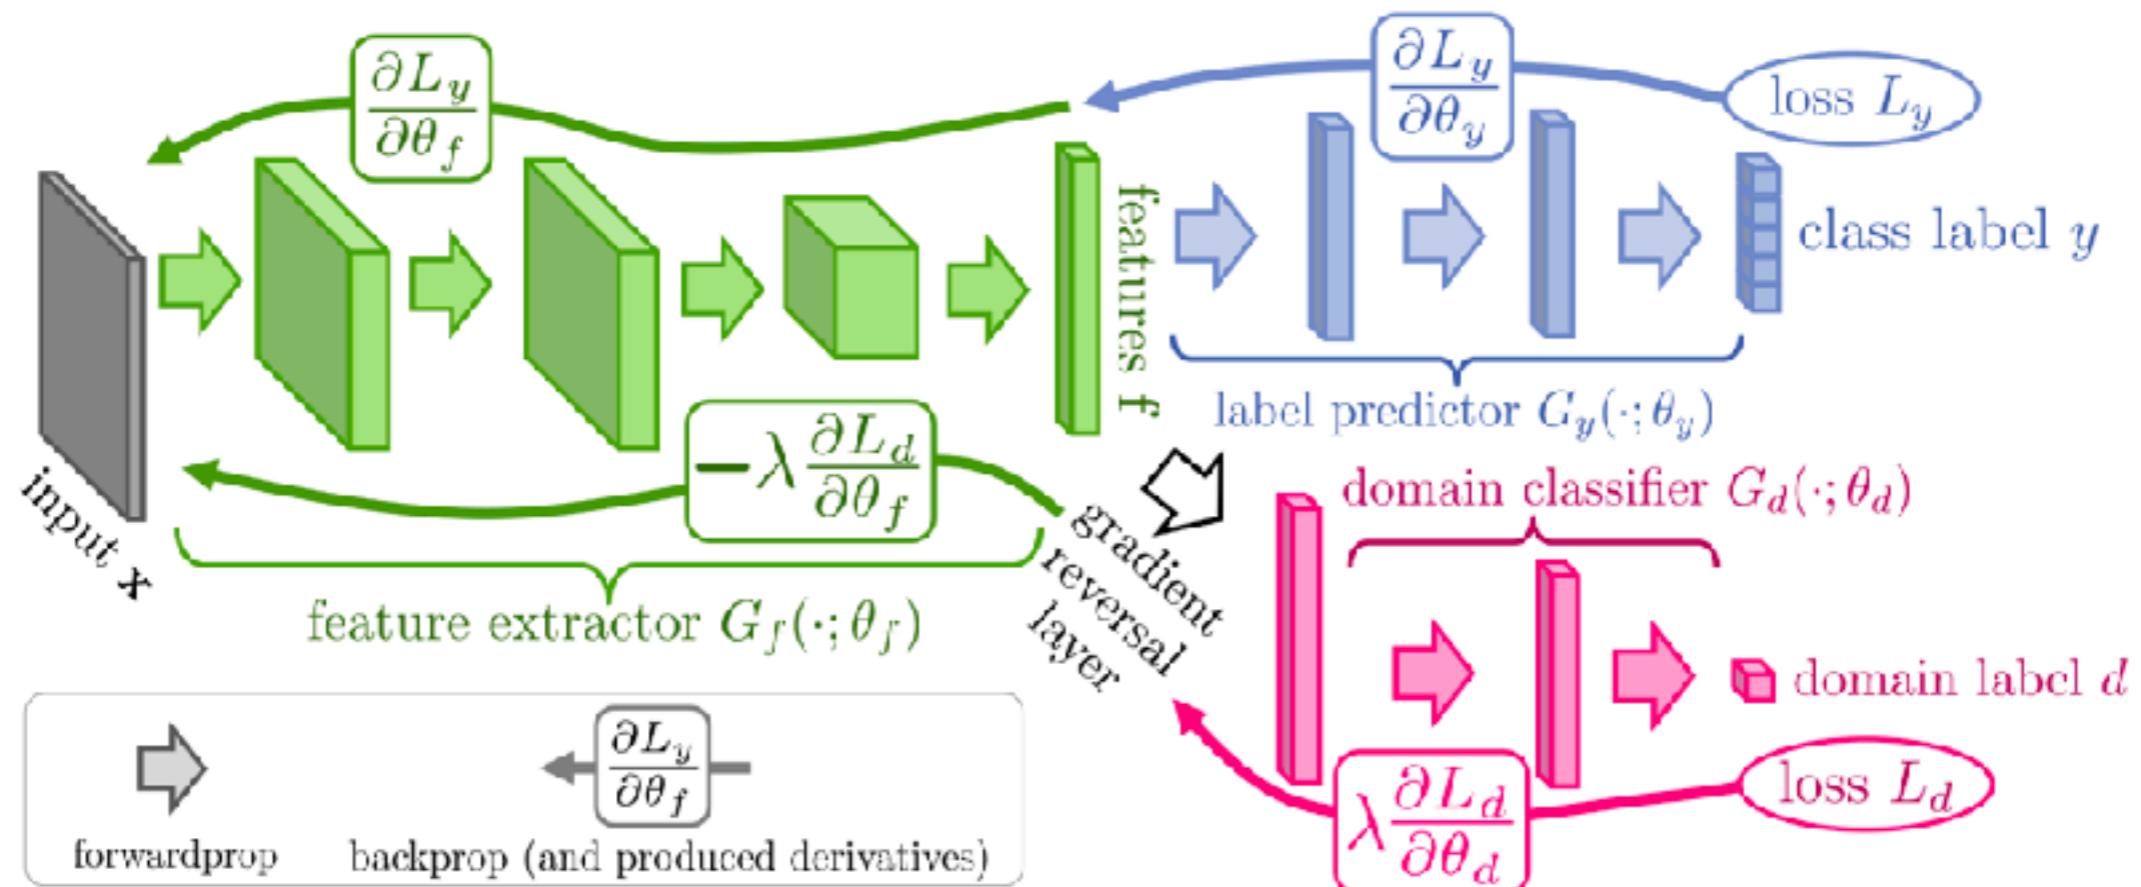
$$\epsilon_T(h) \leq \hat{\epsilon}_S(h) + \frac{4}{m} \sqrt{\left(d \log \frac{2em}{d} + \log \frac{4}{\delta} \right)} + \lambda + d_{\mathcal{H}}(\tilde{\mathcal{U}}_S, \tilde{\mathcal{U}}_T) + 4 \sqrt{\frac{d \log(2m') + \log(\frac{4}{\delta})}{m'}}$$

<http://www.john.blitzer.com/papers/nips06.pdf>



Domain Adaptation

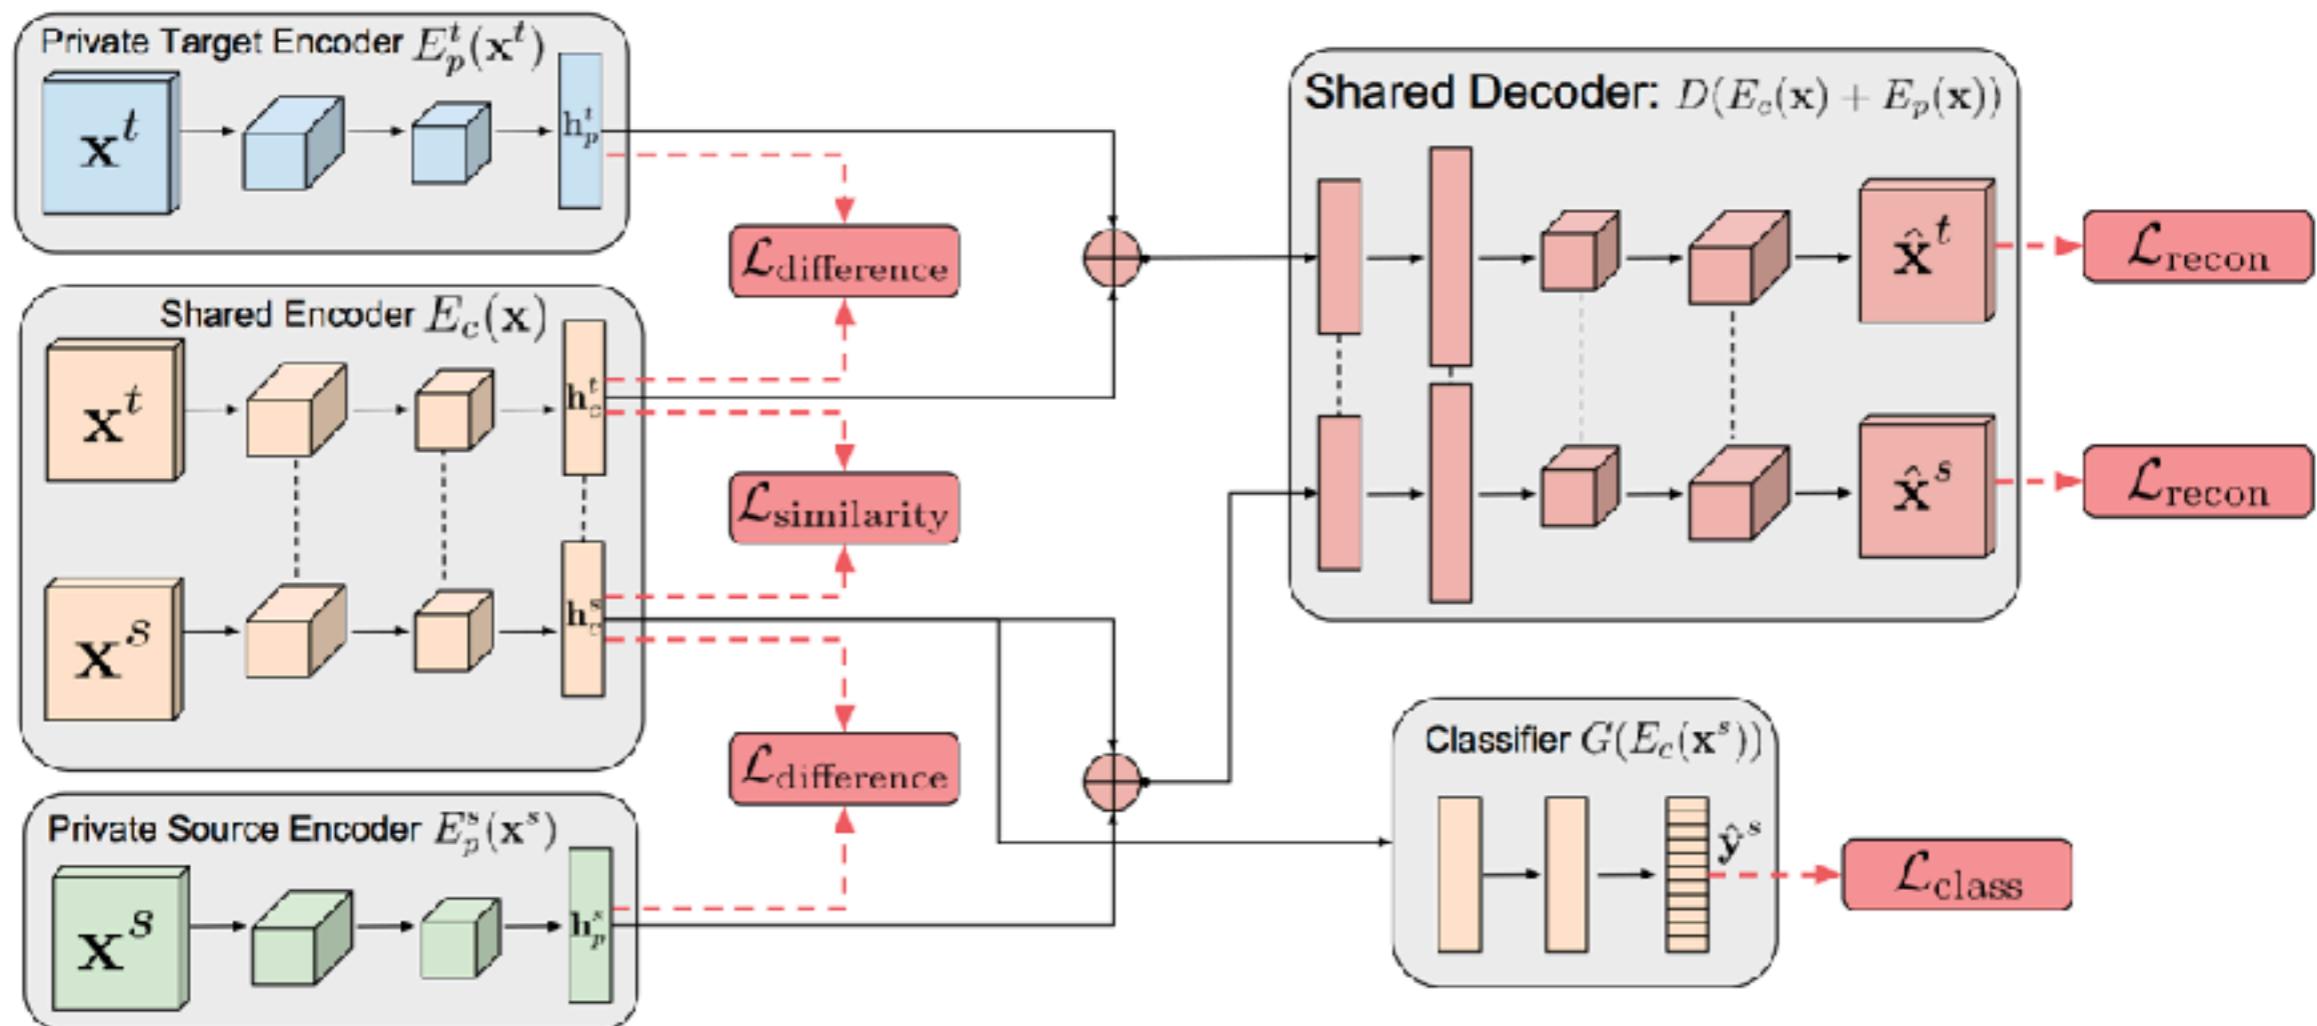
Domain Adversarial Neural Network (DANN)



<https://arxiv.org/abs/1505.07818>

Domain Adaptation

Domain Separation Network (DSN)



<https://arxiv.org/abs/1608.06019>

Domain Adaptation

DANN & DSN

Domain-Adversarial Training of Neural Networks.

$$E(\theta_f, \theta_y, \theta_d) = \frac{1}{n} \sum_{i=1}^n L_g(\theta_f, \theta_y) - \lambda \left(\frac{1}{n} \sum_{i=1}^n L_d(\theta_f, \theta_d) + \frac{1}{m} \sum_{i=1}^m L_t(\theta_f, \theta_d) \right)$$

Domain Separation Networks

$$L = L_{\text{task}} + \alpha L_{\text{recon}} + \beta L_{\text{diff}} + \gamma L_{\text{sim}}$$

$L_{\text{task}} = - \sum_{i=1}^{N_s} y_i^s \log \hat{y}_i$

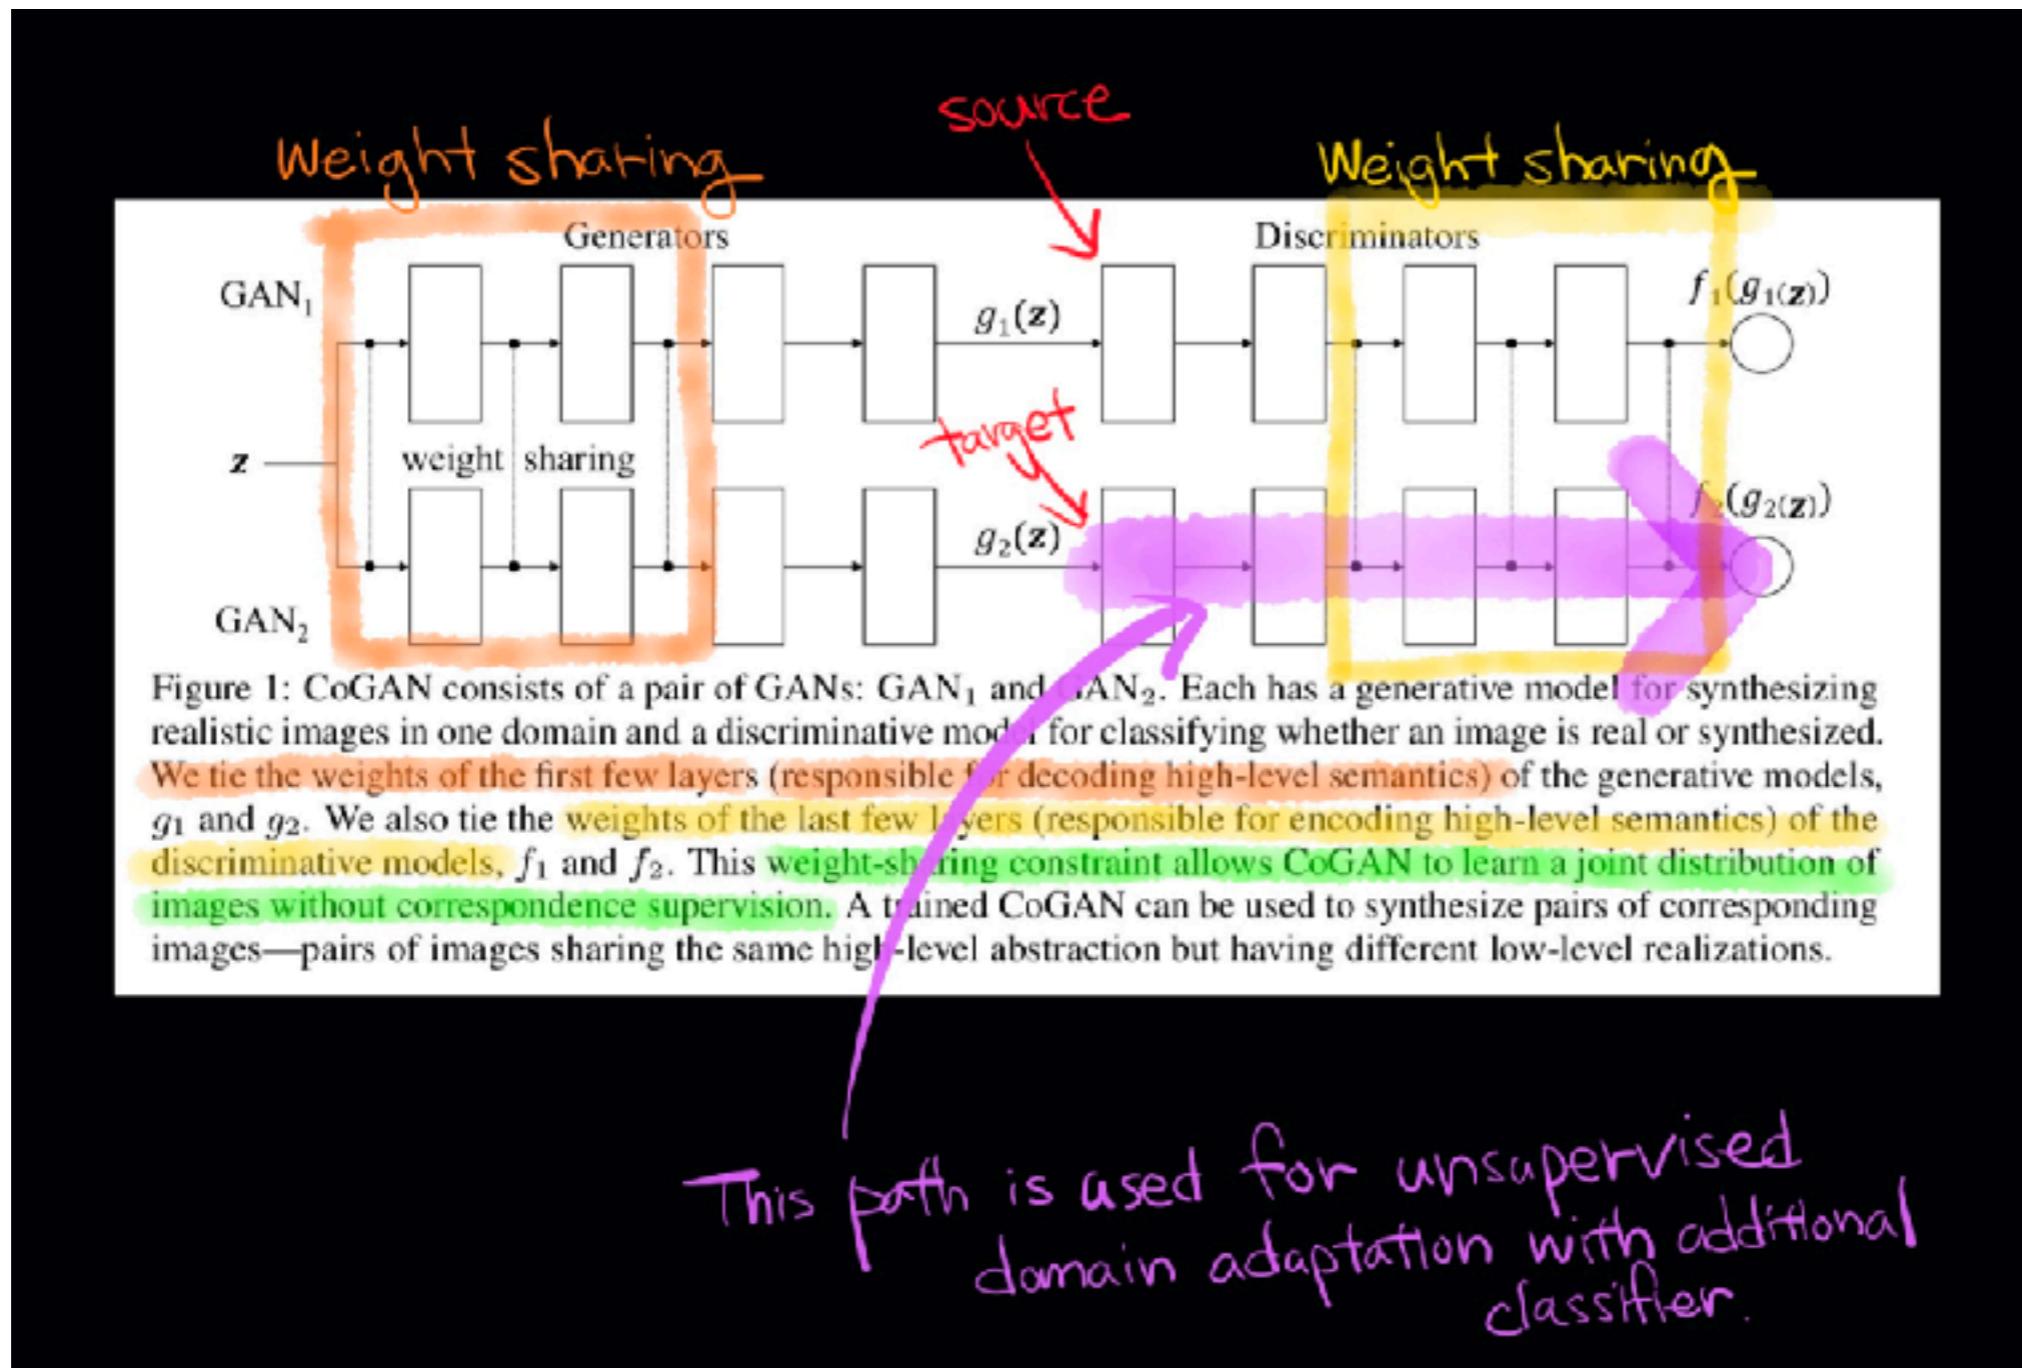
$L_{\text{recon}} = \sum_{i=1}^{N_s} \left(\frac{1}{K} \|y^s - \hat{y}^s\|_2^2 - \frac{1}{K} \|([x^s - \hat{x}^s] \cdot \mathbf{1}_K)\|^2 \right) + \sum_{i=1}^{N_t} \left(\frac{1}{K} \|x^t - \hat{x}^t\|_2^2 - \frac{1}{K} \|([x^t - \hat{x}^t] \cdot \mathbf{1}_K)\|^2 \right)$

$L_{\text{diff}} = \|H_c^{sT} H_p^s\|_F^2 + \|H_c^{tT} H_p^t\|_F^2$ (H_c: shared feat. matrx, H_p: private feat. matrx = $\begin{bmatrix} \vdots \\ \vdots \\ \vdots \end{bmatrix}$)

$L_{\text{sim}} = \sum_{i=1}^{N_{\text{mix}}} \{d_i \log \hat{d}_i + (1-d_i) \log (1-\hat{d}_i)\}$

Domain Adaptation

Coupled GAN

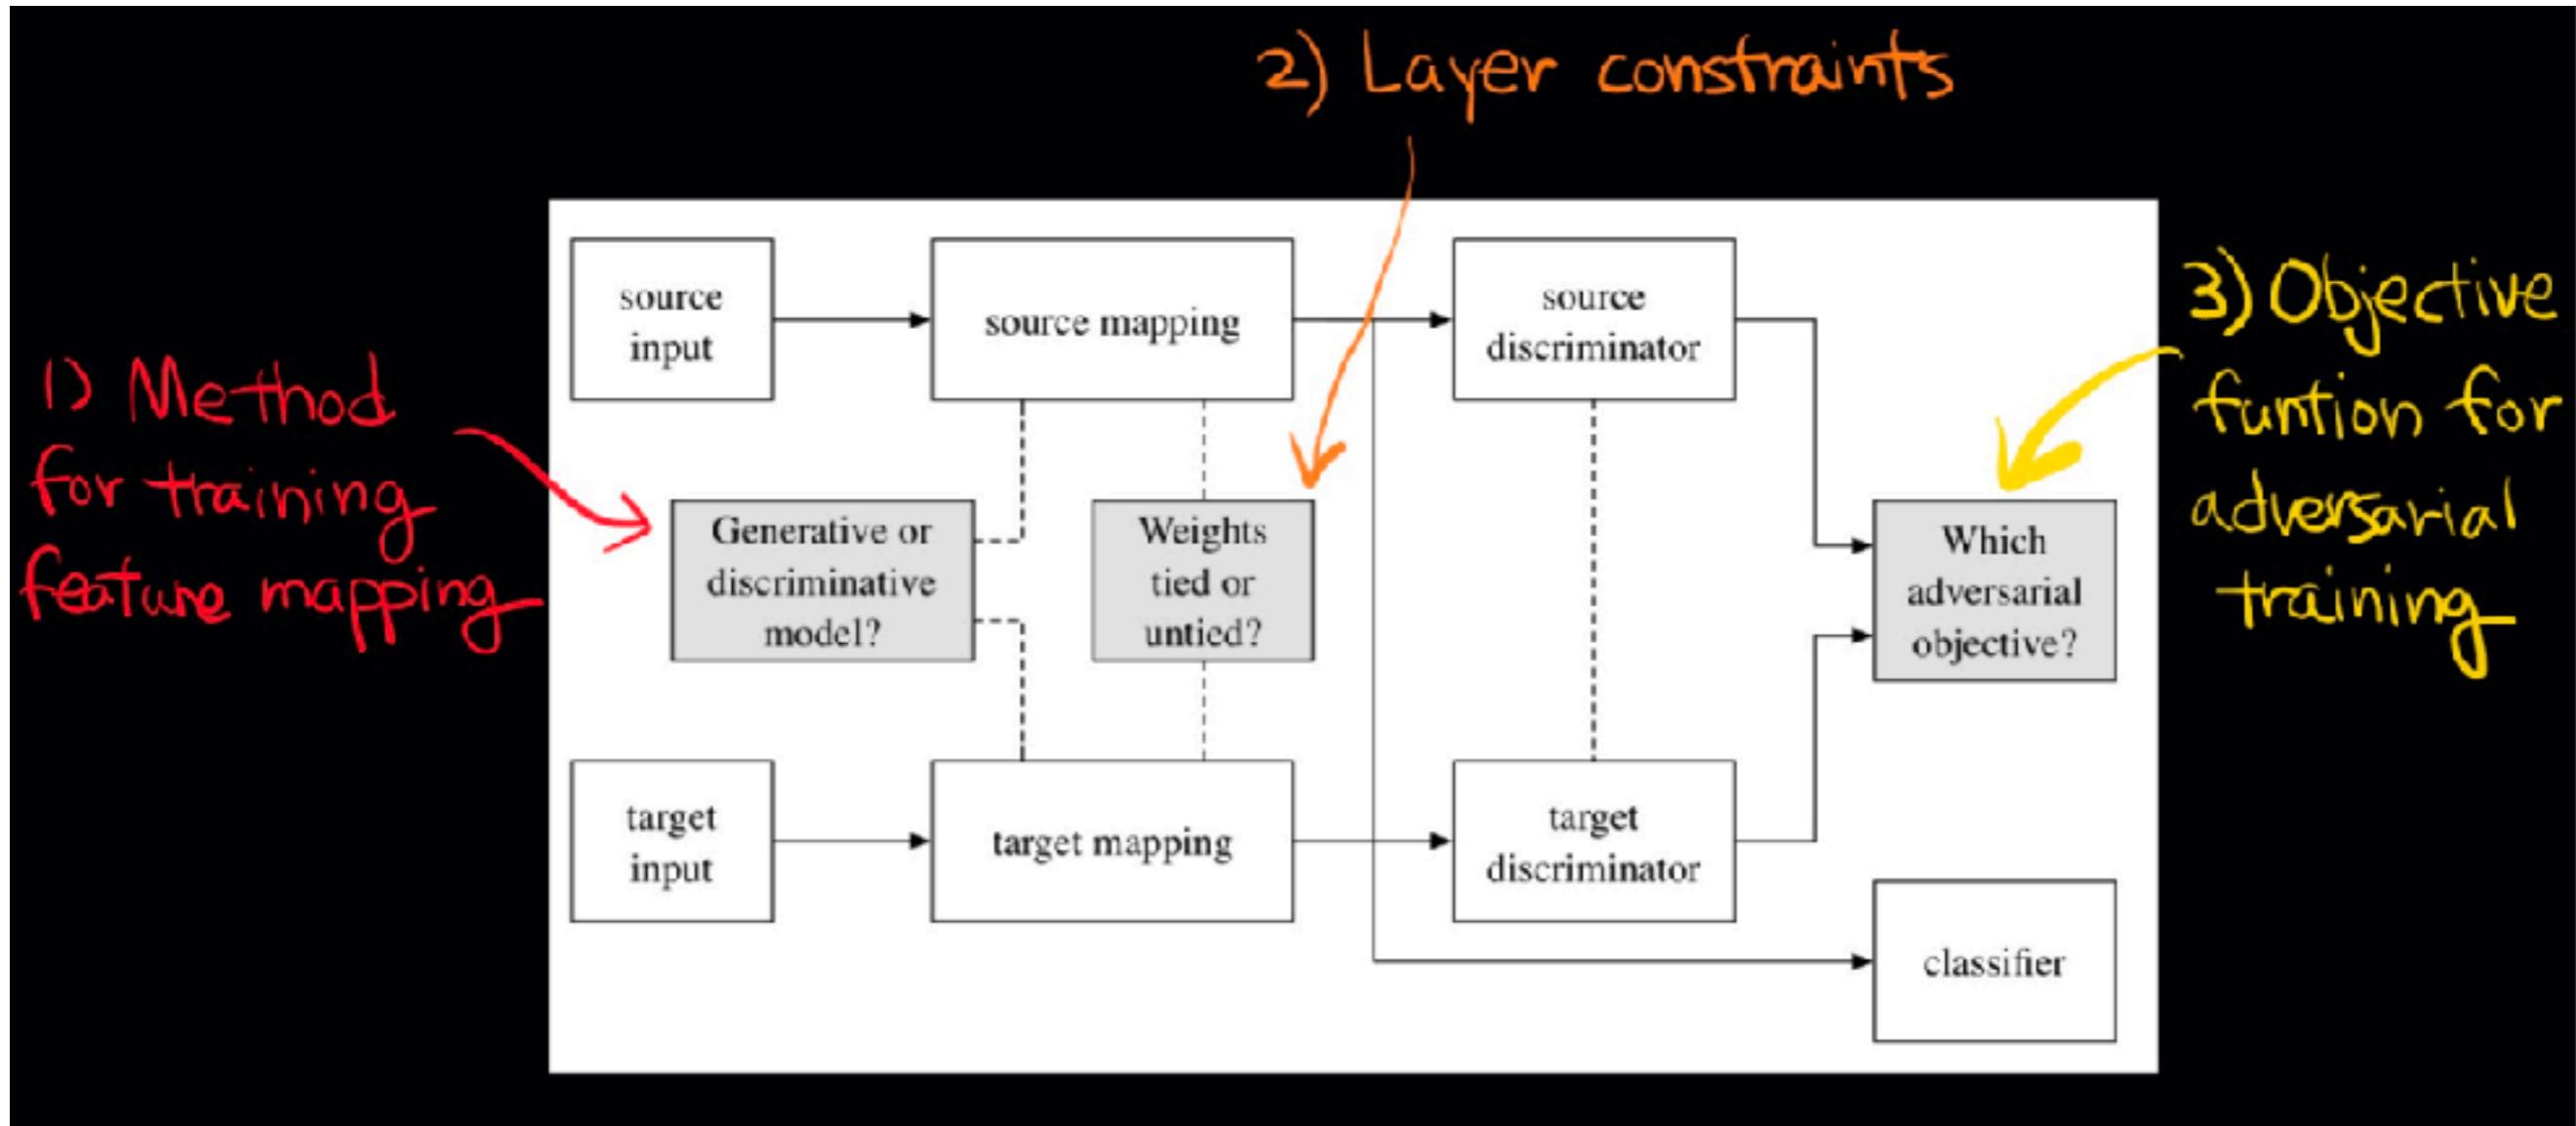


<https://arxiv.org/abs/1606.07536>



Domain Adaptation

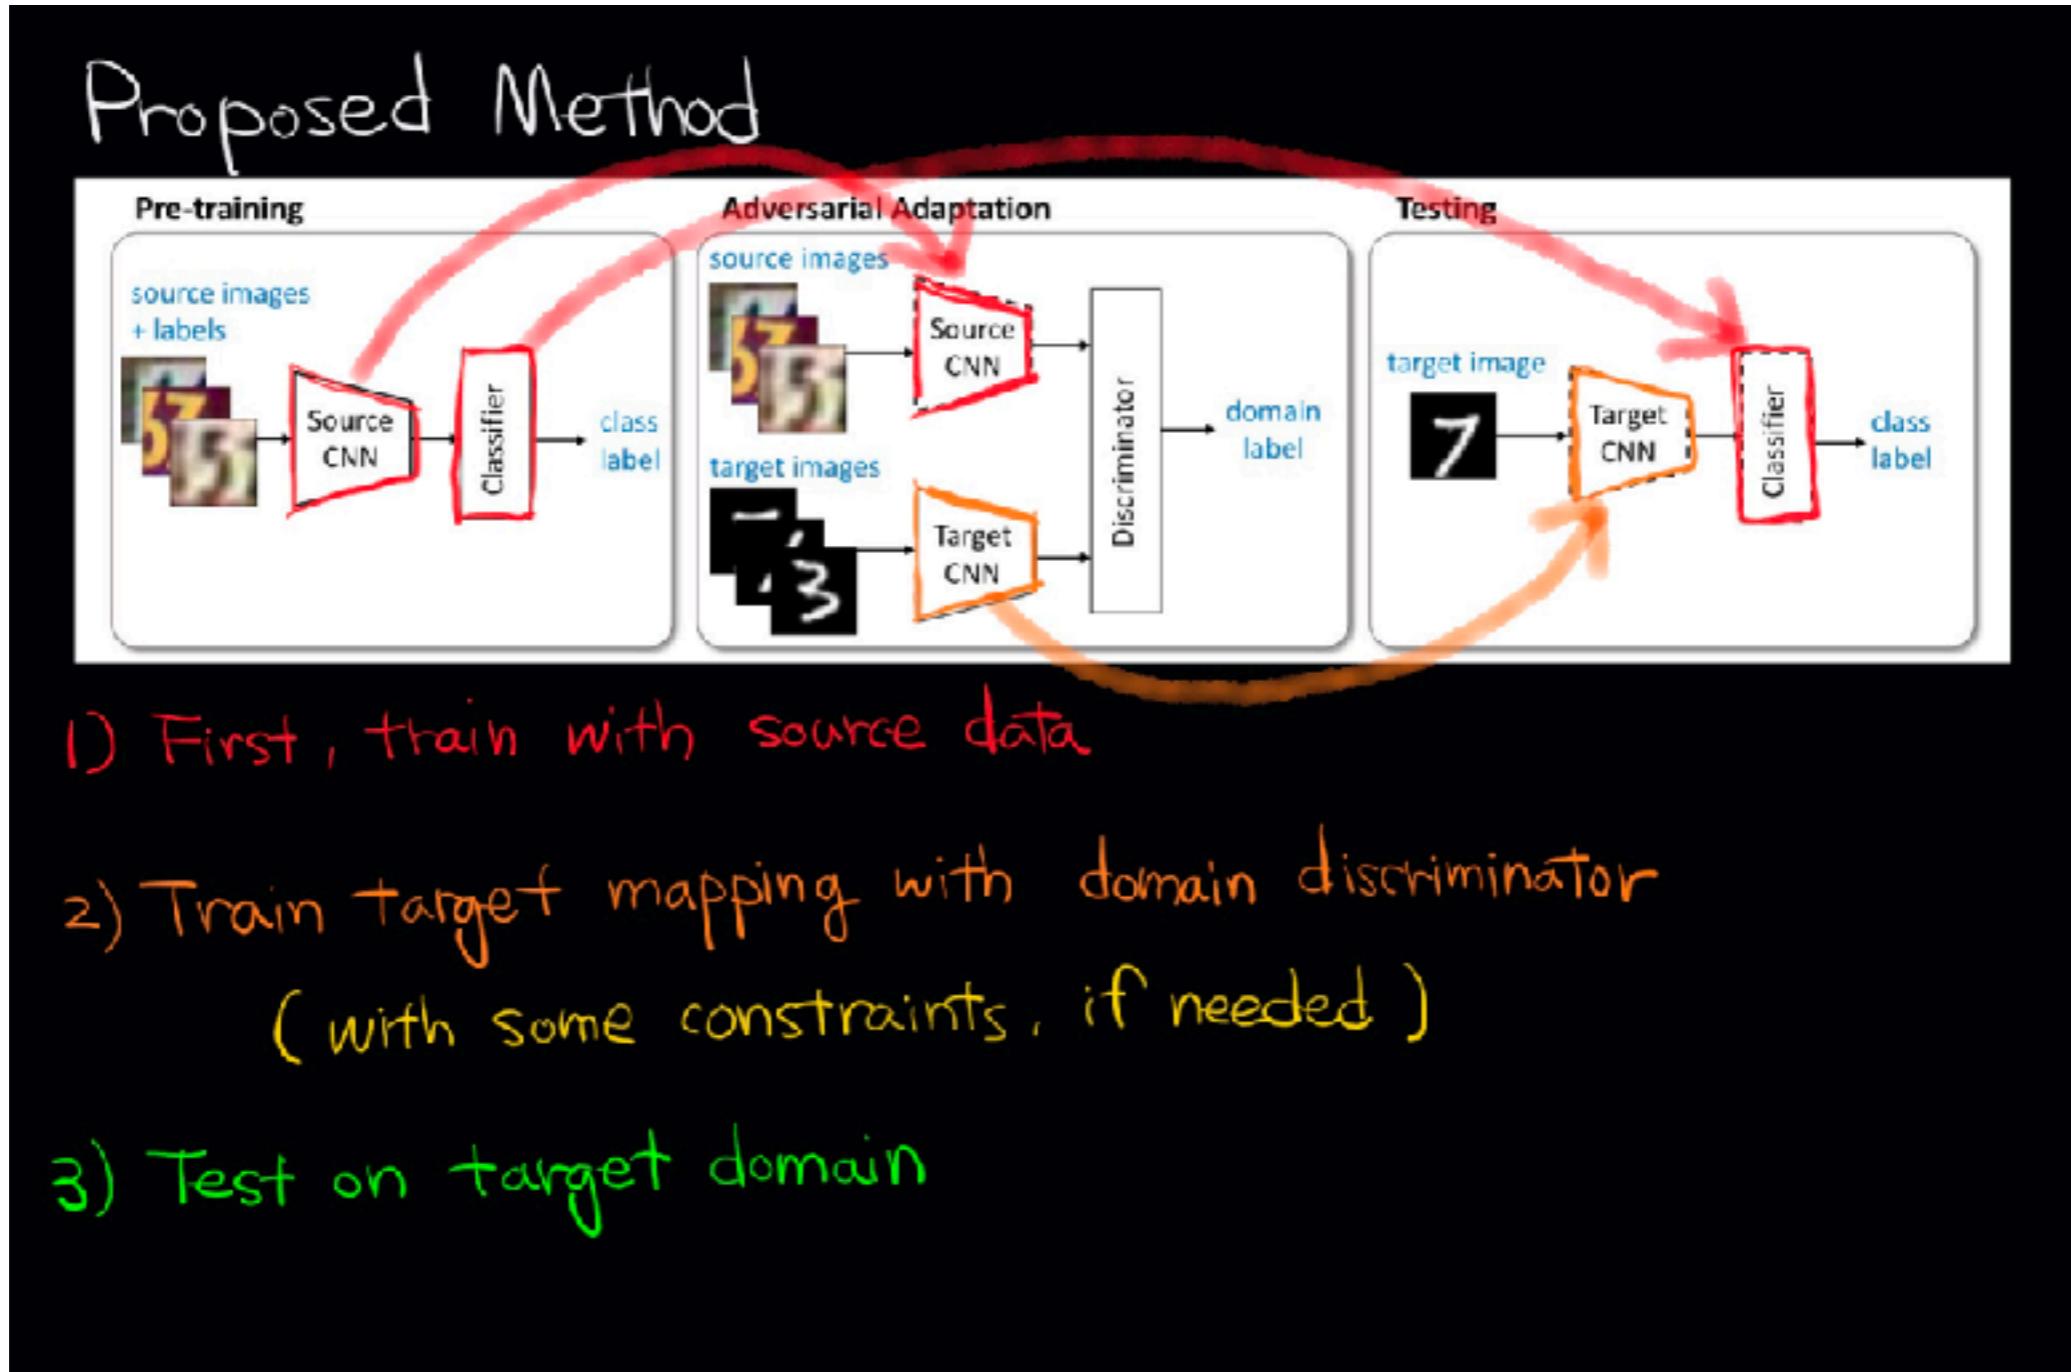
Adversarial Discriminative Domain Adaptation



<https://arxiv.org/abs/1606.07536>

Domain Adaptation

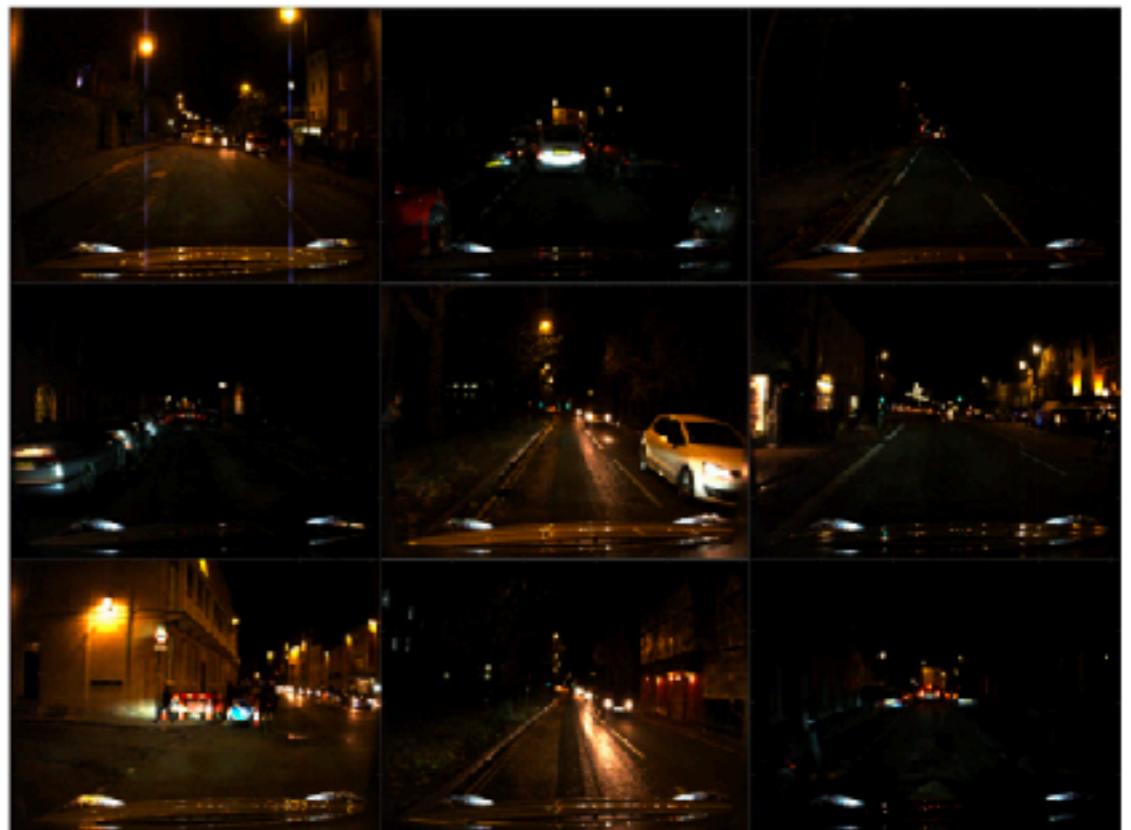
Adversarial Discriminative Domain Adaptation



<https://arxiv.org/abs/1606.07536>

Domain Adaptation

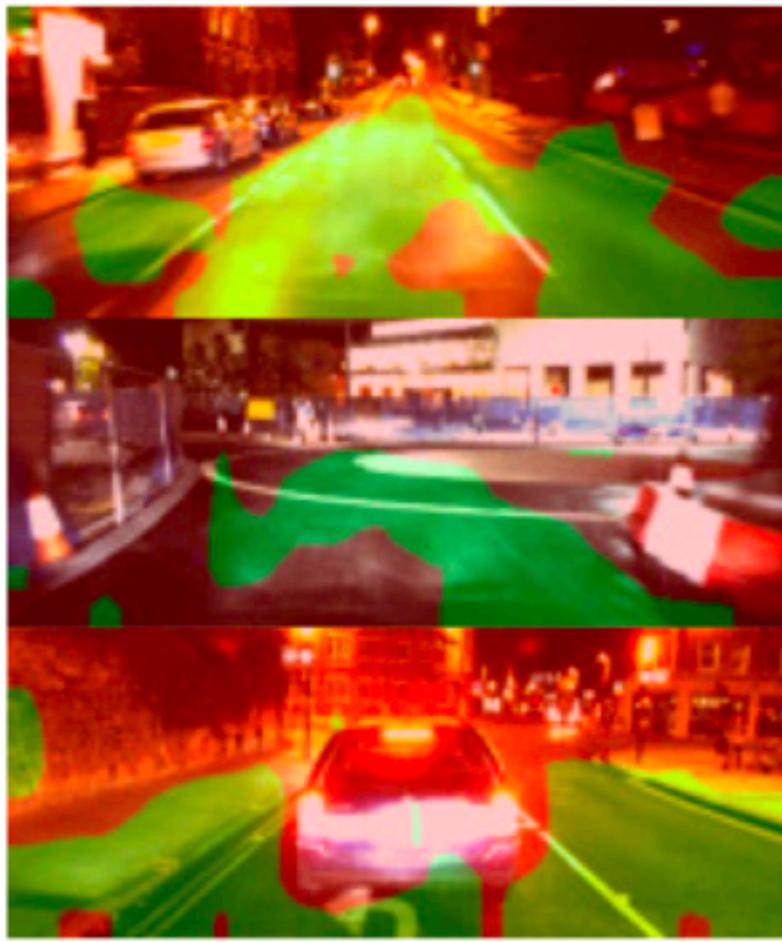
Appearance adaptation



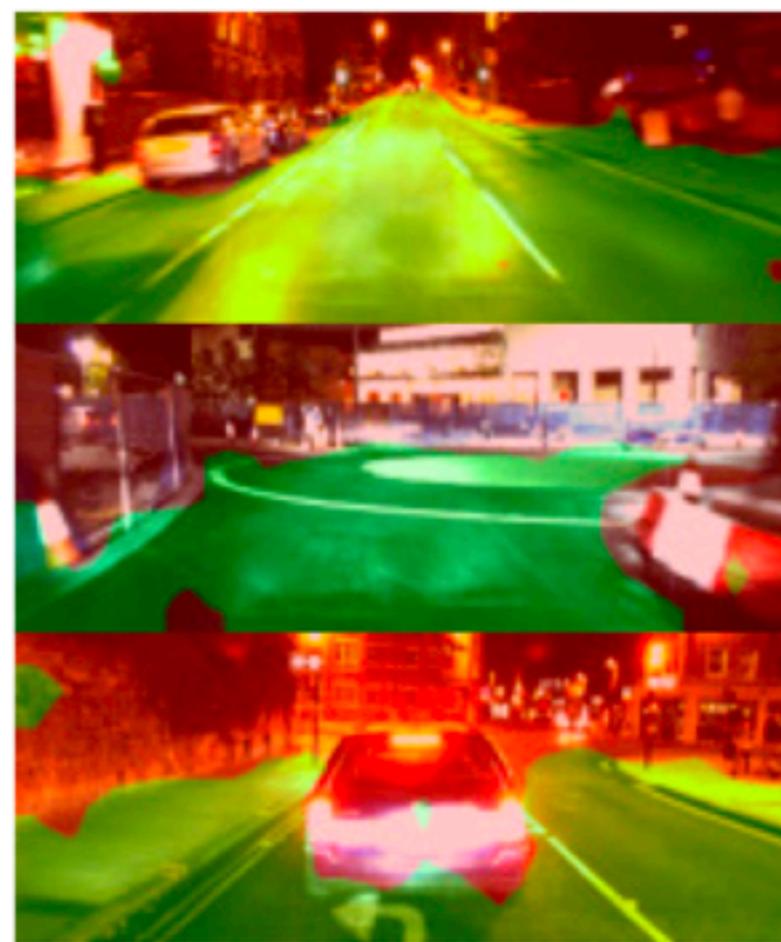
<https://arxiv.org/abs/1703.01461>

Domain Adaptation

Appearance adaptation



FCV-VGG16

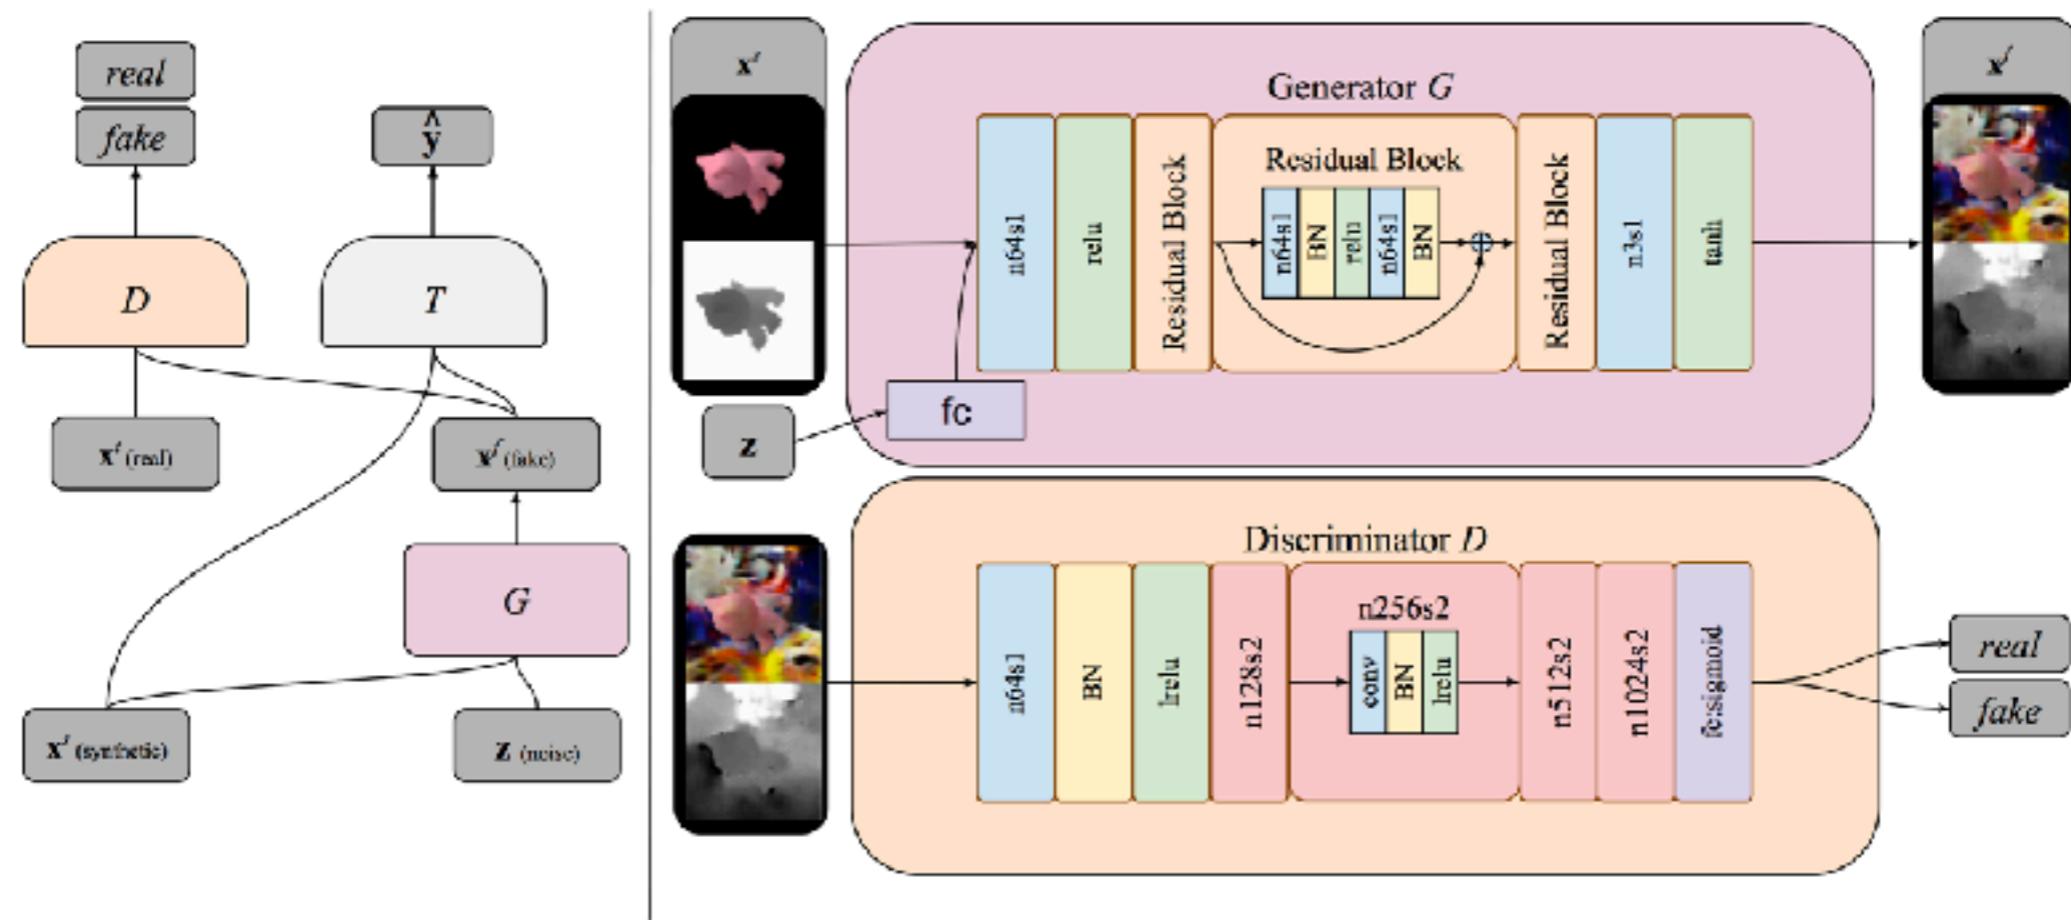


FCV-VGG16 w ADA

<https://arxiv.org/Abs/1703.01461>

Domain Adaptation

DA-GAN



<https://arxiv.org/abs/1612.05424>



Domain Adaptation

DA-GAN



<https://arxiv.org/abs/1612.05424>

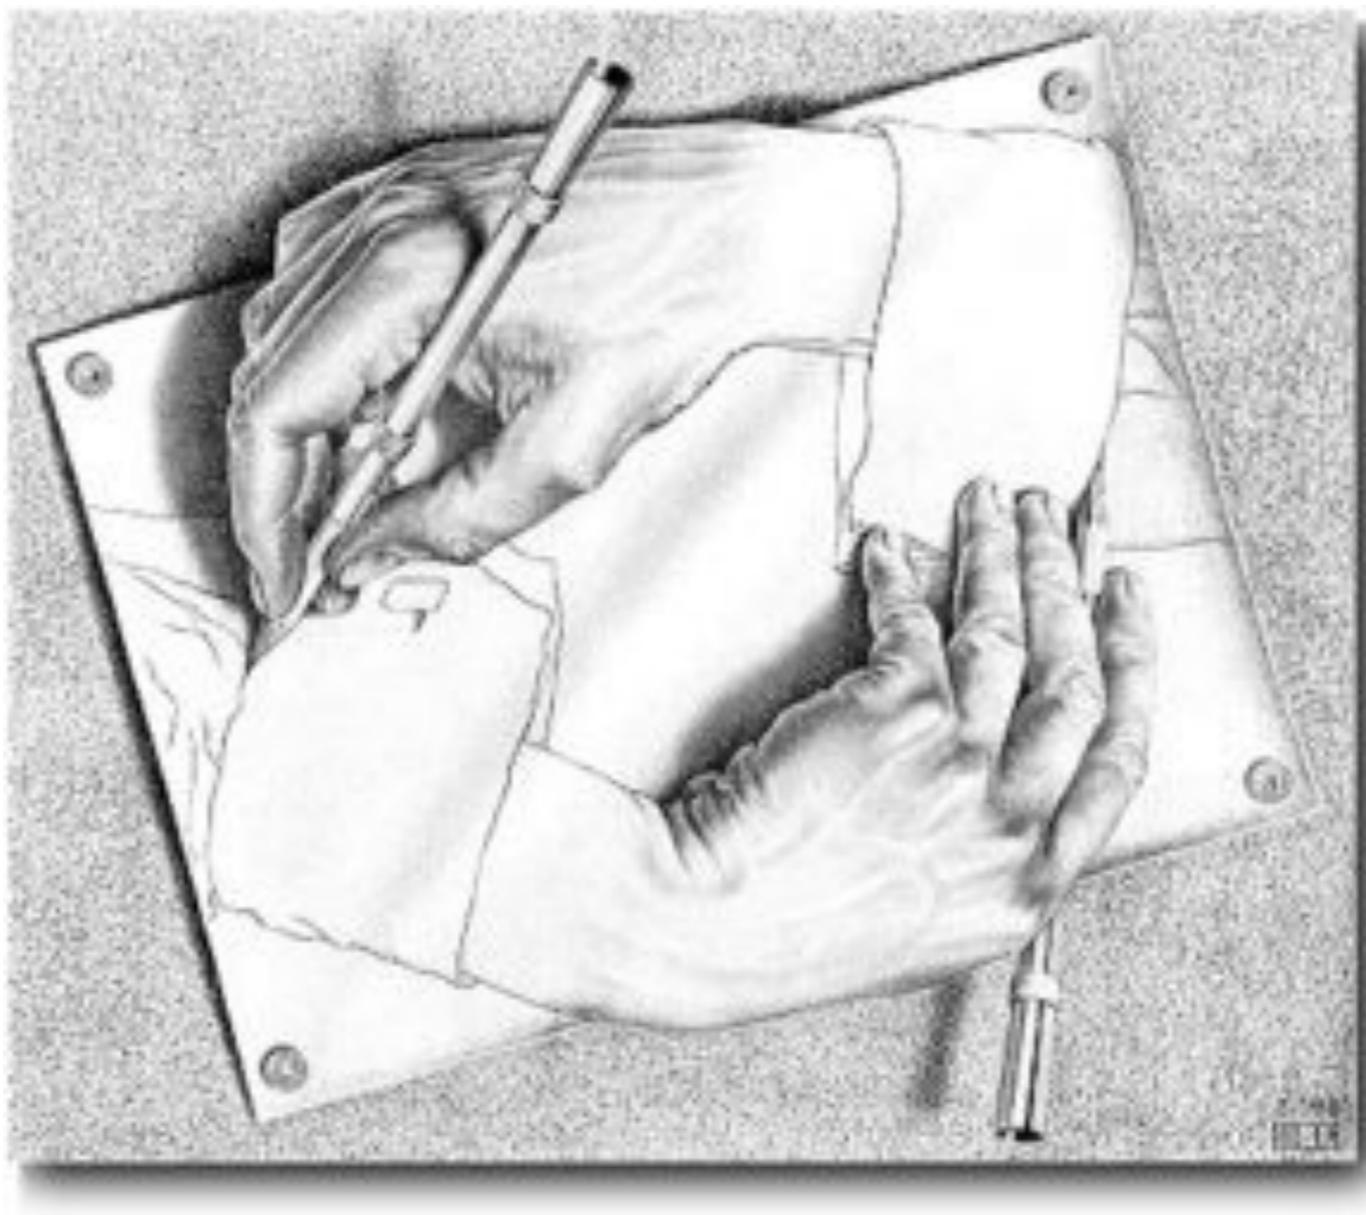


Meta Learning



Meta Learning

Learn to learn



Meta Learning

Learning to Learn without Gradient Descent by Gradient Descent

2. Learning Black-box Optimization

A black-box optimization algorithm can be summarized by the following loop:

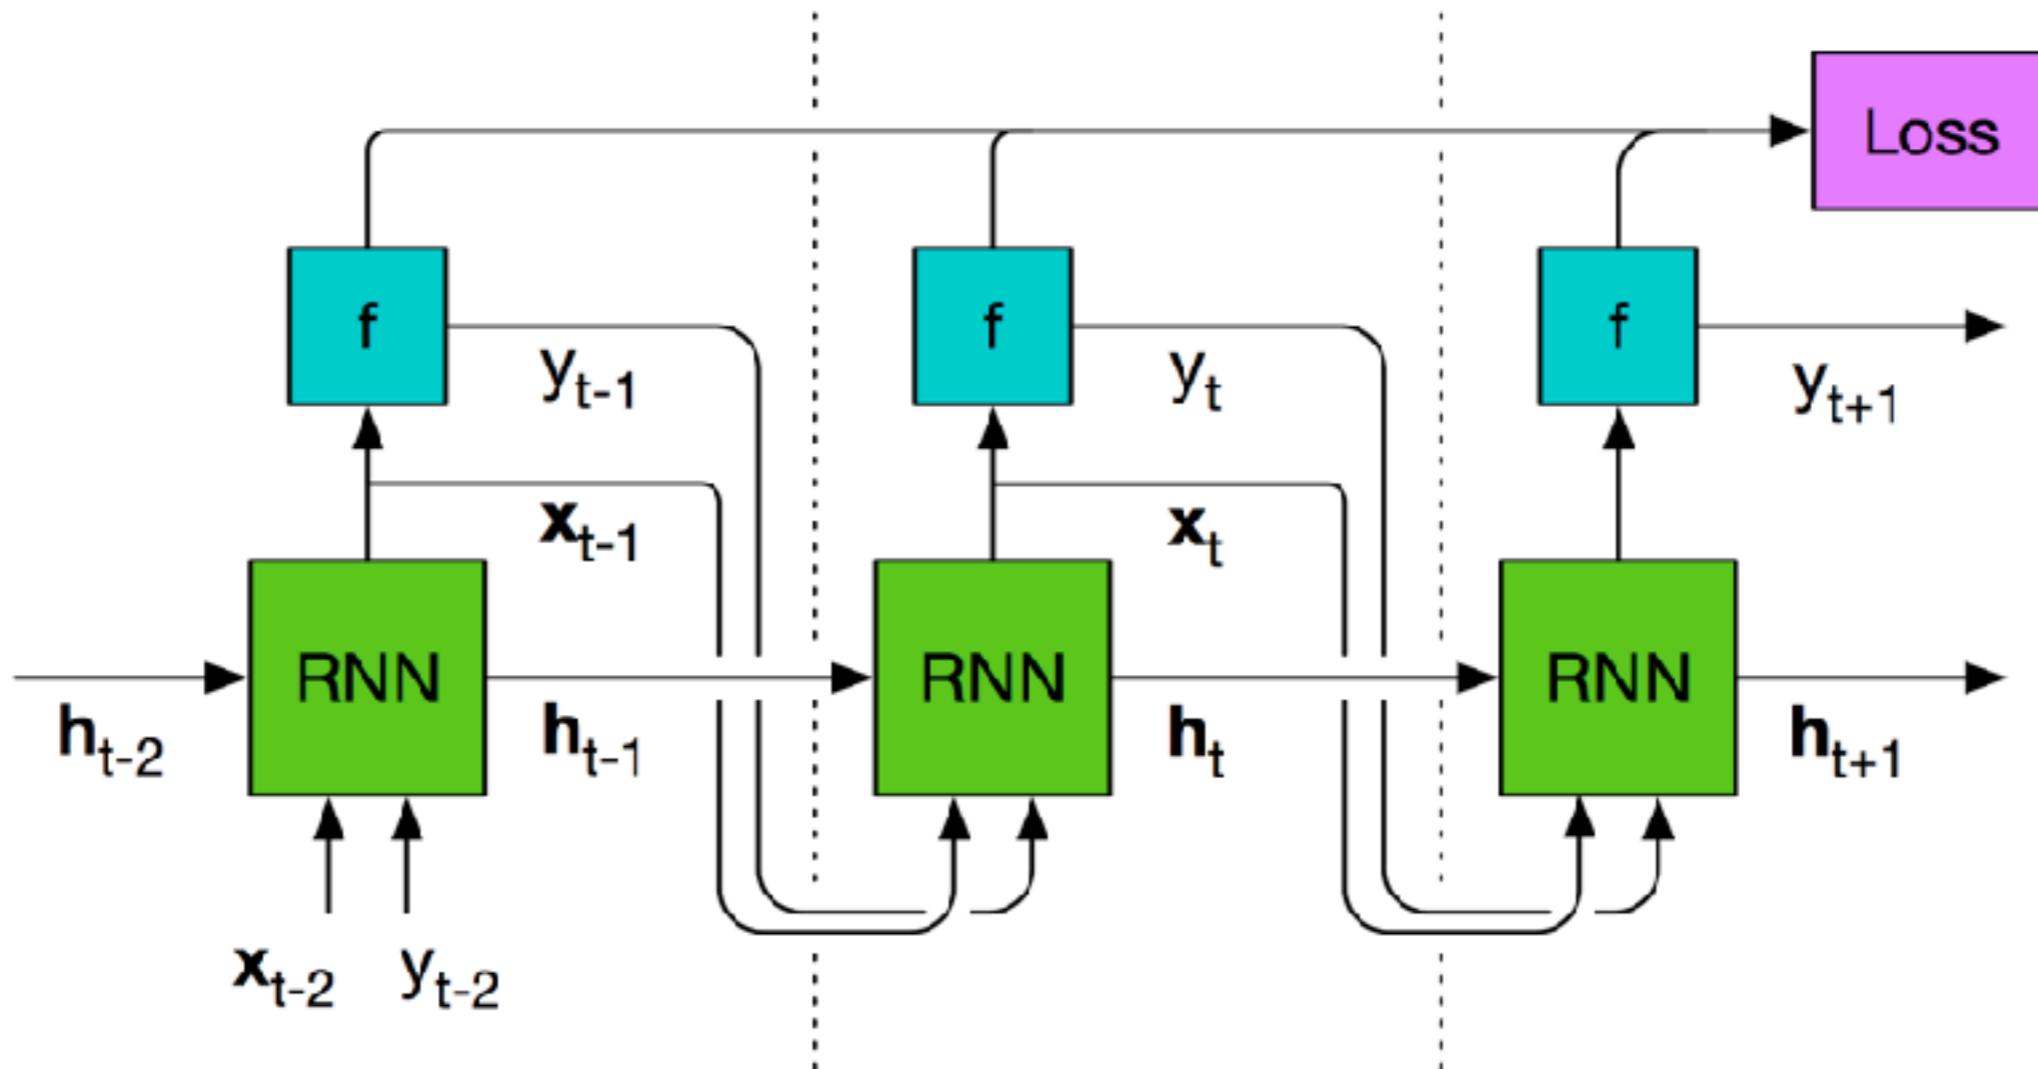
1. Given the current state of knowledge \mathbf{h}_t propose a query point \mathbf{x}_t
2. Observe the response y_t
3. Update any internal statistics to produce \mathbf{h}_{t+1}

<https://arxiv.org/abs/1611.03824>



Meta Learning

Learning to Learn without Gradient Descent by Gradient Descent

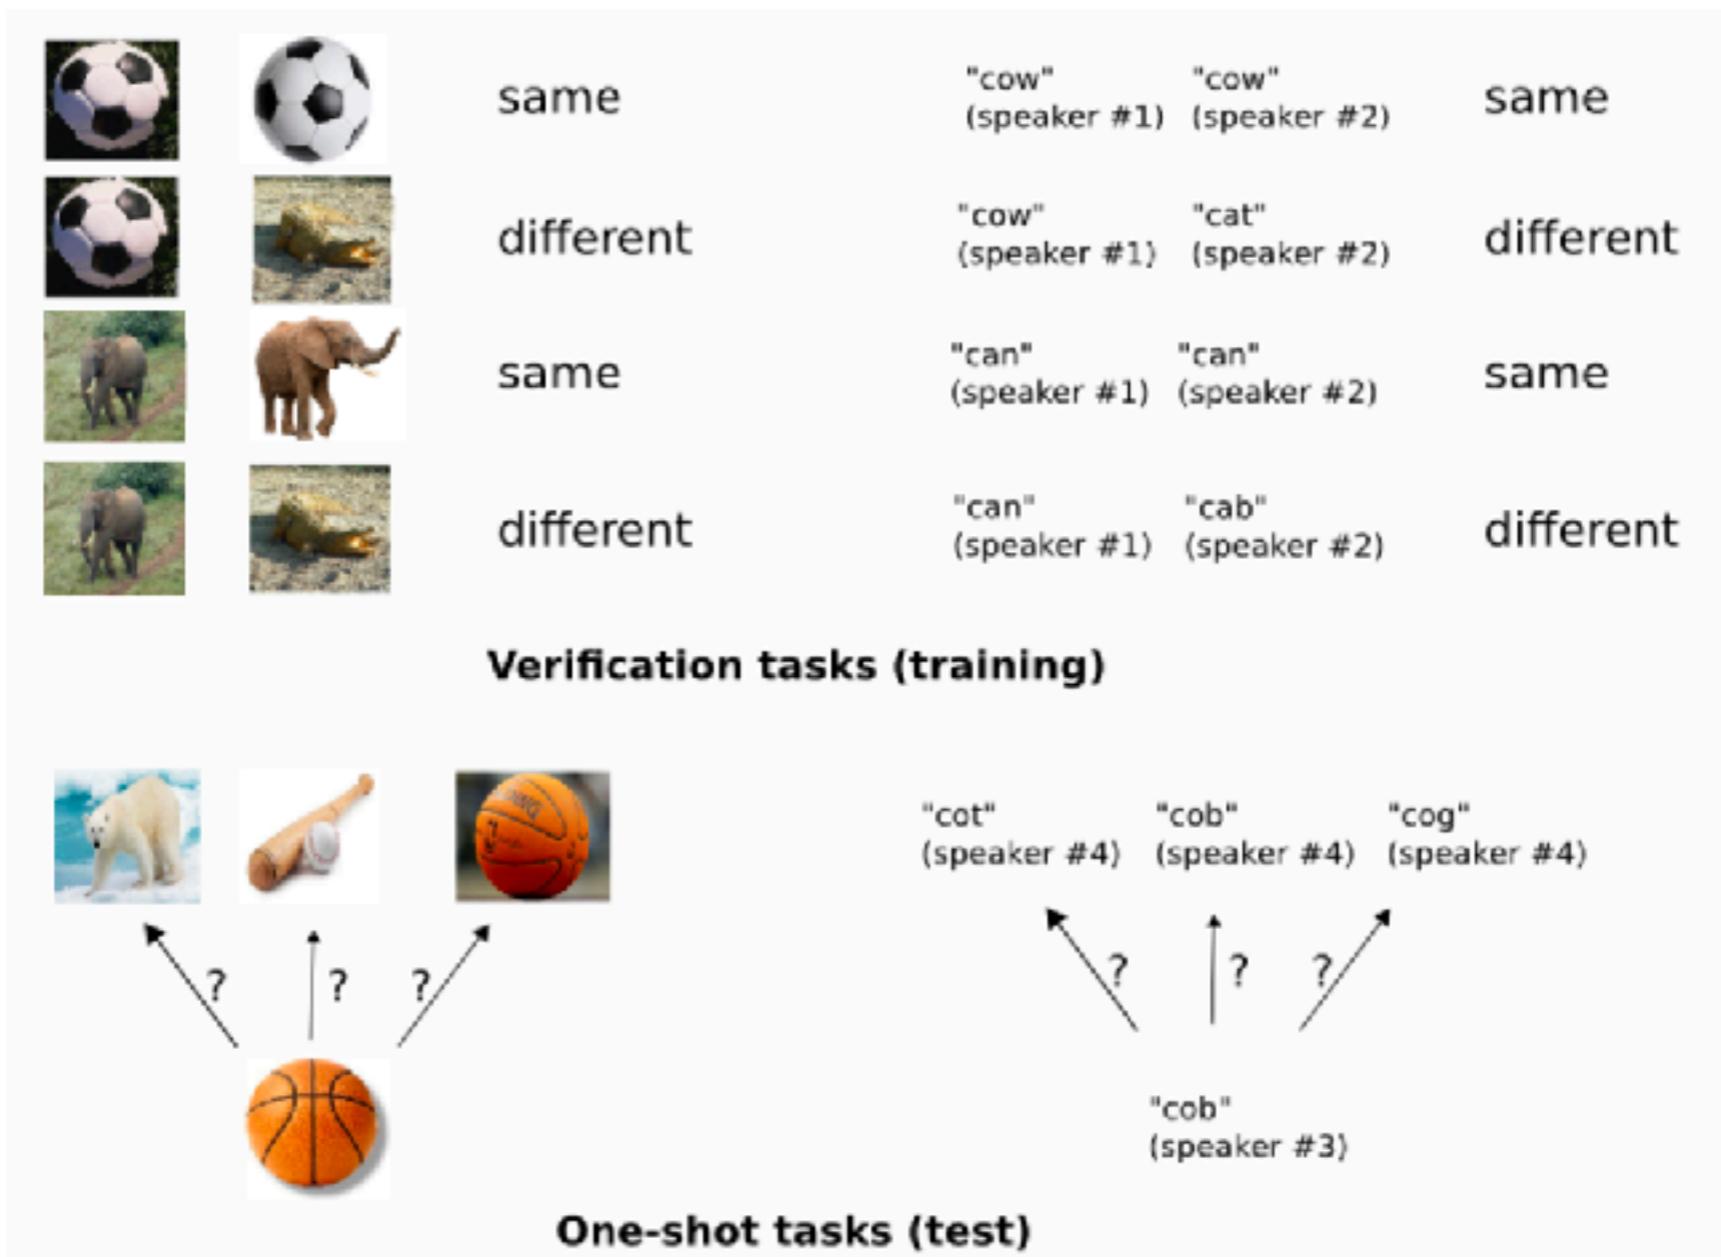


<https://arxiv.org/abs/1611.03824>



Meta Learning

Siamese Neural Networks



<https://www.cs.cmu.edu/~rsalakhu/papers/oneshot1.pdf>



Meta Learning

Siamese Neural Networks

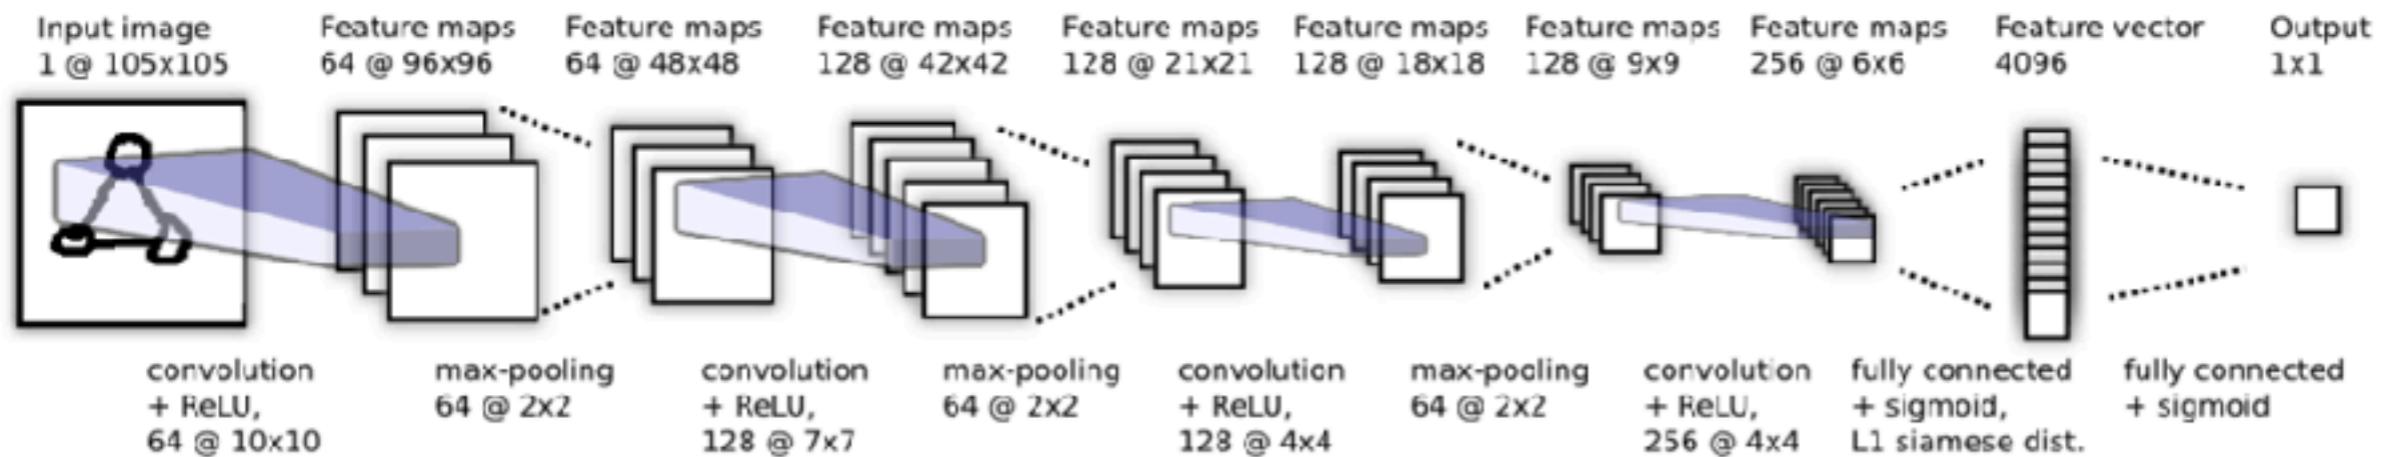


Figure 4. Best convolutional architecture selected for verification task. Siamese twin is not depicted, but joins immediately after the 4096 unit fully-connected layer where the L1 component-wise distance between vectors is computed.

<https://www.cs.cmu.edu/~rsalakhu/papers/oneshot1.pdf>



Meta Learning

Triplet Networks

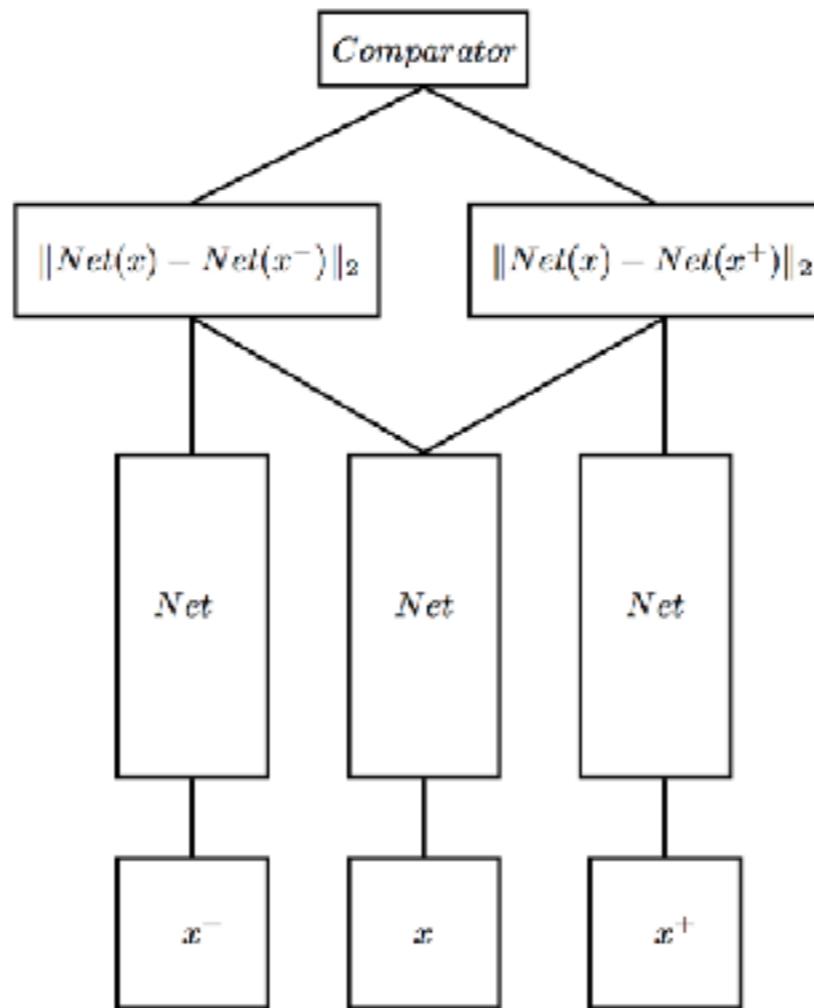


Figure 1: Triplet network structure

<https://arxiv.org/abs/1412.6622>



Meta Learning

Matching Network

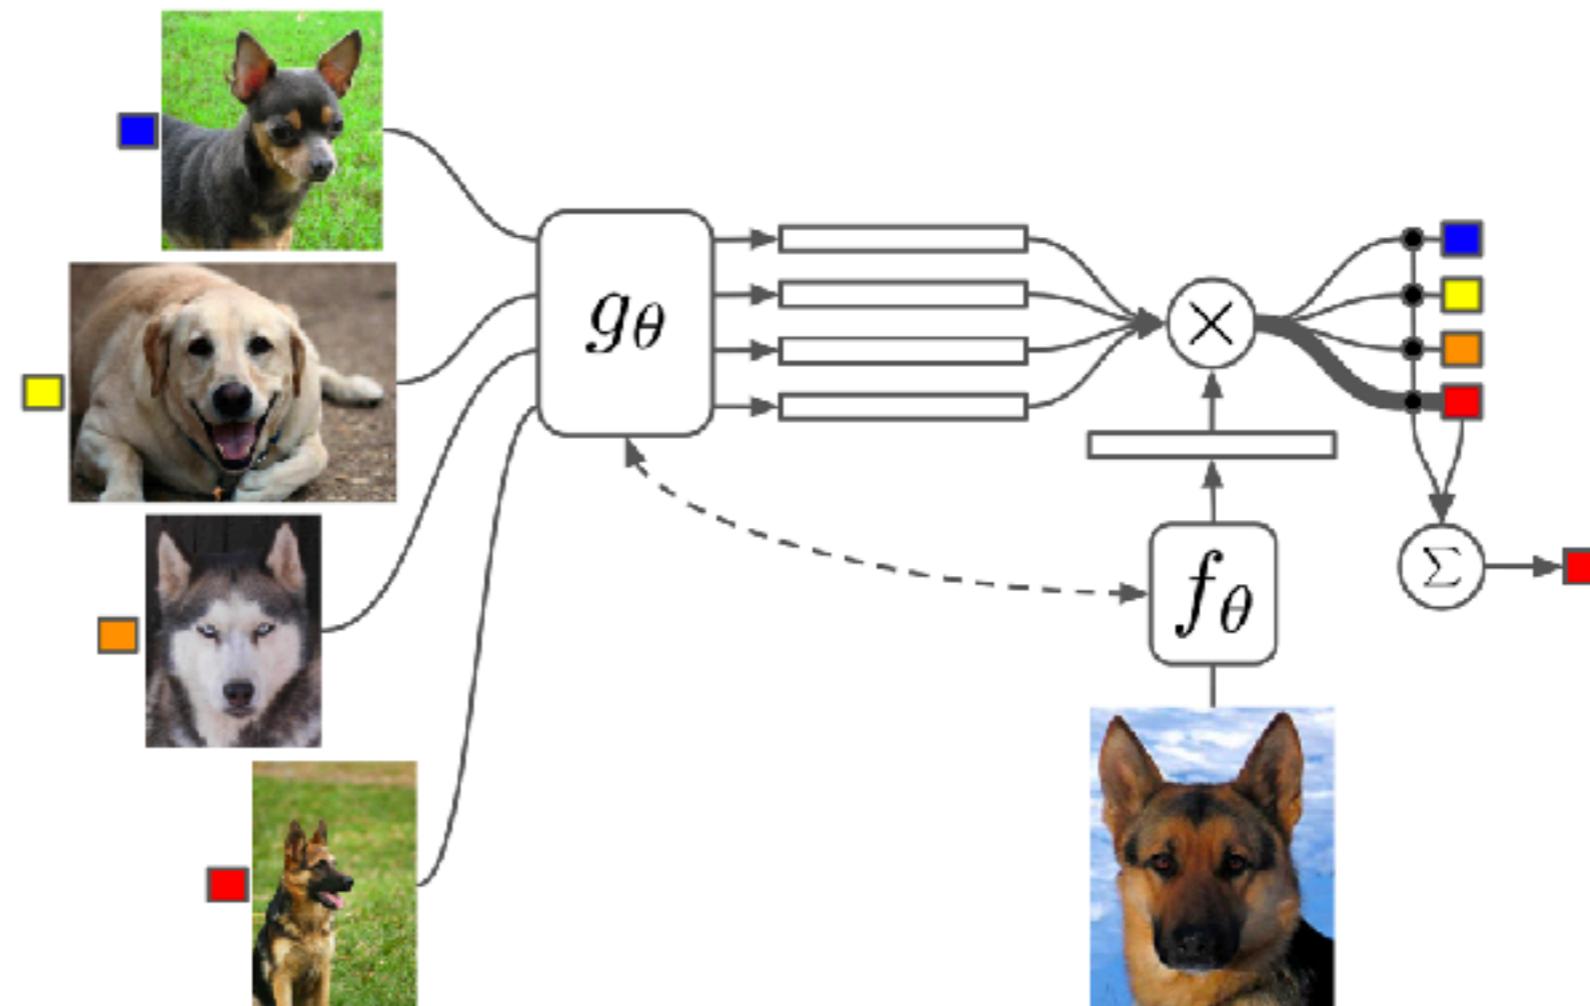


Figure 1: Matching Networks architecture

<https://arxiv.org/abs/1606.04080>

Meta Learning

Matching Network

$$\hat{y} = \sum_{i=1}^k a(\hat{x}, x_i) y_i$$

<https://arxiv.org/abs/1606.04080>



Meta Learning

Prototypical Networks

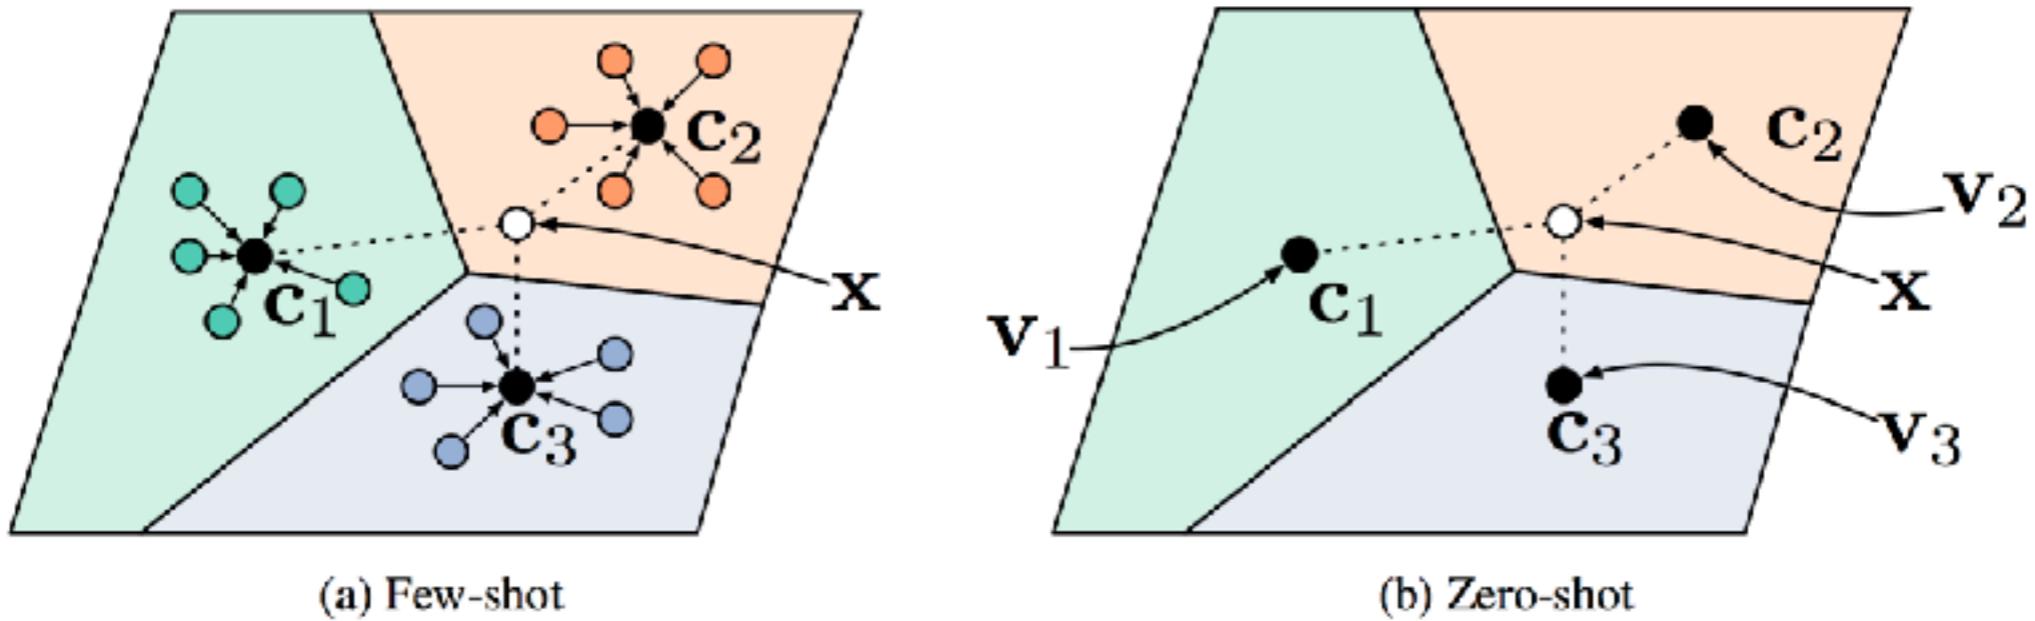


Figure 1: Prototypical networks in the few-shot and zero-shot scenarios. **Left:** Few-shot prototypes \mathbf{c}_k are computed as the mean of embedded support examples for each class. **Right:** Zero-shot prototypes \mathbf{c}_k are produced by embedding class meta-data \mathbf{v}_k . In either case, embedded query points are classified via a softmax over distances to class prototypes: $p_\phi(y = k|\mathbf{x}) \propto \exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_k))$.

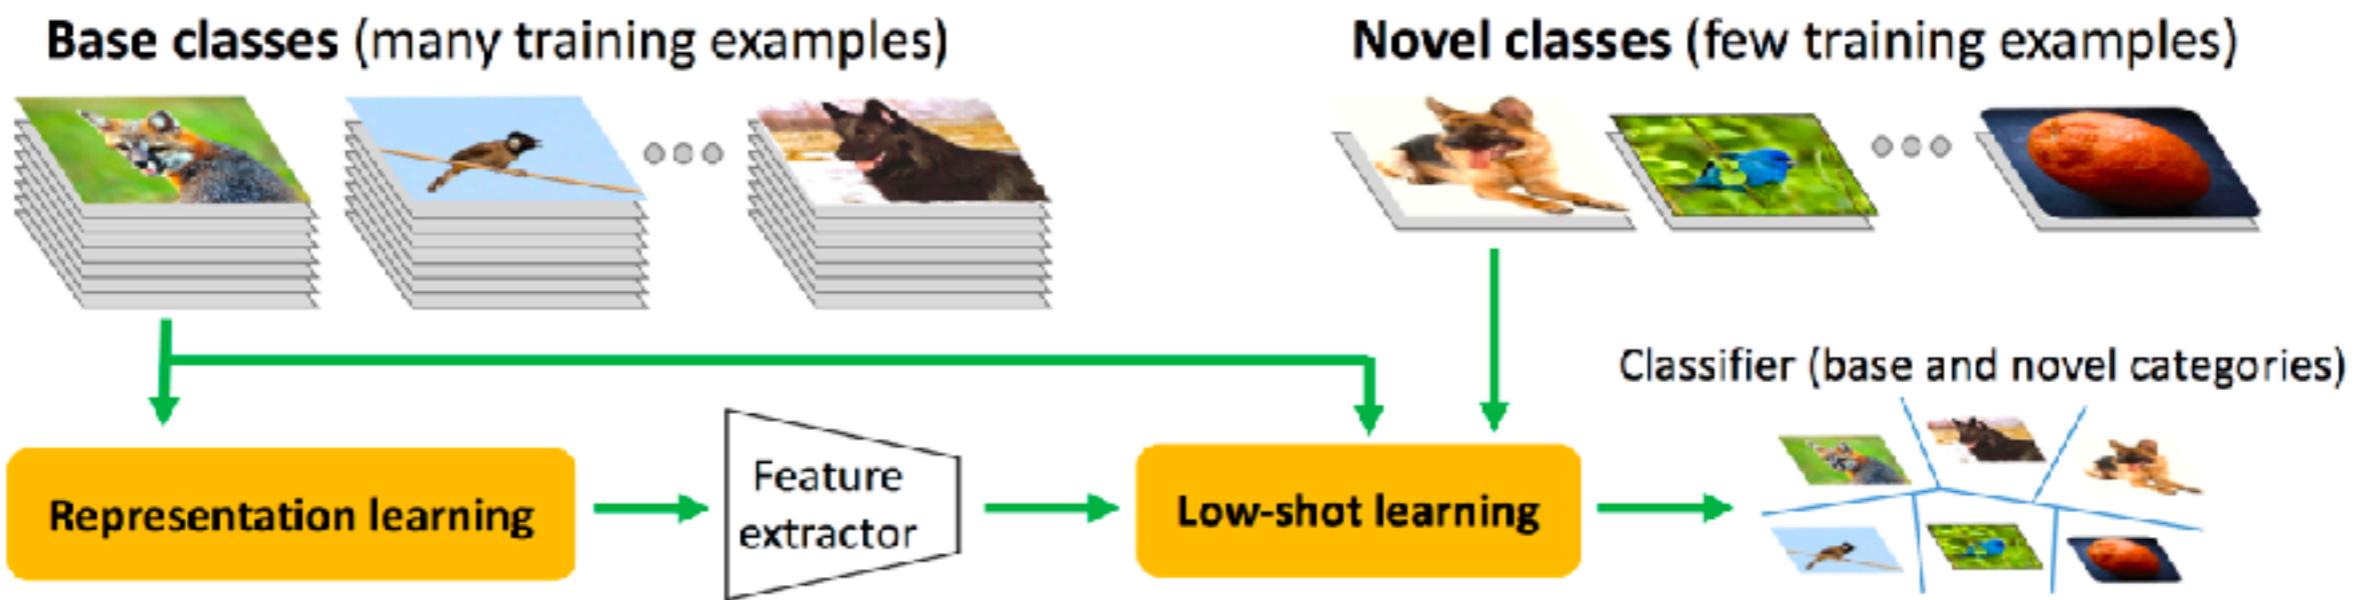
Prototypical networks learn **a metric space** in which classification can be performed by computing distances to prototype representations of each class.

<https://arxiv.org/abs/1703.05175>



Meta Learning

Low-shot learning



Low-shot visual learning—the ability to recognize novel object categories from very few examples—is a hallmark of human visual intelligence.

<https://arxiv.org/abs/1606.02819>



Meta Learning

Low-shot learning

Learning to **generate** new examples



<https://arxiv.org/abs/1606.02819>

Meta Learning

Design a CNN with RL

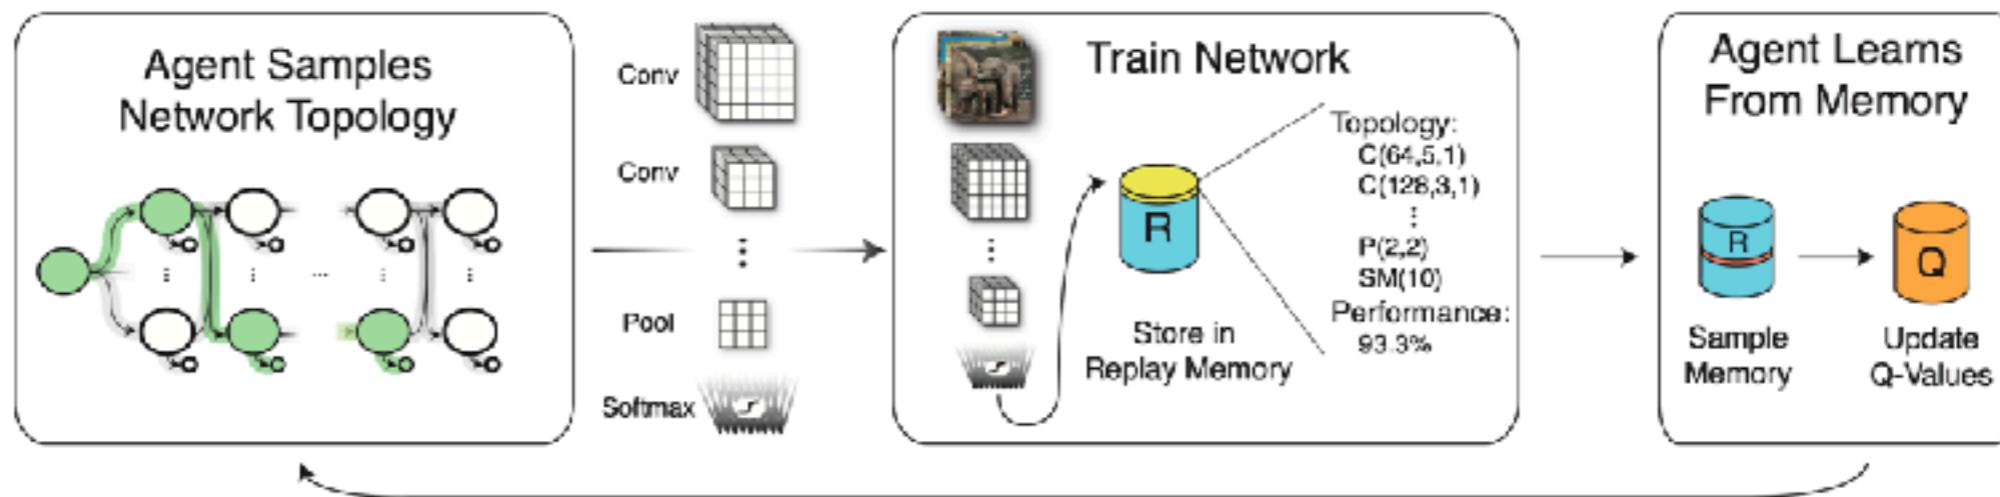


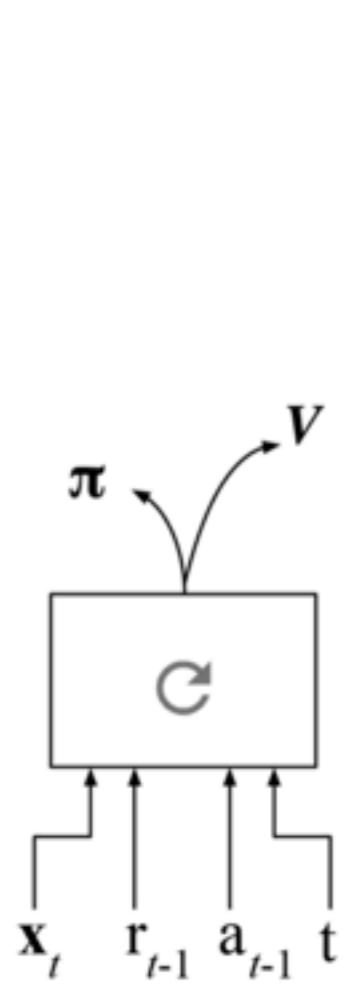
Figure 1: **Designing CNN Architectures with Q -learning:** The agent begins by sampling a Convolutional Neural Network (CNN) topology conditioned on a predefined behavior distribution and the agent's prior experience (left block). That CNN topology is then trained on a specific task; the topology description and performance, e.g. validation accuracy, are then stored in the agent's memory (middle block). Finally, the agent uses its memories to learn about the space of CNN topologies through Q -learning (right block).

<https://arxiv.org/pdf/1611.02167>

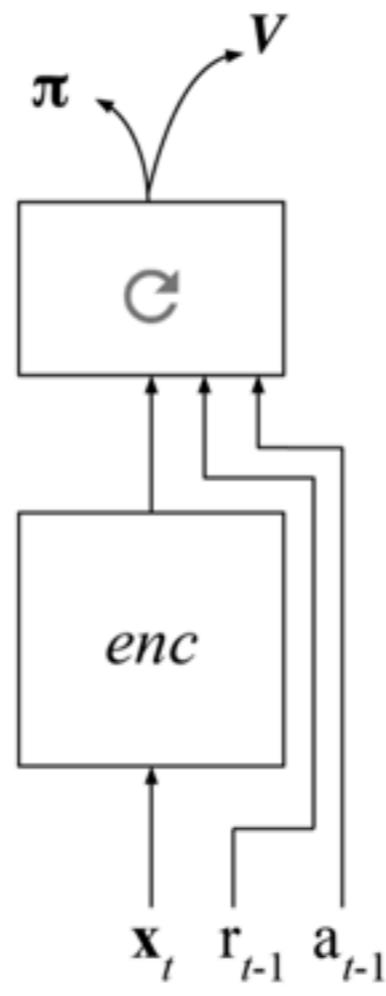


Meta Learning

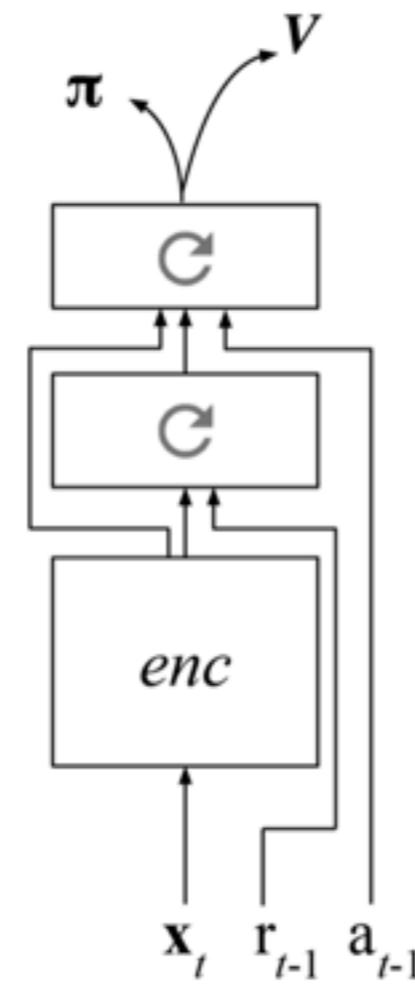
Deep meta RL



(a) LSTM A2C



(b) LSTM A3C



(c) Stacked-LSTM A3C

<https://arxiv.org/abs/1611.05763>



Meta Learning

Deep meta RL

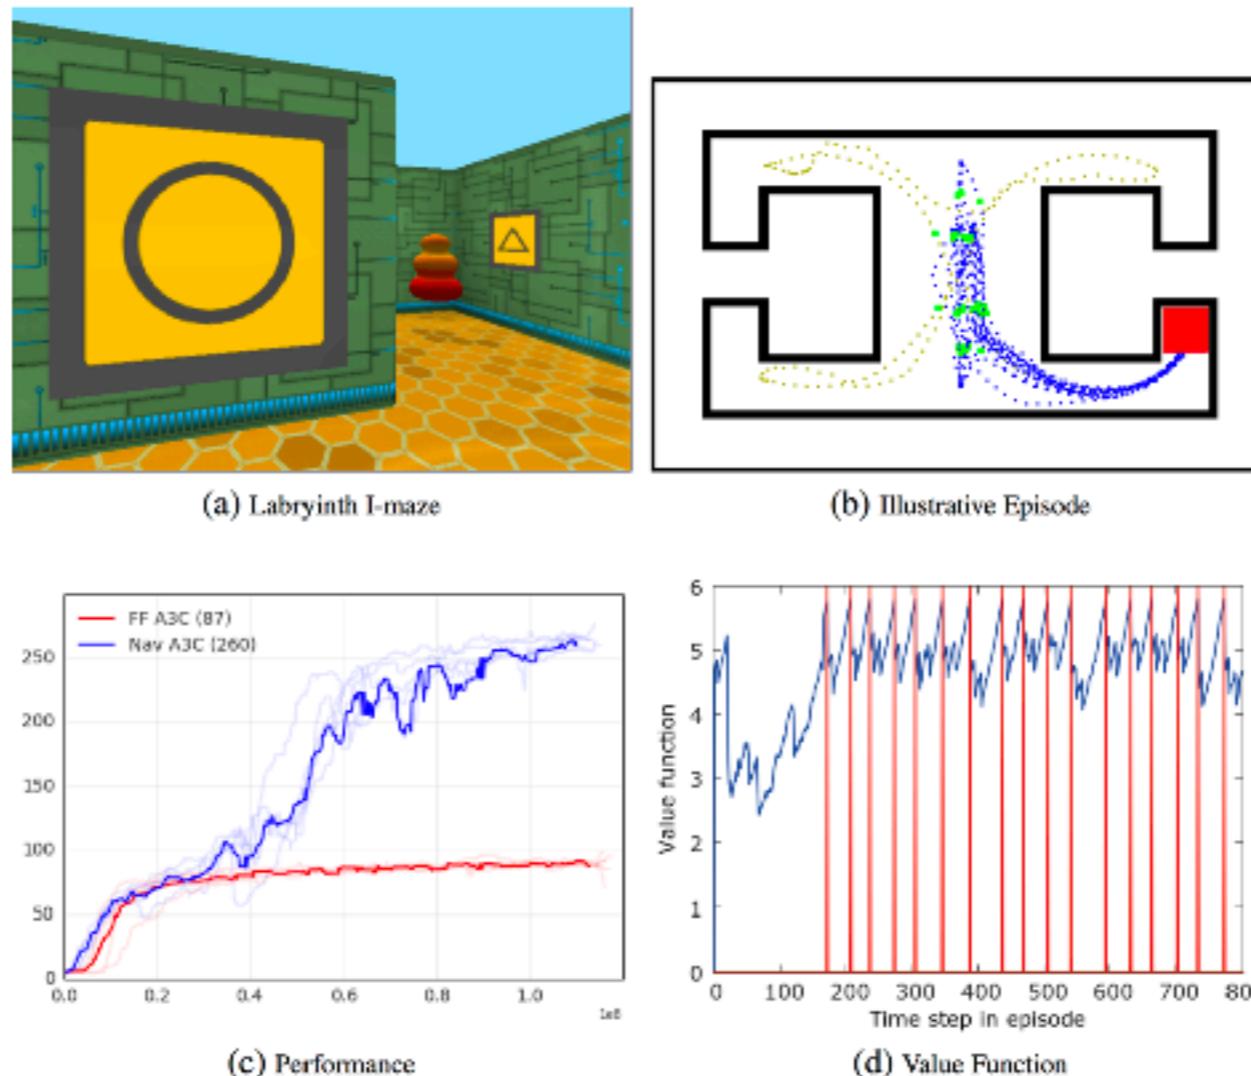


Figure 7: a) view of I-maze showing goal object in one of the 4 alcoves b) following initial exploration (light trajectories), agent repeatedly goes to goal (blue trajectories) c) Performance of stacked LSTM (termed “Nav A3C”) and feedforward (“FF A3C”) architectures, per episode (goal = 10 points) averaged across top 5 hyperparameters. e) following initial goal discovery (goal hits marked in red), value function occurs well in advance of the agent seeing the goal which is hidden in an alcove. Figure used with permission from [Mirowski et al. \(2016\)](#).

<https://arxiv.org/abs/1611.05763>

Meta Learning

Model-Agnostic Meta-Learning

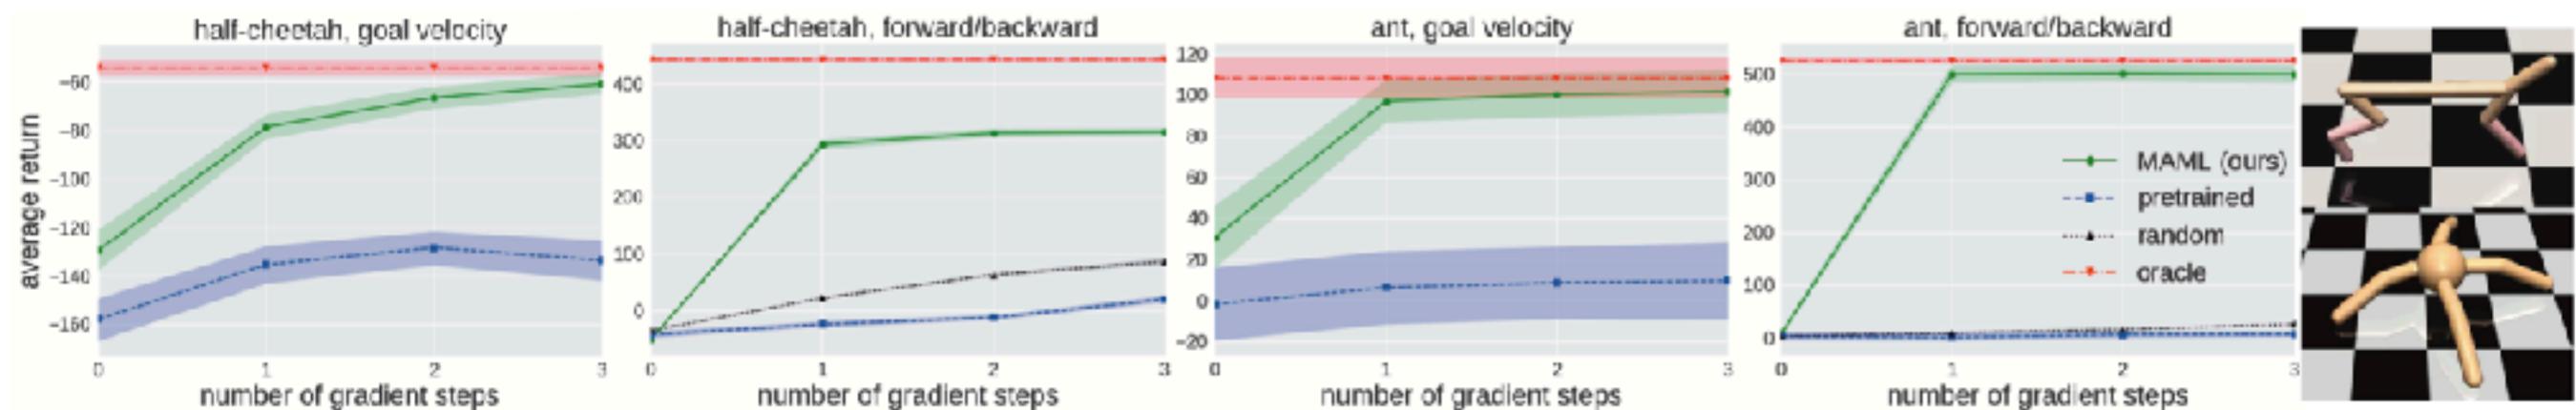


Figure 5. Reinforcement learning results for the half-cheetah and ant locomotion tasks, with the tasks shown on the far right. Each gradient step requires additional samples from the environment, unlike the supervised learning tasks. The results show that MAML can adapt to new goal velocities and directions substantially faster than conventional pretraining or random initialization, achieving good performances in just two or three gradient steps. We exclude the goal velocity, random baseline curves, since the returns are much worse (< -200 for cheetah and < -25 for ant).

<https://arxiv.org/abs/1703.03400>



Meta Learning

Model-Agnostic Meta-Learning

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
 - 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 7: **end for**
 - 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
 - 9: **end while**
-

<https://arxiv.org/abs/1703.03400>



Meta Learning

Model-Agnostic Meta-Learning

Algorithm 3 MAML for Reinforcement Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

```
1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Sample  $K$  trajectories  $\mathcal{D} = \{(\mathbf{x}_1, \mathbf{a}_1, \dots, \mathbf{x}_H)\}$  using  $f_\theta$ 
        in  $\mathcal{T}_i$ 
6:     Evaluate  $\nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)$  using  $\mathcal{D}$  and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation 4
7:     Compute adapted parameters with gradient descent:
         $\theta'_i = \theta - \alpha \nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)$ 
8:     Sample trajectories  $\mathcal{D}'_i = \{(\mathbf{x}_1, \mathbf{a}_1, \dots, \mathbf{x}_H)\}$  using  $f_{\theta'_i}$ 
        in  $\mathcal{T}_i$ 
9:   end for
10:  Update  $\theta \leftarrow \theta - \beta \nabla_\theta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$  using each  $\mathcal{D}'_i$ 
      and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation 4
11: end while
```

In reinforcement learning (RL), the goal of few-shot meta-learning is to enable an agent to **quickly acquire a policy for a new test task** using only a small amount of experience in the test setting

<https://arxiv.org/abs/1703.03400>



Uncertainty in Deep Learning

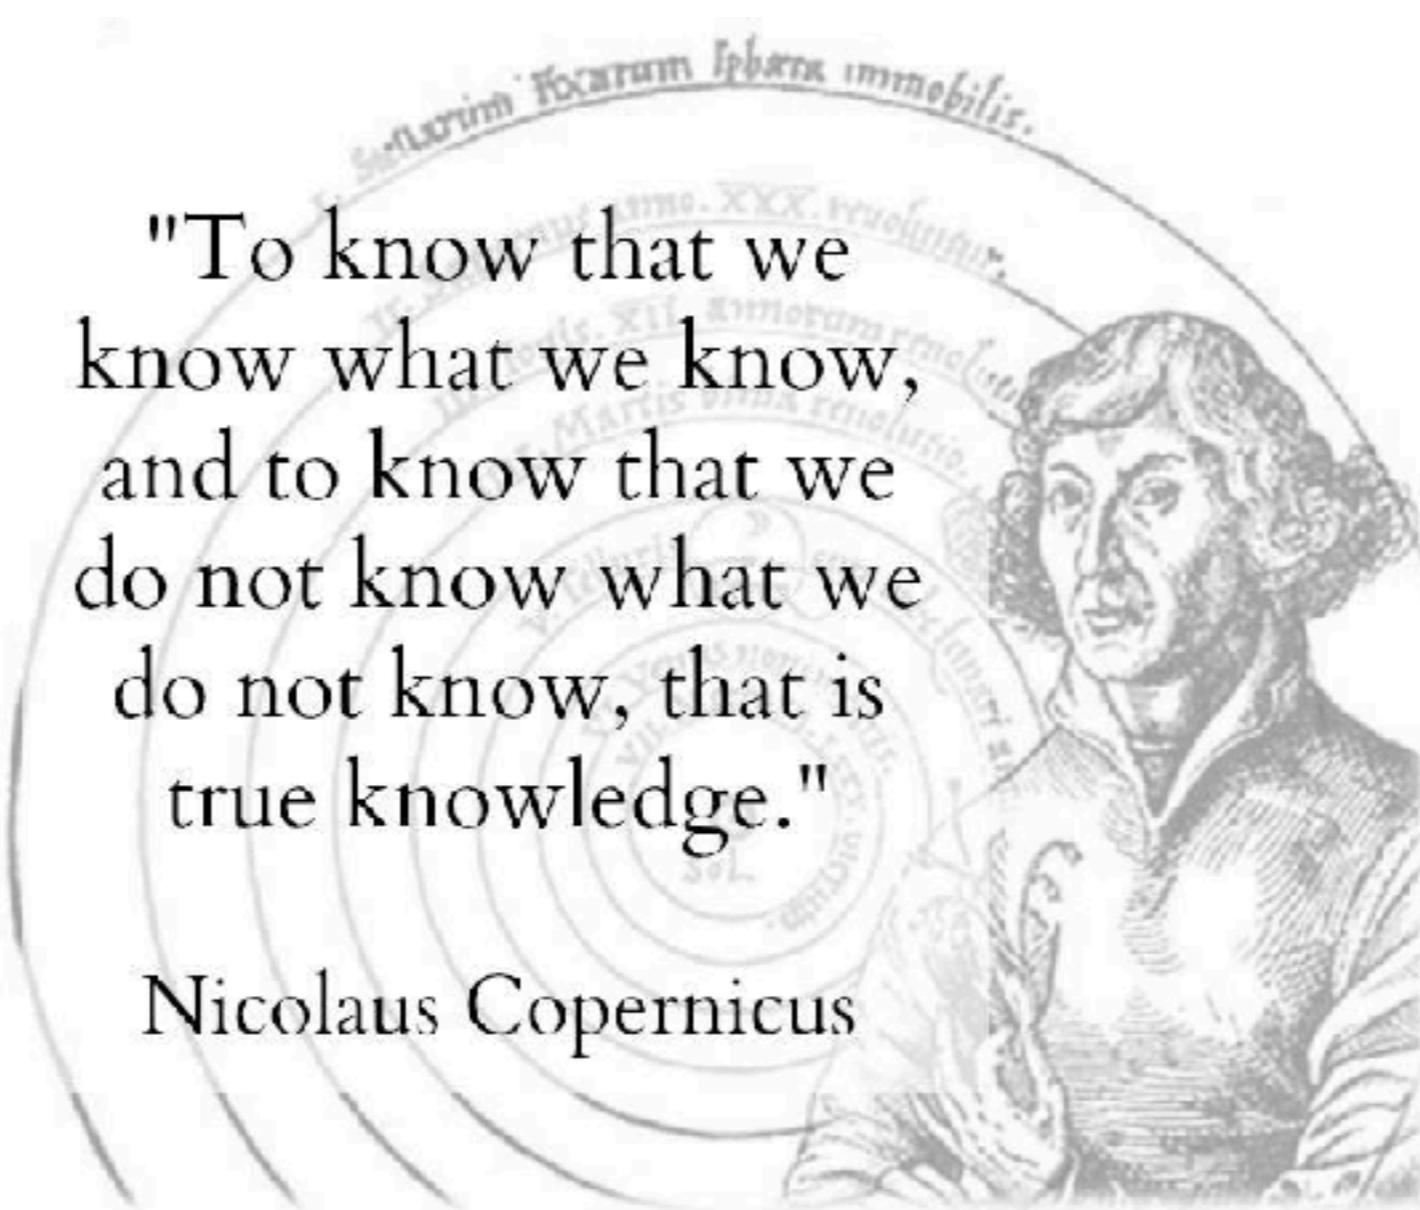


Uncertainty in Deep Learning

What we do not know

"To know that we
know what we know,
and to know that we
do not know what we
do not know, that is
true knowledge."

Nicolaus Copernicus



Uncertainty in Deep Learning

What we do not know



Uncertainty in Deep Learning

What we do not know



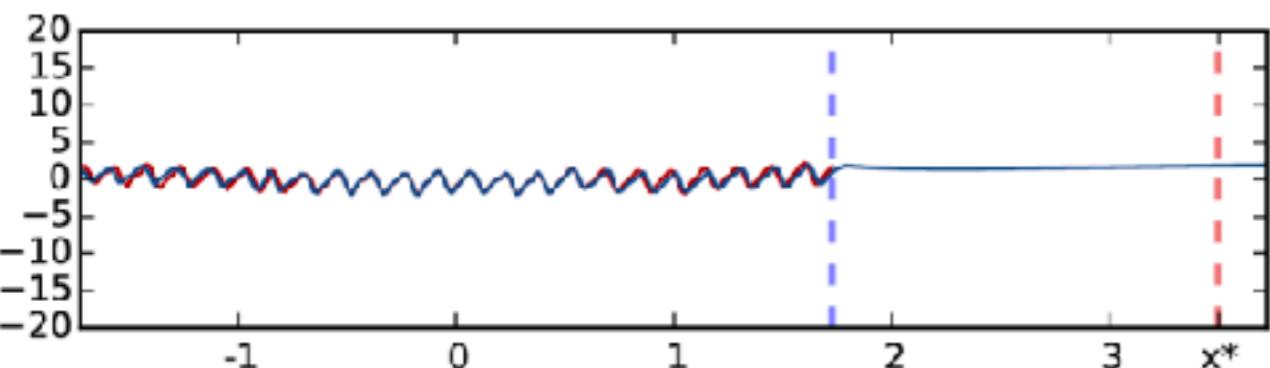
Uncertainty in Deep Learning

What we do not know

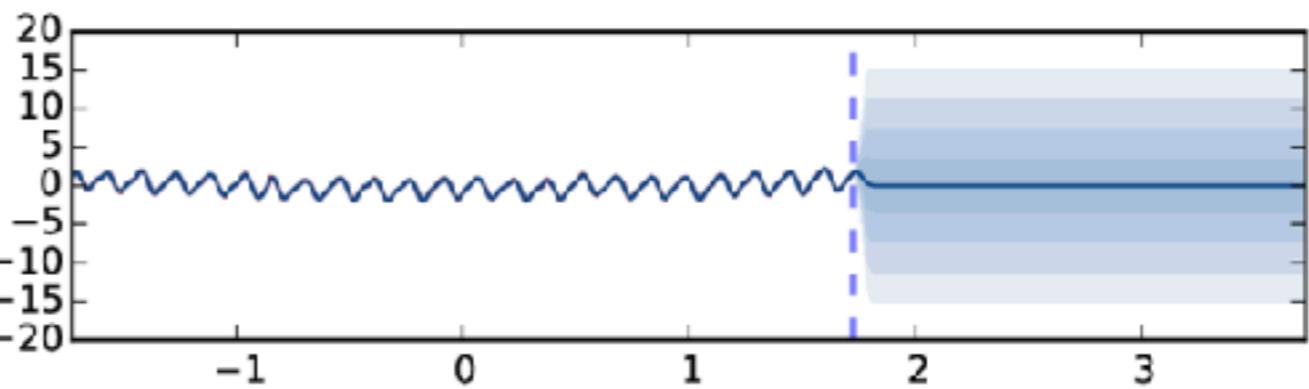


Uncertainty in Deep Learning

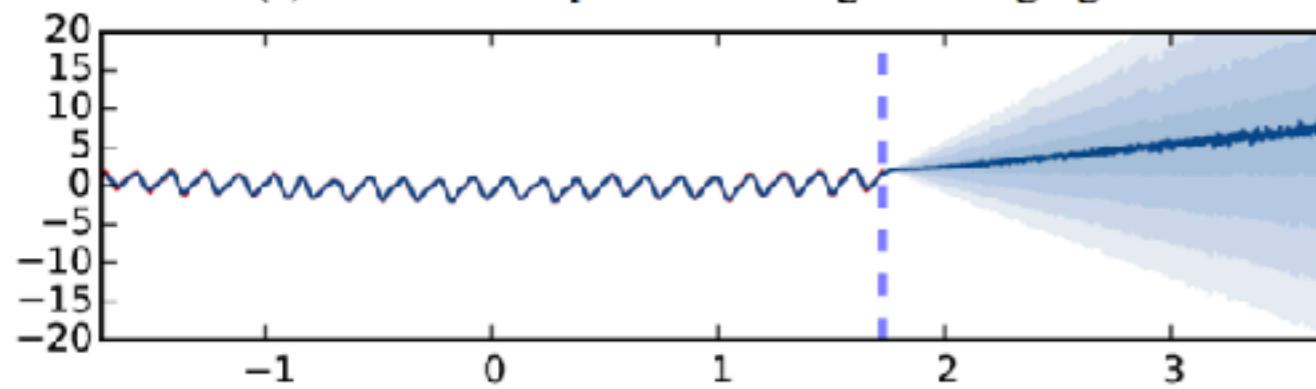
Dropout as Bayes Approximation



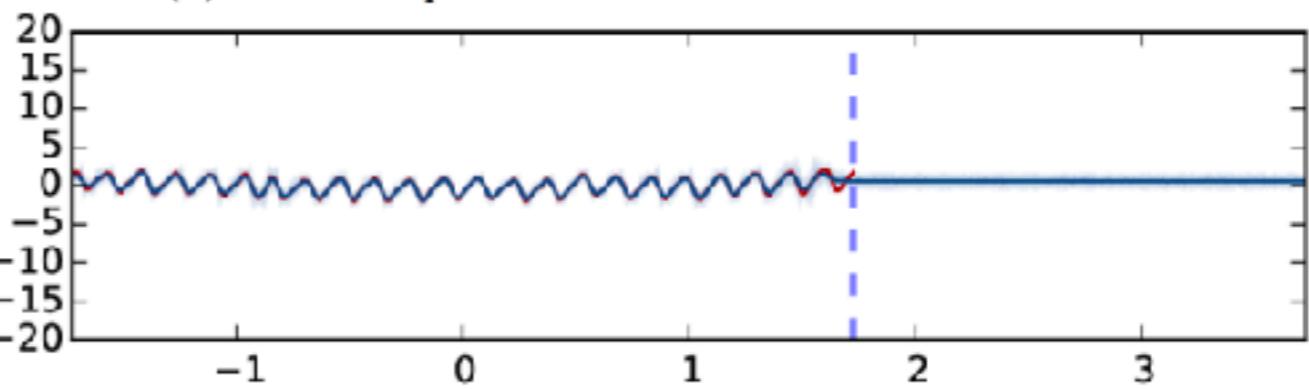
(a) Standard dropout with weight averaging



(b) Gaussian process with SE covariance function



(c) MC dropout with ReLU non-linearities



(d) MC dropout with TanH non-linearities

<https://arxiv.org/abs/1506.02142>



Uncertainty in Deep Learning

Dropout as Bayes Approximation

$$\begin{aligned} \text{Var}_{q(\mathbf{y}^*|\mathbf{x}^*)}(\mathbf{y}^*) &\approx \tau^{-1} \mathbf{I}_D \\ &+ \frac{1}{T} \sum_{t=1}^T \hat{\mathbf{y}}^*(\mathbf{x}^*, \mathbf{W}_1^t, \dots, \mathbf{W}_L^t)^T \hat{\mathbf{y}}^*(\mathbf{x}^*, \mathbf{W}_1^t, \dots, \mathbf{W}_L^t) \\ &- \mathbb{E}_{q(\mathbf{y}^*|\mathbf{x}^*)}(\mathbf{y}^*)^T \mathbb{E}_{q(\mathbf{y}^*|\mathbf{x}^*)}(\mathbf{y}^*) \end{aligned}$$

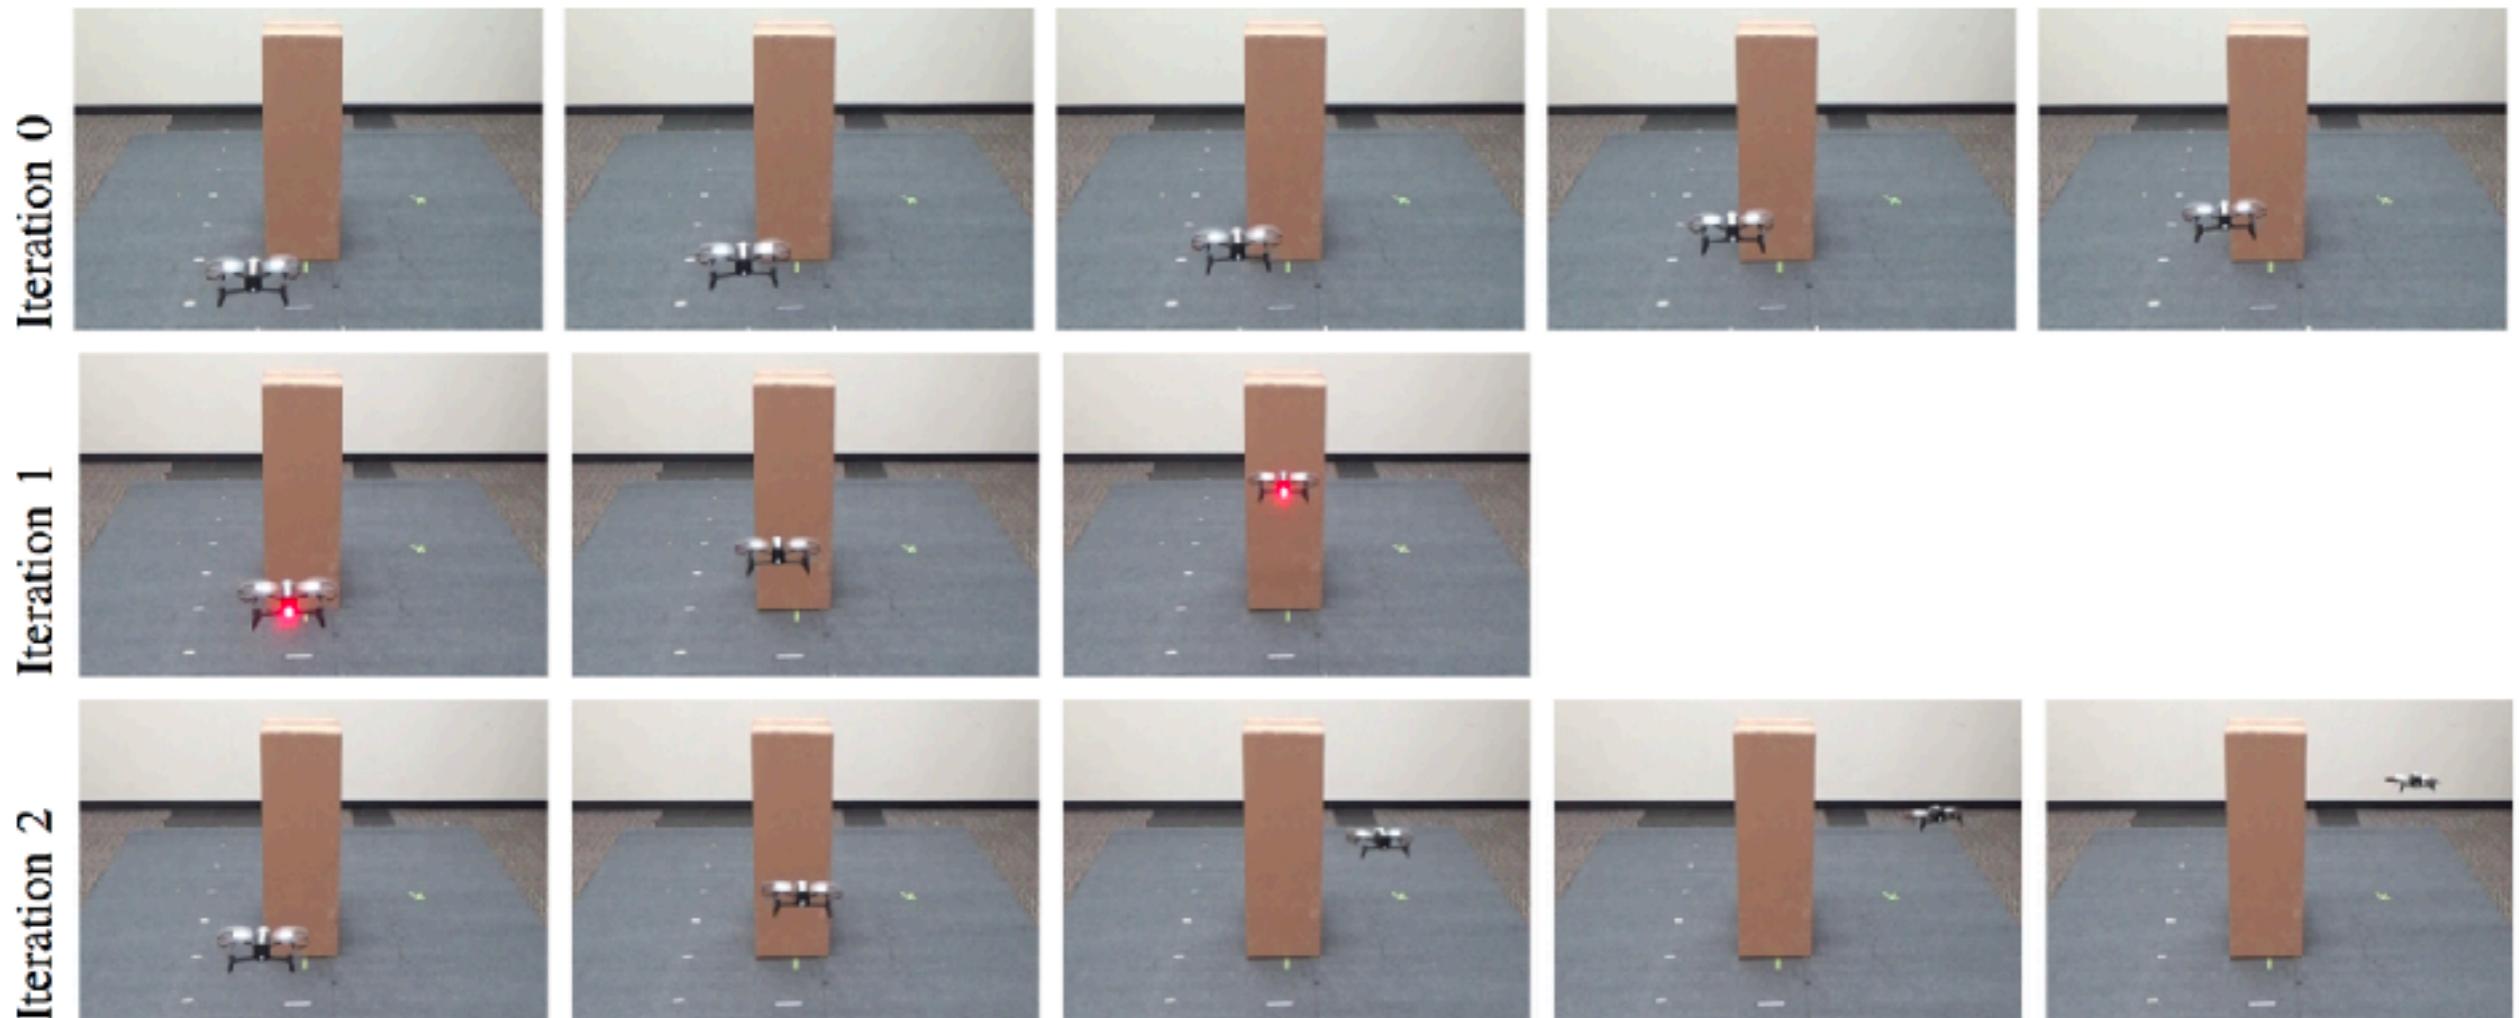
which equals the sample variance of T stochastic forward passes through the NN plus the inverse model precision.

<https://arxiv.org/abs/1506.02142>



Uncertainty in Deep Learning

Uncertainty-Aware RL



<https://arxiv.org/abs/1702.01182>

Uncertainty in Deep Learning

Uncertainty-Aware RL

Algorithm 1 Neural net training with bootstrapping and dropout

- 1: **input:** dataset $\mathcal{D} = \{\mathbf{x}_t^{(i)}, \mathbf{u}_{t:t+H}^{(i)}, \mathbf{o}_t^{(i)}\}$, neural network model NN
- 2: **for** $b = 1$ to B **do**
- 3: Sample a dataset of subsequences $\mathcal{D}^{(b)}$ from the full dataset \mathcal{D} with replacement
- 4: Initialize neural network $\text{NN}^{(b)}$ with random weights
- 5: **for** number of SGD iterations **do**
- 6: Sample datapoint $(\mathbf{x}_t, \mathbf{u}_{t:t+H}, \mathbf{o}_t)$ from $\mathcal{D}^{(b)}$
- 7: Sample $\text{NN}_d^{(b)}$ by masking the units in $\text{NN}^{(b)}$ using dropout
- 8: Run forward pass on $\text{NN}_d^{(b)}$ using $(\mathbf{x}_t, \mathbf{u}_{t:t+H}, \mathbf{o}_t)$
- 9: Run backward pass on $\text{NN}_d^{(b)}$ to get gradient $g_d^{(b)}$
- 10: Update model $\text{NN}^{(b)}$ parameters using $g_d^{(b)}$
- 11: **end for**
- 12: **end for**

<https://arxiv.org/abs/1702.01182>



Uncertainty in Deep Learning

Novelty Detection



(a) LIDAR-mapped hallway environment for training and testing.



(b) Xtion-mapped novel environment used for LIDAR-free testing.

<http://roboticsproceedings.org/rss13/p64.pdf>

Uncertainty in Deep Learning

Novelty Detection



(a) Input.



(b) Output.



(c) Input.



(d) Output.



(a) Input.



(b) Output.



(c) Input.

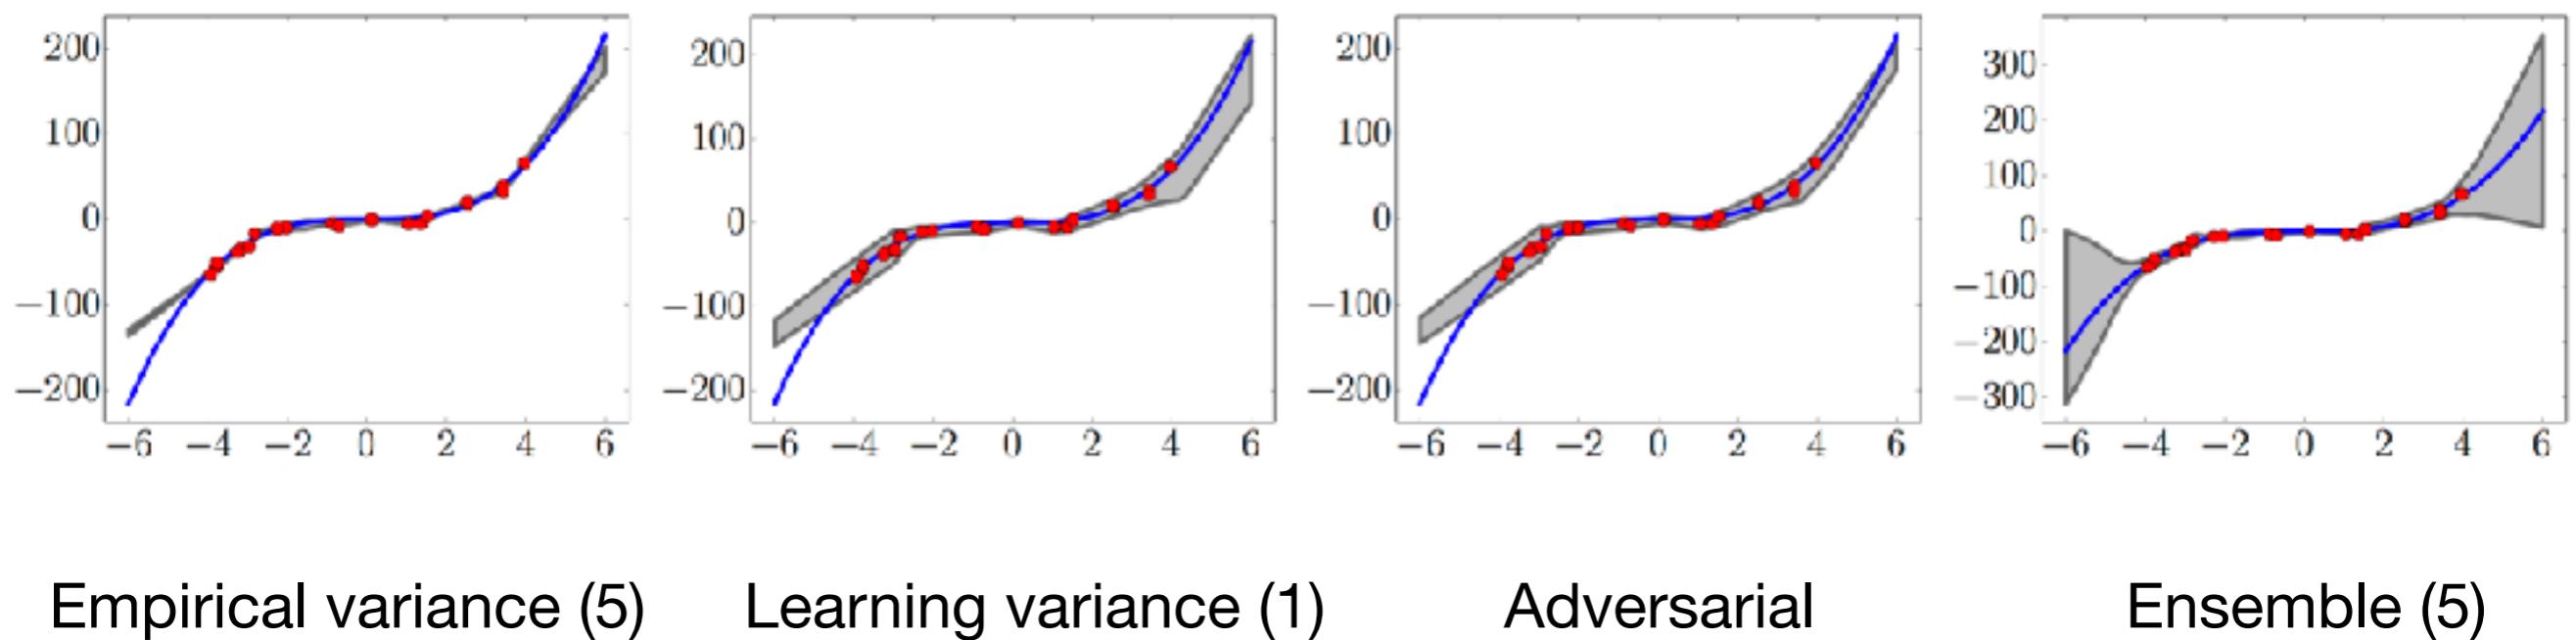


(d) Output.

<http://roboticsproceedings.org/rss13/p64.pdf>

Uncertainty in Deep Learning

Deep Ensemble



Empirical variance (5)

Learning variance (1)

Adversarial

Ensemble (5)

<https://arxiv.org/abs/1612.01474>



Uncertainty in Deep Learning

Deep Ensemble

Algorithm 1 Pseudocode of the training procedure for our method

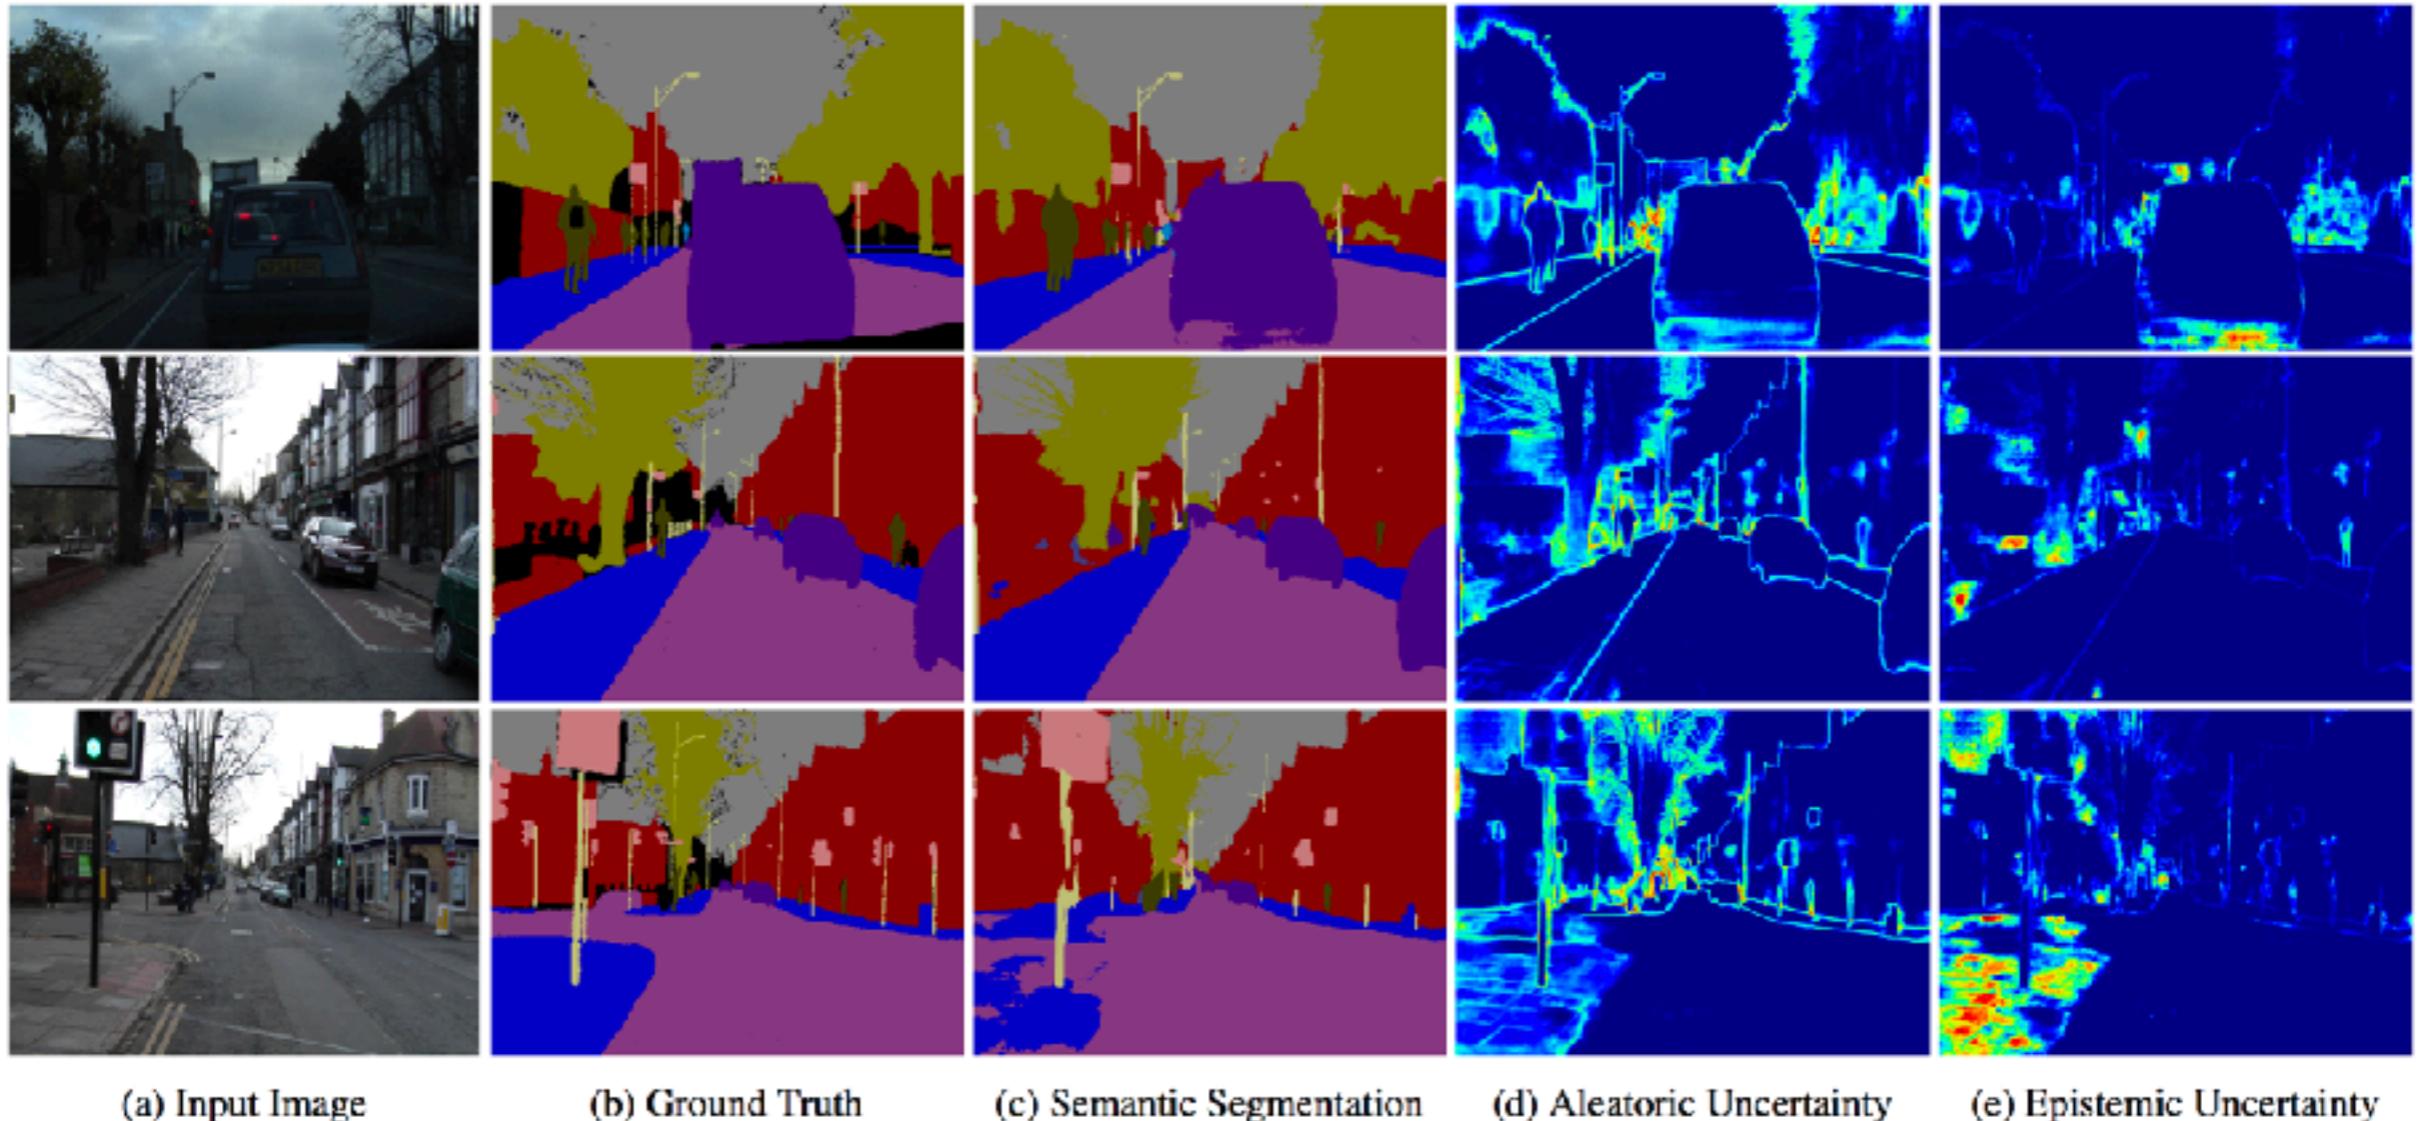
- 1: Initialise $\theta_1, \theta_2, \dots, \theta_M$ randomly
 - 2: **for** $m = 1 : M$ **do** \triangleright *train networks independently in parallel*
 - 3: Sample data point n_m randomly for each net \triangleright *single n_m for clarity, minibatch in practice*
 - 4: Generate adversarial example using $\mathbf{x}'_{n_m} = \mathbf{x}_{n_m} + \epsilon \text{ sign}(\nabla_{\mathbf{x}_{n_m}} \ell(\theta_m, \mathbf{x}_{n_m}, y_{n_m}))$
 - 5: Minimise $\ell(\theta_m, \mathbf{x}_{n_m}, y_{n_m}) + \ell(\theta_m, \mathbf{x}'_{n_m}, y_{n_m})$ w.r.t. θ_m \triangleright *adversarial training*
-

<https://arxiv.org/abs/1612.01474>



Uncertainty in Deep Learning

Epistemic and Aleatoric Uncertainty



<https://arxiv.org/abs/1703.04977>



Uncertainty in Deep Learning

Epistemic and Aleatoric Uncertainty

$$\text{Var}(x) \approx \frac{1}{T} \sum_{t=1}^T \hat{x}_t^2 - \left(\frac{1}{T} \sum_{t=1}^T \hat{x}_t \right)^2 + \frac{1}{T} \sum_{t=1}^T \hat{\sigma}_t^2$$

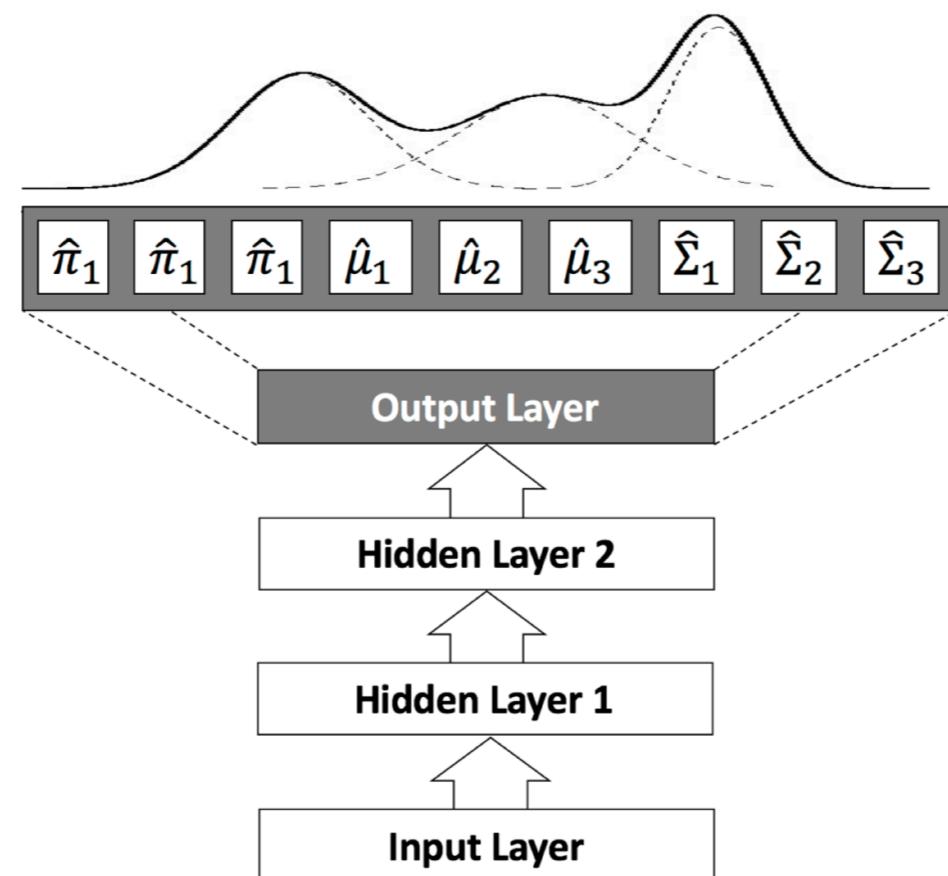
with $\{\hat{x}_t, \hat{\sigma}_t^2\}_{t=1}^T$ a set of T sampled outputs: $\hat{x}_t, \hat{\sigma}_t^2 = f^{\widehat{\mathbf{W}}_t}(I)$ for randomly masked weights $\widehat{\mathbf{W}}_t \sim q(\mathbf{W})$.

<https://arxiv.org/abs/1703.04977>



Uncertainty in Deep Learning

Uncertainty-Aware Learning from Demonstration

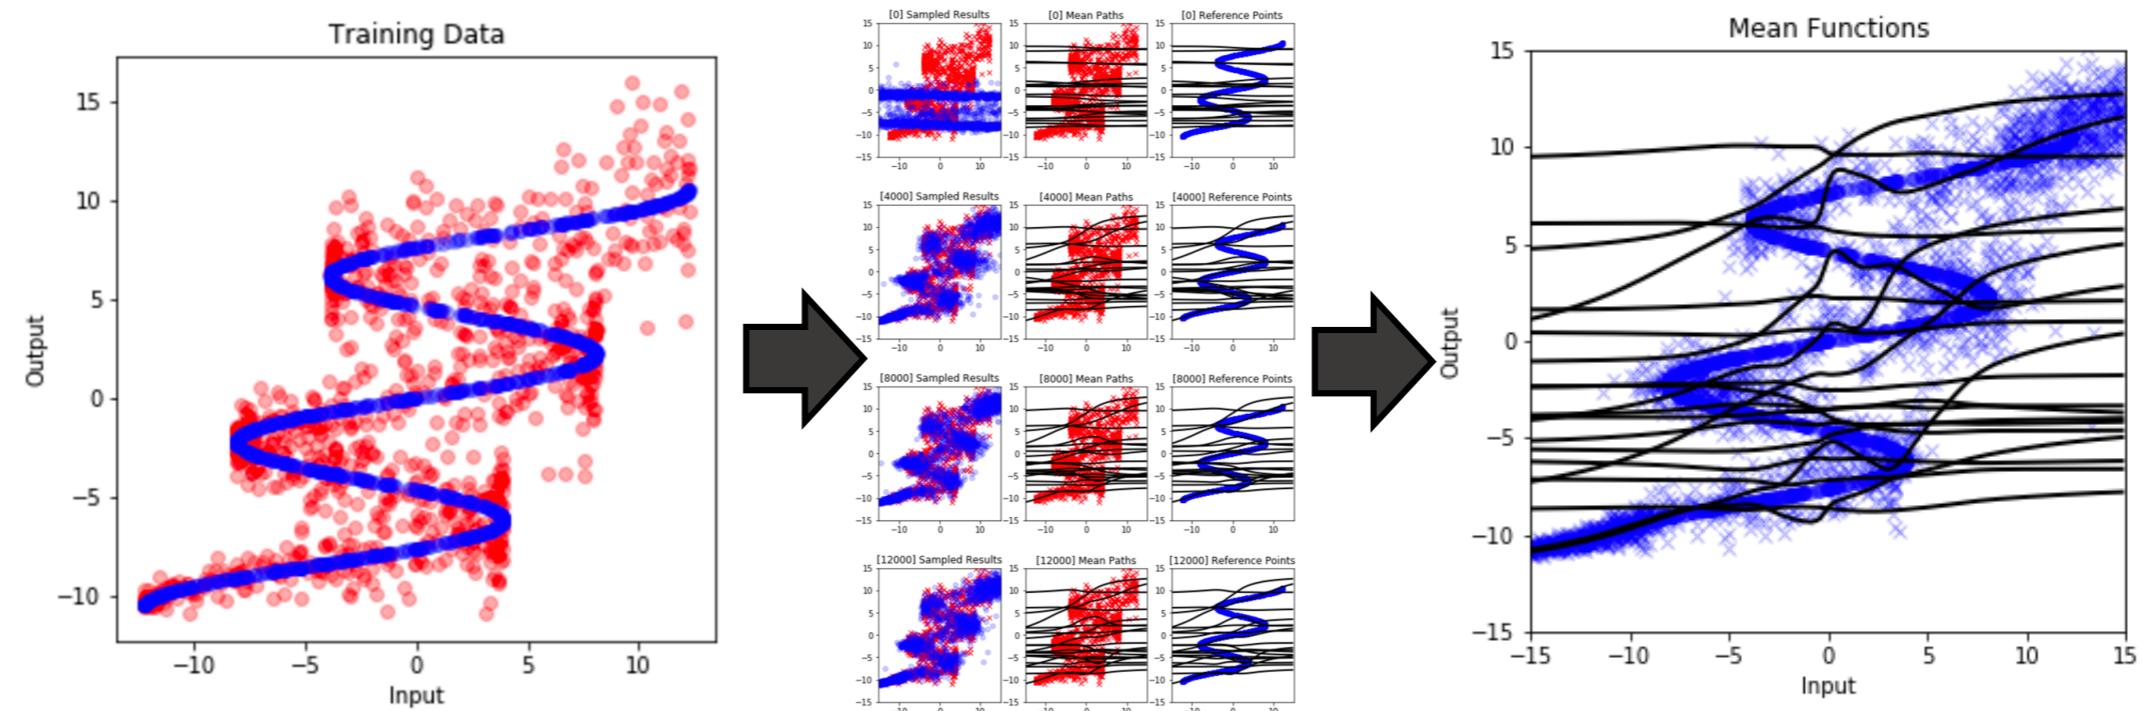


<https://arxiv.org/abs/1709.02249>



Uncertainty in Deep Learning

Uncertainty-Aware Learning from Demonstration



<https://arxiv.org/abs/1709.02249>



Uncertainty in Deep Learning

Uncertainty-Aware Learning from Demonstration

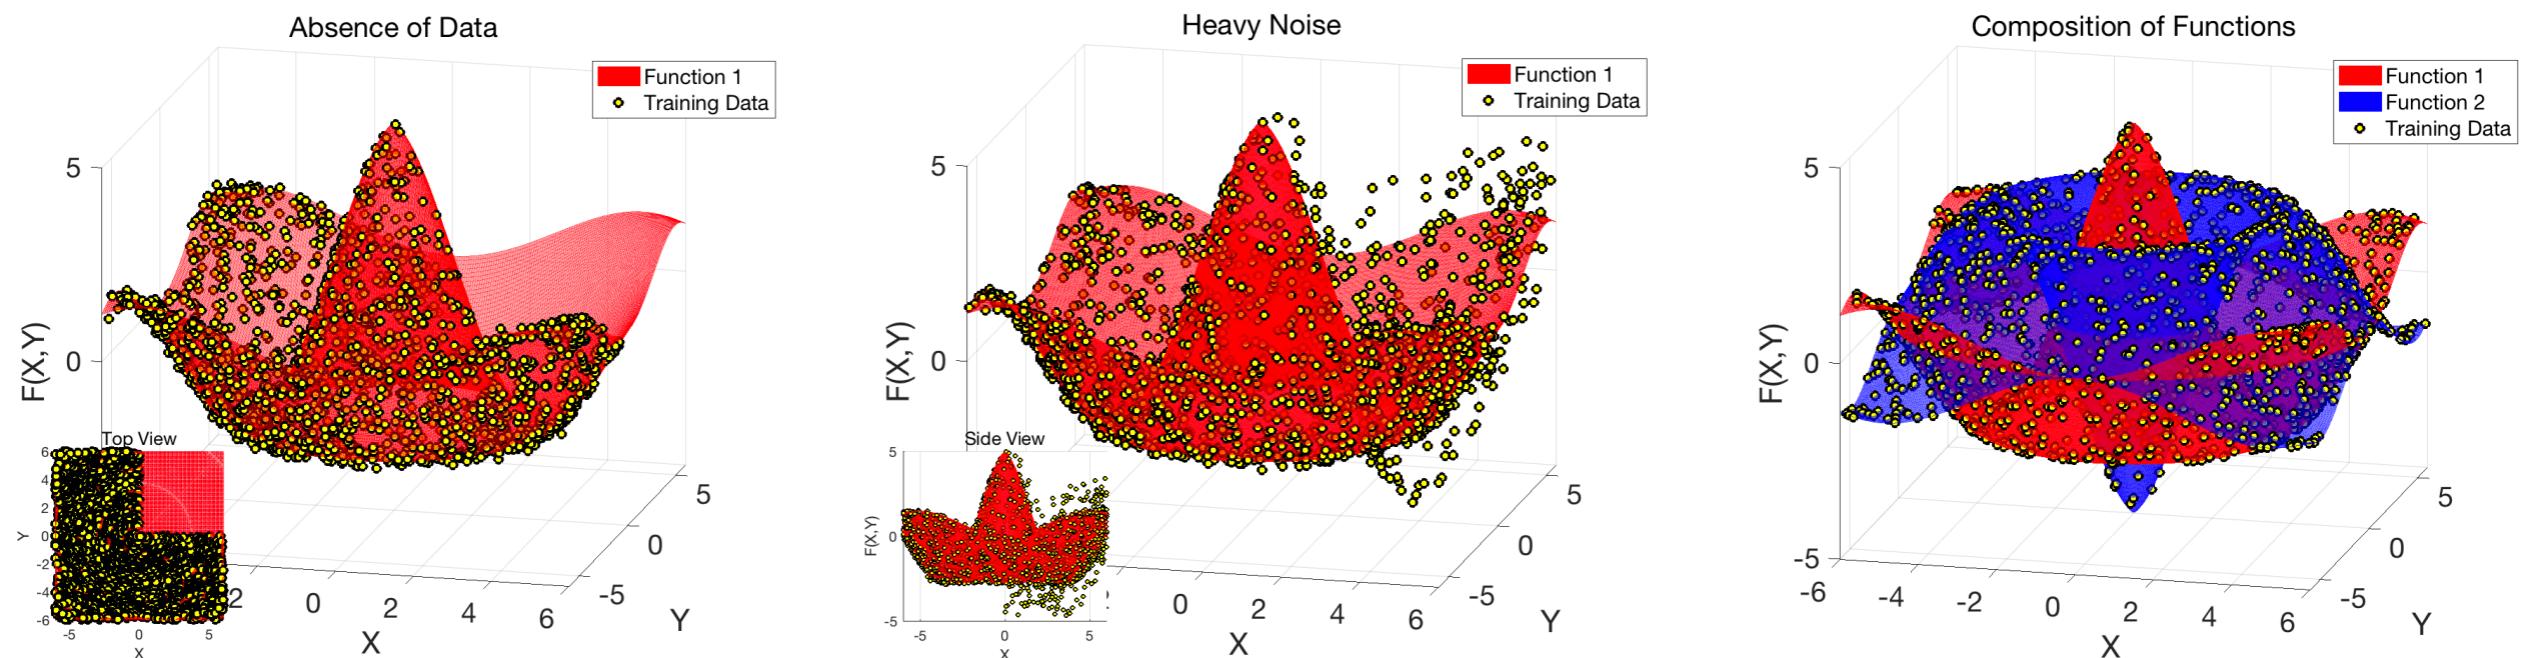
$$\begin{aligned} \mathbb{V}(\mathbf{y}|\mathbf{x}) &= \underbrace{\sum_{j=1}^K \pi_j(\mathbf{x}) \Sigma_j(\mathbf{x})}_{\text{unexplained variance}} \\ &\quad + \underbrace{\sum_{j=1}^K \pi_j(\mathbf{x}) \left\| \mu_j(\mathbf{x}) - \sum_{k=1}^K \pi_k(\mathbf{x}) \mu_k(\mathbf{x}) \right\|^2}_{\text{explained variance}}. \end{aligned}$$

<https://arxiv.org/abs/1709.02249>



Uncertainty in Deep Learning

Uncertainty-Aware Learning from Demonstration

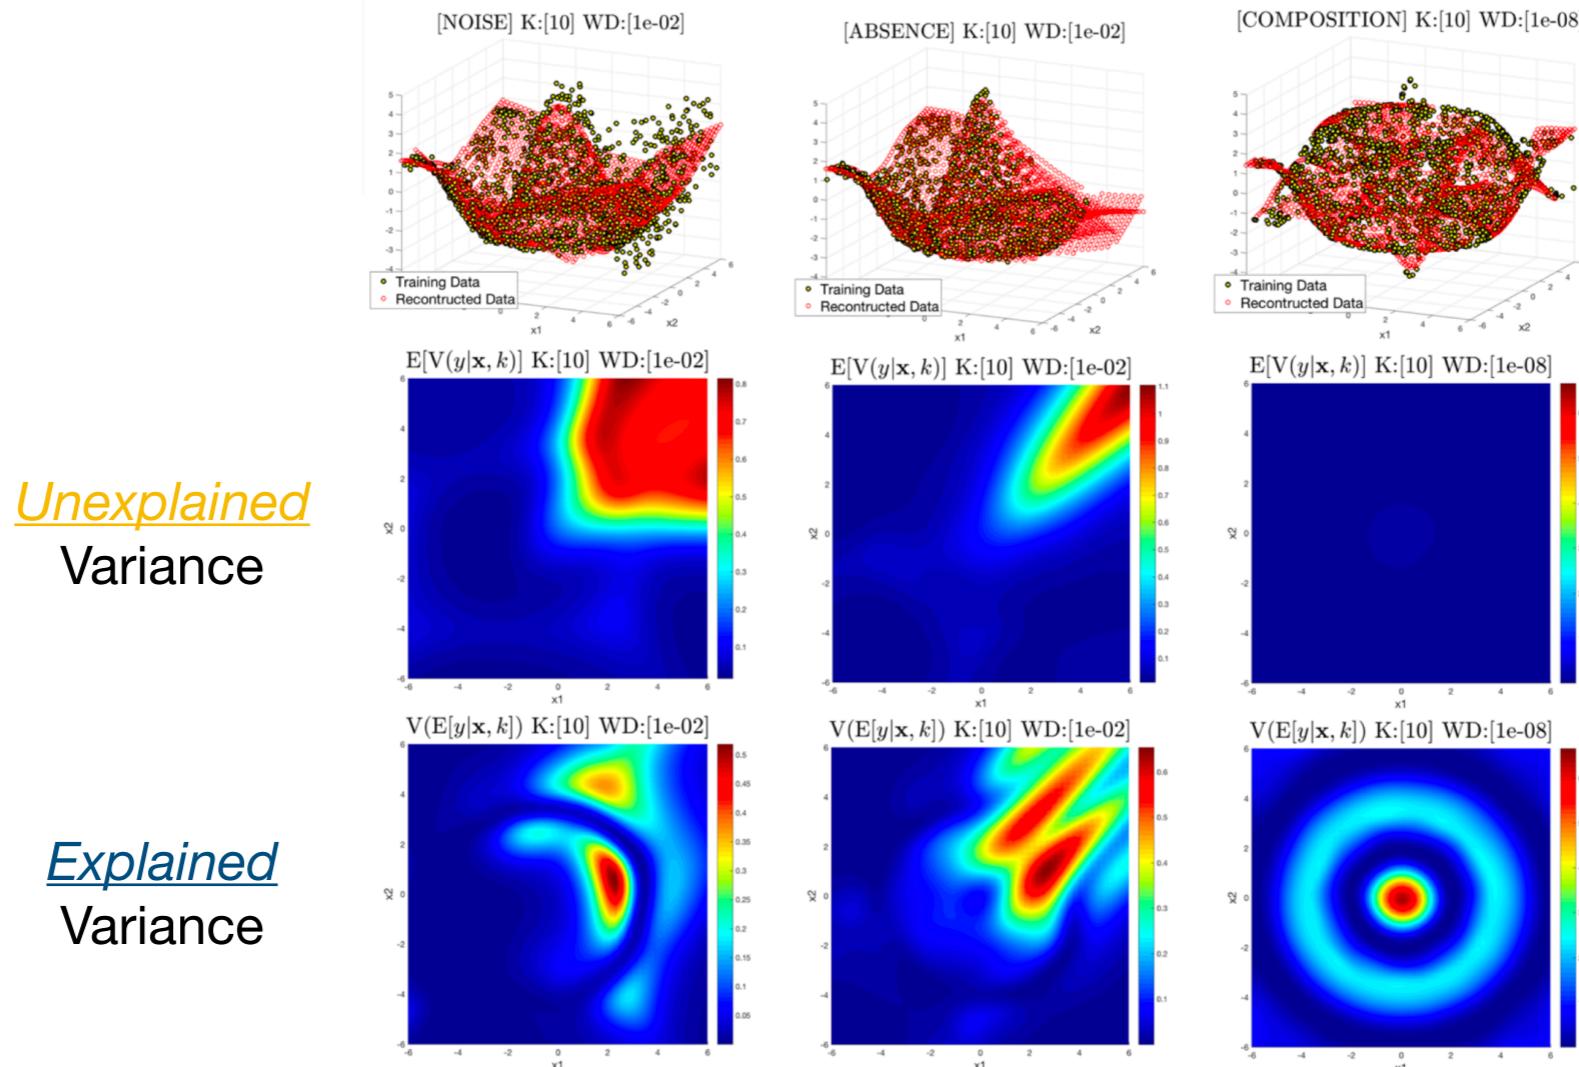


<https://arxiv.org/abs/1709.02249>



Uncertainty in Deep Learning

Uncertainty-Aware Learning from Demonstration



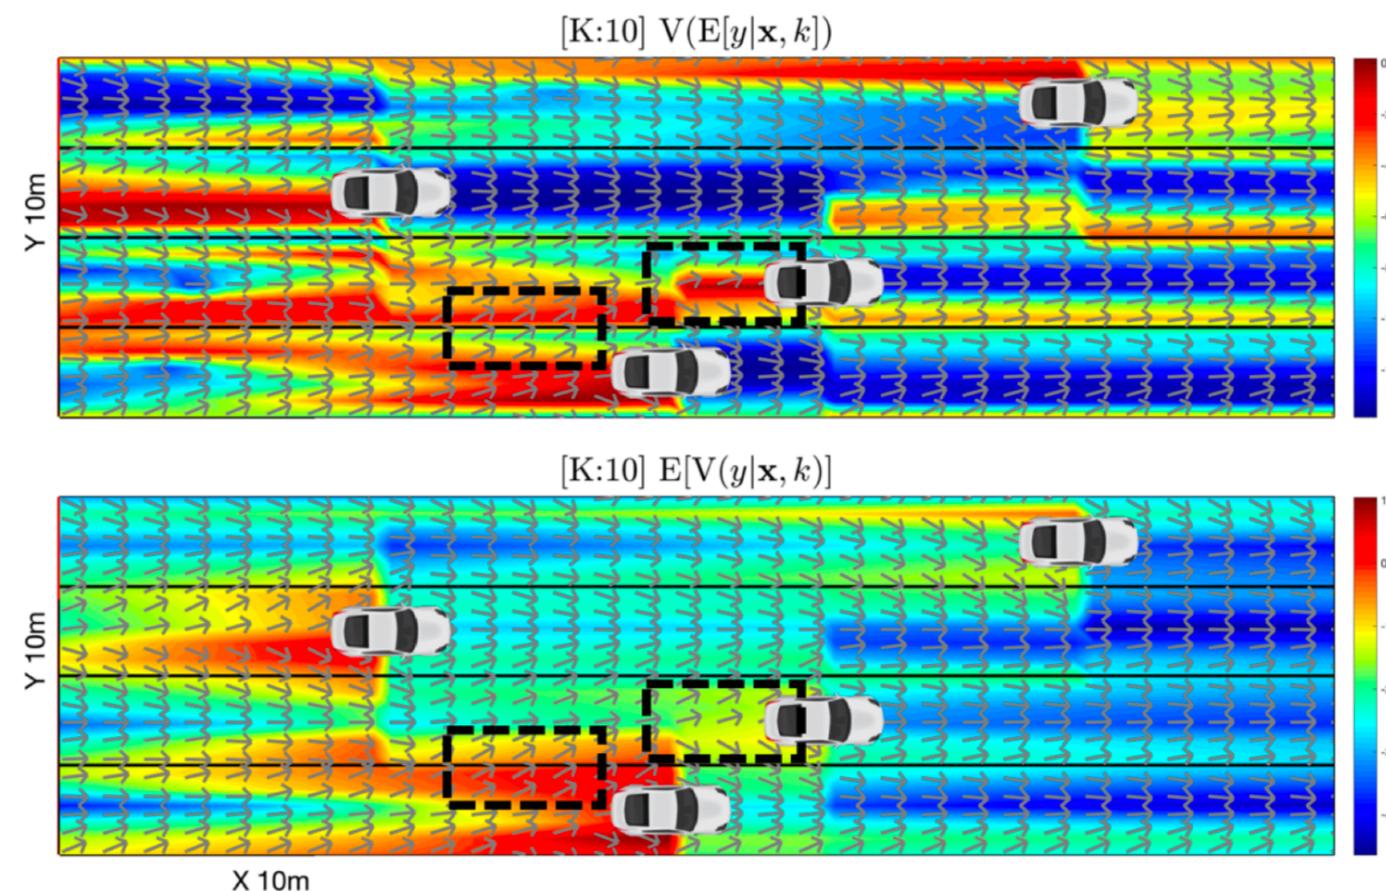
<https://arxiv.org/abs/1709.02249>



Uncertainty in Deep Learning

Uncertainty-Aware Learning from Demonstration

Explained
Variance



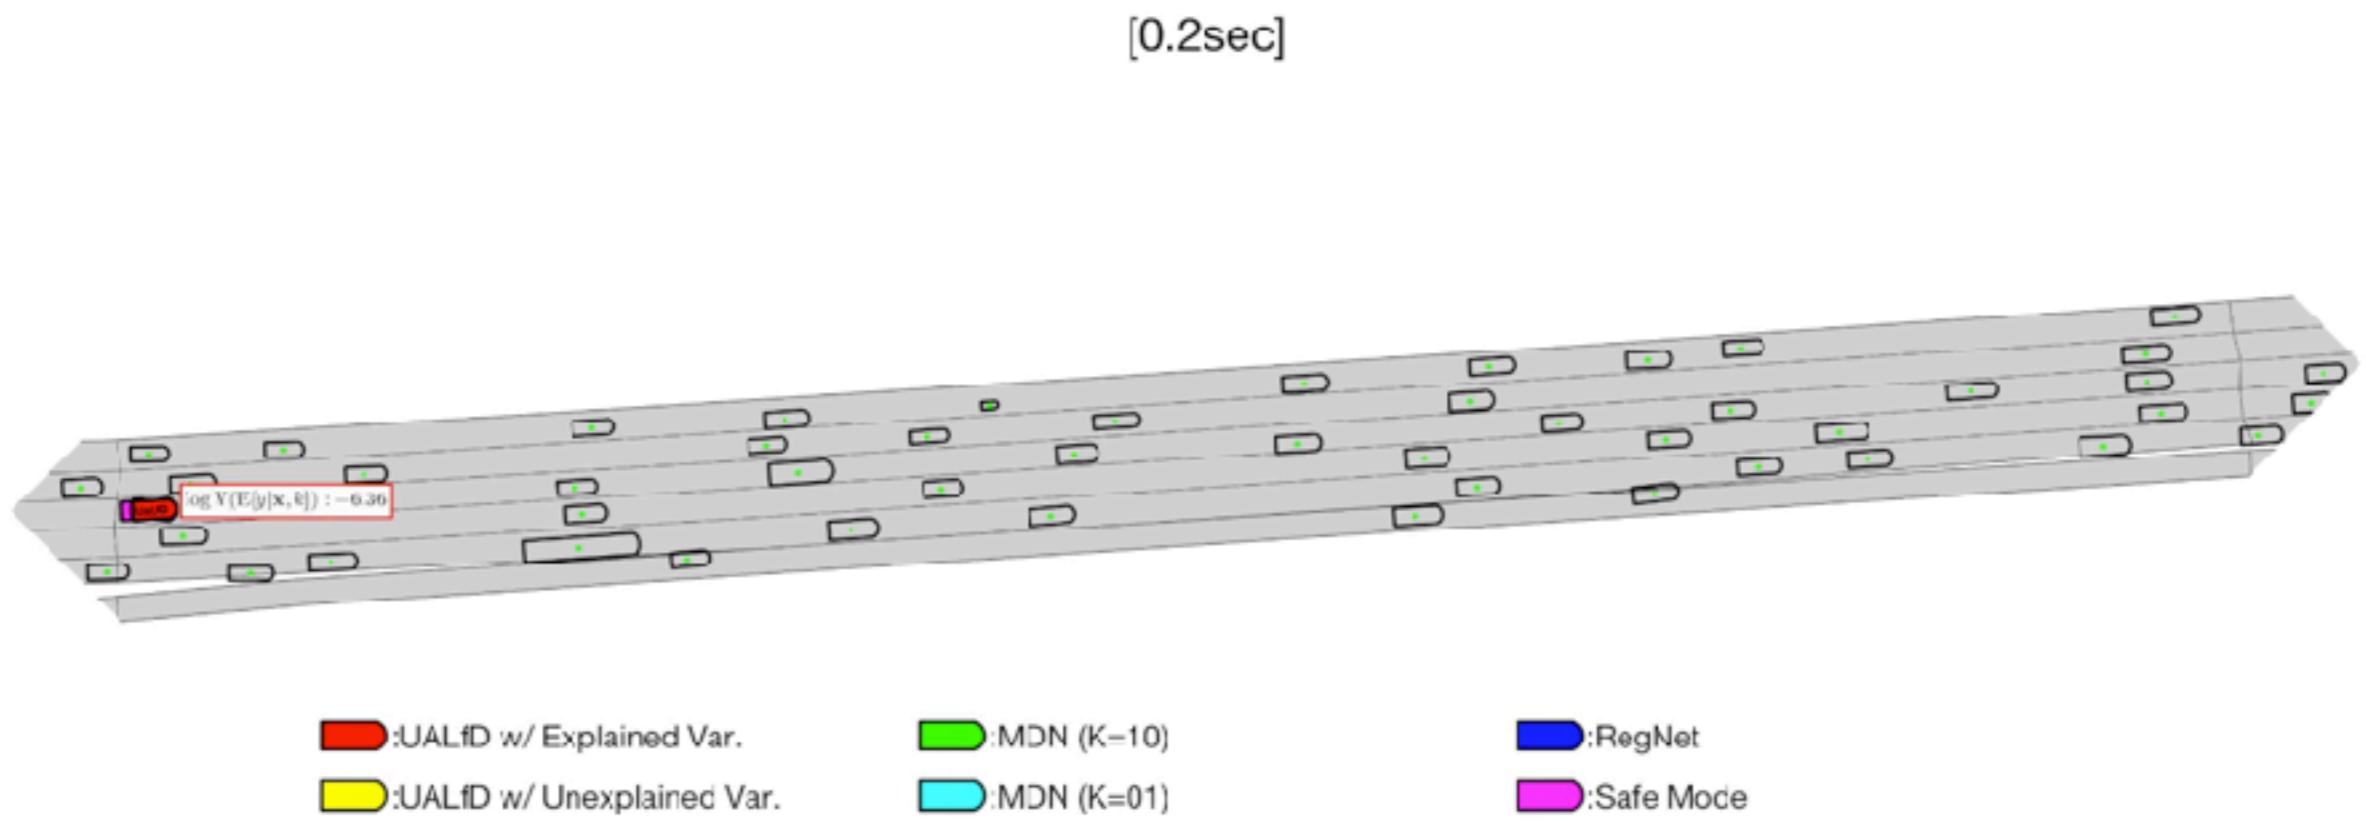
Unexplained
Variance

<https://arxiv.org/abs/1709.02249>



Uncertainty in Deep Learning

Uncertainty-Aware Learning from Demonstration



<https://arxiv.org/abs/1709.02249>



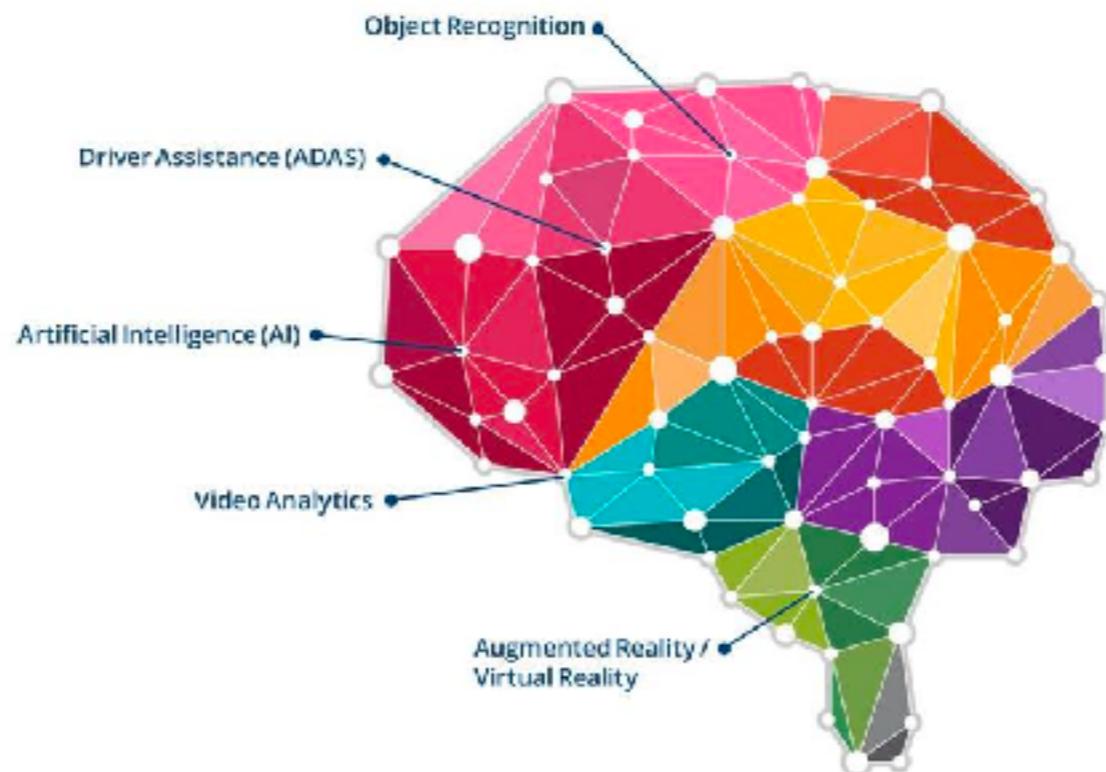
Conclusion

No more ImageNet challenge



Conclusion

Focus on applications



Conclusion

Consider real problems





CPS
LAB

Thank you for your attention.

Contact information (sungjoon.choi@cplab.snu.ac.kr)

